



Scheduling

- Scheduling is the “problem of allocating scarce resources to activities over time.” [Baker 1974]
- Typically, planning is deciding *what* to do, and scheduling is deciding *when* to do it.
- Generally, scheduling is determining how to accomplish some set of tasks while:
 - obeying task constraints (e.g., due dates, task interactions) [feasibility criteria]
 - minimizing resource usage (e.g., time, personnel) [optimization criteria]

CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Feasibility

- Overconstrained:
 - Sometimes not all constraints can be honored (e.g., *soft* constraints)
 - decide which constraints to *relax* then decide when to do the tasks
- Oversubscribed:
 - Sometimes not all tasks can be accommodated given fixed resources.
 - decide *which* tasks to discard then decide when to do them

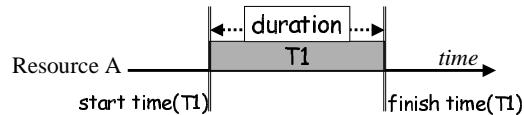
CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Task Characteristics

- Non-preemptive:
 - once started, must be allowed to finish.



- Preemptive:
 - can be interrupted and re-started later.

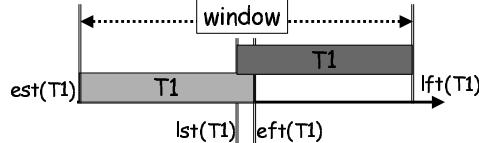
CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Task Characteristics (cont.)

- Time windows:
 - must be scheduled within some interval



- release dates: beginning of window
- deadlines: end of window
- slack: LST - EST

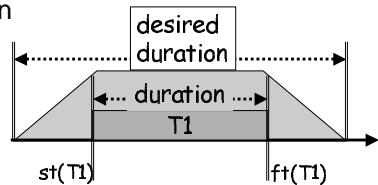
CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Task Characteristics (cont.)

- Processing time:
 - fixed duration
 - flexible: penalty function for less than desired duration



CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Task Characteristics (cont.)

- Precedence Constraints:
 - Task A before Task B
 - $start(B)-start(A) \geq 0$
- Resource Constraints:
 - must be on a particular resource or particular type of resource
- Alternatives
 - resources
 - times

CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Task Characteristics (cont.)

- Set-up Times
 - task dependent interaction effects
 - depending on what preceded task in schedule, it may take extra time to prepare resource to execute the task
 - e.g., making different colored plastics in a vat, extruding different viscosity materials, re-orienting antenna to receive signal

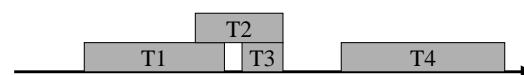
CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Resource Characteristics: Capacity

- Unit capacity: one task at a time
- Multi-capacity:
 - n tasks at a time
 - certain sum of task size requirements at a time



CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Resource Characteristics: Capacity Types

- consumable:
 - monotonically decreasing capacity
- renewable:
 - variable (up and down) capacity
- reusable:
 - fixed, unchangable capacity

CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Optimization Metrics

- Resource based
 - minimize number of resources used/cost of resources allocated
 - maximize resource utilization: percentage of time that resources are not idle
 - minimum peak resource usage
- Task based
 - cumulative value of tasks scheduled
 - cumulative value of time slots for tasks
 - maximum weighted # of tasks

CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Optimization Metrics (cont.)

- Time based
 - Makespan: duration between start of first task and completion of last task
 - Total Weighted Flow Time: sum of weighted finish times
 - Maximum Tardiness: largest delay from desired finish time to actual end time for any task
 - Total Weighted Tardiness: Sum of the delays between desired and actual finish times
 - Total Weighted Number of Late/Discarded Tasks

CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Scheduling Problem

- Given:
 - set of resources
 - set of tasks with constraints, temporal descriptions and resource requirements
 - an objective function
- find a solution (mapping of tasks to times on resources) that
 - satisfies constraints
 - minimizes objective function

CS 540, Artificial Intelligence

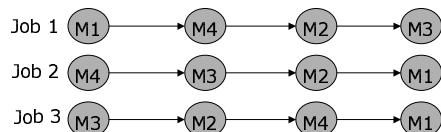
© Adele Howe, Spring 2004



Benchmark Application: Manufacturing Scheduling

Job Shop (JSP):

- n jobs on m machines
- each job must be processed on each machine exactly once for a fixed duration in a pre-defined order (job routing order).
- unit capacity, non-preemptive, no time windows or setups
- minimize makespan



CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



JSP (cont.)

- Solution:
 - is processing order for all jobs on each machine.
 - specifies est for each job/machine s.t. precedence and resource constraints are satisfied.
- $(n!)^m$ possible solutions of which a subset are feasible.

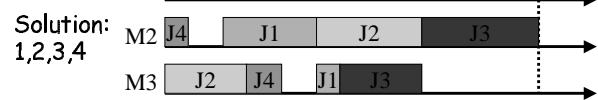
CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Example JSP

	O _{i,1}		O _{i,2}		O _{i,3}	
Job	Dur	Rout	Dur	Rout	Dur	Rout
1	5	1	8	2	2	3
2	7	3	3	1	9	2
3	1	1	7	3	10	2
4	2	2	3	3	1	1

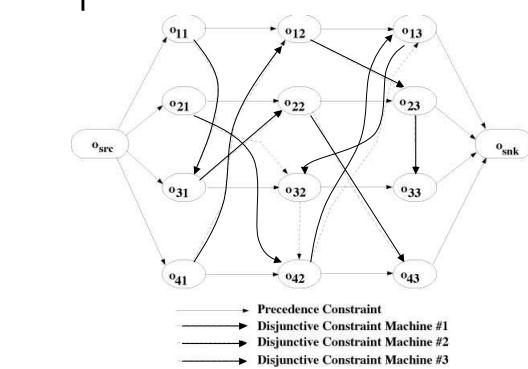


CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Example JSP (cont.)



CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



JSP Taxonomy of Schedules

Feasible Solution: a job processing order on each machine without any cyclic dependencies

Feasible Schedule: specifies a start time s.t. all constraints are satisfied

Inadmissible: operations start later than their est

Semi-Active: some operations are scheduled at their est

Global left shift: move an operation earlier into machine idle time

Active: no global left shifts are possible (difficult to determine)

Non-Delay: no machine is idle if an operation is available

Optimal: minimizes the makespan

CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Algorithms for Scheduling

- Heuristic Constructive (e.g., LDS, HBSS)
- Constraint Programming (e.g., dynamic programming, edge-finding)
- Local Search
- Genetic Algorithms

[Note: makespan computation tends to be most expensive aspect]

CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Slack Based Heuristics

- order operations in increasing (LST – EST)
- greedy version:
 - until all operations scheduled:
 - schedule open operation that has predecessors scheduled and minimizes slack at its EST
 - update ESTs and LSTs
- minslack: impose ordering on pairs of unsequenced tasks with minimal maximal slack:
 - $(\max(lft(A)-est(B), lft(B)-est(A)) - d(A) - d(B))$

CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Texture Based Heuristics

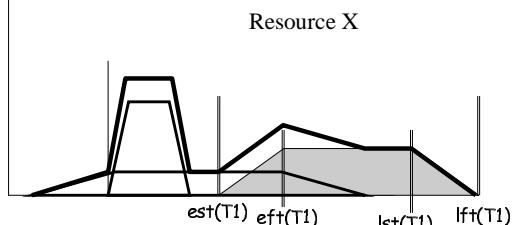
- construct a profile of resource usage over time to identify and reduce contention
- SumHeight:
 - estimate probabilistically the competition of tasks on resources for each time
 - identify the point of highest contention
 - make ordering decision to reduce the contention

CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Texture Based Heuristics (cont.)



black lines are task demands, red line is aggregate demand

CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Constraint Programming

Edge Finding:

- determining whether an operation can, cannot, or must execute before (or after) other operations on a machine.
- Rules: (O is set of operations on machine, A is a particular operation, D is duration)

$$\forall O, \forall A \notin O, LFT_{O \cup \{A\}} - EST_O < \sum_{i \in O} D_i + D_A \Rightarrow A \ll O$$

$$\forall O, \forall A \notin O, LFT_O - EST_{O \cup \{A\}} < \sum_{i \in O} D_i + D_A \Rightarrow A \gg O$$

$$A \ll O \Rightarrow end(A) \leq \min_{i \in O} (LFT_i - D_i)$$

$$A \gg O \Rightarrow start(A) \geq \max_{i \in O} (EST_i + D_i)$$

CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Edge-Finding (cont.)

1. Compute a schedule (S) by iteratively applying priority rule:
 - whenever a resource is free and an operation is available, schedule the operations with the smallest LFT
2. Given S, for each operation A,
 1. compute the set of operations (P) not finished at $EST(A)$.
 2. For each P (in decreasing order of $LFT(P)$), update ordering and start/end times based on rules.

CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Local Search

- state of the art [Nowicki & Smutnicki 2003]
- solutions represented as disjunctive graph with edges only between adjacent operations
- JSP specific move operators (N5 for N&S03)
- initiate search from random semi-active or greedy solutions
- tabu search: checks for cycles in recent solutions (tabu tenure=8) and re-starts if cycles or if no progress over some time
- path relinking to generate new solutions for re-starts

CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



An Aside: Critical Operations, Paths and Blocks

Critical Operation: operations in which $est=lst$.

If start time is delayed, makespan will increase.

Critical Path: contiguous sequence of critical operations starting at $t=0$ and ending at makespan with no machine idle time.

Critical Blocks: subsequences of critical paths with all operations on the same machine.

CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Critical Operations, Paths and Blocks Example

Greedy Solution: 4,3,2,1



Critical operations:

J12, J13, J22, J32, J33, J42, J43

Critical path:

J42, J43, J33, J32, J22, J12, J13

Critical blocks:

{J43, J33}, {J32, J22, J12}

CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Critical Blocks Based Move Operators

■ N1:

- invert the order of a pair of adjacent operations on the same critical block
- focuses on only 1) feasible neighbors and 2) those operations that can change the makespan
- fully connected search space: there exists a path from any feasible solution to a global optimal



CS 540, Artificial Intelligence

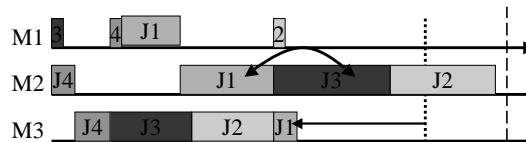
© Adele Howe, Spring 2004



Critical Blocks Based Move Operators (cont.)

■ N5, like N1 except

- swap only adjacent operations that include either first or last operations in the critical block
- do not swap either the first two operations in the first critical block or the last two operations in the last critical block
- not completely connected

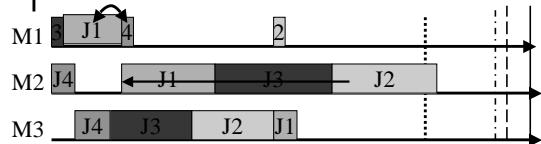


CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Effect of Moves (cont.)



Recall the optimal...



CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Genetic Algorithms

- Represent solution as permutation of tasks
- Requires a schedule builder to convert solution into schedule and assess objective function. [Indirect search]
- Syswerda's permutation crossover was used for scheduling.

CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



Variations on JSP

- WorkFlow:
 - machines are partitioned into wf subsets
 - every job must be processed on all machines in a partition before moving to the next partition
- Flow Shop:
 - all jobs have identical order on all machines
 - workflow with $wf = \# \text{ machines}$

CS 540, Artificial Intelligence

© Adele Howe, Spring 2004



JSP Difficulty

- Analysis:
 - Optimization is NP-hard
 - No polynomial time algorithm can guarantee a solution within 20% of optimal makespan.
- Observations:
 - Given fixed n and m , workflow JSPs are typically more difficult than random JSPs.
 - Given fixed n and m , flowshop JSPs are typically more difficult than workflow JSPs.
 - Given fixed m and wf , square JSPs are typically more difficult than rectangular JSPs (more jobs than machines).
 - Given fixed n, m and wf , relative problem difficulty tends to be algorithm independent.

CS 540, Artificial Intelligence

© Adele Howe, Spring 2004