


An introduction to CP Optimizer

Philippe Laborie
IBM Analytics, Decision Optimization



May 25, 2016

Outline


- 1) Overview of CP Optimizer
- 2) Modeling concepts
- 3) Automatic search
- 4) Development tools

Overview of CP Optimizer

- A component of **IBM ILOG CPLEX Optimization Studio**
- A **Constraint Programming** engine for combinatorial problems (including **scheduling problems**)
- Implements a **Model & Run** paradigm (like CPLEX)
 - Model: **Concise yet expressive** modeling language
 - Run: **Powerful automatic search procedure**
Search algorithm is **complete**
- Available through the following interfaces:
 - OPL
 - C++ (native interface)
 - Python, Java, .NET (wrapping of the C++ engine)
- Set of **tools** to support the development of efficient models


Overview of CP Optimizer

- First version in 2007/2008
- Search “CP Optimizer” on Google Scholar to get an idea of how CP Optimizer is being used across academy and industry



Overview of CP Optimizer

- Model & Run paradigm



Overview of CP Optimizer

IBM ILOG CPLEX Optimization Studio

File Edit Navigate Search Run Window Help

Quick Access

sched_rcpsp.mod

```

1 using CP;
2 // Classical Resource-Constrained Project Scheduling Problem (RCPSP)
3 int NbTasks = ...;
4 int NbRsrcs = ...;
5 range RsrcIds = 0..NbRsrcs-1;
6 int Capacity[r in RsrcIds] = ...;
7 tuple Task {
8   key int id;
9   int pt;
10  int dmds[RsrcIds];
11  {int} succs;
12 }
13 {Task} Tasks = ...;

15 dvar interval itvs[t in Tasks] size t.pt;
16

17 cumulFunction rsrcUsage[r in RsrcIds] =
18   sum (t in Tasks: t.dmds[r]>0) pulse(itvs[t], t.dmds[r]);
19

20 minimize max(t in Tasks) endOf(itvs[t]);
21 subject to {
22   forall (r in RsrcIds)
23     rsrcUsage[r] <= Capacity[r];
24   forall (t1 in Tasks, t2id in t1.succs)
25     endBeforeStart(itvs[t1], itvs[<t2id>]);
26 }

```


Value for itvs

Tasks (size 122)	Start	End	Size
[<1>]	0	0	0
[<2>]	0	10	10
[<3>]	0	10	10
[<4>]	4	10	6
[<5>]	14	20	6
[<6>]	10	17	7
[<7>]	10	15	5
[<8>]	17	20	3

Value for rsrcUsage

RsrcIds (size 4) | Value

Chart mode: Summed



Engine log

```

! -----  

! Search terminated normally, 6 solutions found.  

! Best objective : 103 (optimal - effective tol. is 0)  

! Number of branches : 185,324  

! Number of fails : 75,620  

! Total memory usage : 8.5 MB (8.0 MB CP Optimizer + 0.6 MB Concert)  

! Time spent in solve : 1.12s (1.12s engine + 0.00s extraction)  


! Search speed (br. / s) : 164,586.1  

! -----

```

Statistics

Objective Solution Proven solution




Time (seconds)

00:00:01:18

Overview of CP Optimizer

Interval variables



The screenshot shows the IBM ILOG CPLEX Optimization Studio interface. On the left, the code editor displays a model file named `sched_rcpsp.mod` for a Classical Resource-Constrained Project Scheduling Problem (RCPSP). A specific line of code, `dvar interval itvs[t in Tasks] size t.pt;`, is highlighted with a red oval. The interface includes several windows:

- Value for itvs:** A Gantt chart showing task intervals. A red oval highlights the first few tasks.
- Value for rsrcUsage:** A chart showing resource usage over time for four resources.
- Statistics:** A plot showing the objective value (green dashed line) and solution progress (yellow squares) over time (seconds).
- Engine log:** A window displaying solver logs, indicating 6 solutions found and various performance metrics.


```

1 using CP;
2 // Classical Resource-Constrained Project Scheduling Problem (RCPSP)
3 int NbTasks = ...;
4 int NbRsrcs = ...;
5 range RsrcIds = 0..NbRsrcs-1;
6 int Capacity[r in RsrcIds] = ...;
7 tuple Task {
8   key int id;
9   int pt;
10  int dmds[RsrcIds];
11  {int} succs;
12 }
13 {Task} Tasks = ...;
14
15 dvar interval itvs[t in Tasks] size t.pt;
16
17 cumulFunction rsrcUsage[r in RsrcIds] =
18   sum (t in Tasks: t.dmds[r]>0) pulse(itvs[t], t.dmds[r]);
19
20 minimize max(t in Tasks) endOf(itvs[t]);
21 subject to {
22   forall (r in RsrcIds)
23     rsrcUsage[r] <= Capacity[r];
24   forall (t1 in Tasks, t2id in t1.succs)
25     endBeforeStart(itvs[t1], itvs[<t2id>]);
26 }

```

Overview of CP Optimizer

Functions



The screenshot shows the IBM ILOG CPLEX Optimization Studio interface. On the left, the code editor displays a model named `sched_rcpsp.mod` for a Classical Resource-Constrained Project Scheduling Problem (RCPSP). The code includes declarations for tasks, resources, and constraints, such as calculating resource usage and minimizing the end-of-task value.

The interface features several windows:

- Value for itvs**: A window showing a Gantt chart and a table of task details. The table includes columns for Task ID, Start, End, and Size. The Gantt chart visualizes the start and end times of tasks across multiple resources.
- Value for rsrcUsage**: A window showing a resource usage chart over time. The chart displays the cumulative usage of four resources (RsrcIds) over 100 time units, with resource capacities represented by horizontal bars at the bottom.
- Engine log**: A window displaying the solver's log output, which includes information about the search process, number of solutions found, and performance metrics.
- Statistics**: A window showing a plot of the objective function value versus solve time. The plot includes a green dashed line for the objective value, yellow squares for intermediate solutions, and red diamonds for proven solutions.

Modeling concepts


- Two main types of decision variables:
 - Integer variables
 - Interval variables

Modeling concepts (integer variables)

Variables	Expressions	Constraints
<p>Variables are <i>discrete integer</i></p> <p>Domains can be specified as a range [1..50] or as a set of values {1, 3, 5, 7, 9}</p> <p>dvar int x in 1..50</p>	<p>Expressions can be integer or floating-point, for example <code>0.37*y</code> is allowed</p> <p>Basic arithmetic (<code>+, -, *, /</code>) and more complex operators (<code>min, max, log, pow etc.</code>) are supported</p> <p>Relational expressions can be treated as 0-1 expressions. e.g. <code>x = (y < z)</code></p> <p>Special expressions:</p> <ul style="list-style-type: none"> <code>x == a[y]</code> <code>x == count(Y, 3)</code> <code>y == cond ? y : z</code> 	<p>Rich set of constraints:</p> <p>Standard relational constraints (<code>==, !=, <, >, <=, >=</code>)</p> <p>Logical combinators (<code>&&, , !, =></code>)</p> <p>Specialized (global) constraints</p> <ul style="list-style-type: none"> <code>allDifferent(X)</code> <code>allowedAssignments(X, tuples)</code> <code>forbiddenAssignments(X, tuples)</code> <code>pack(load, container, size)</code> <code>lexicographic(X, Y)</code> <code>inverse(X, Y)</code>

Modeling concepts (interval variables for scheduling)

- Scheduling (our definition of):
 - Scheduling consist of assigning starting and completion times to a set of activities while satisfying different types of constraints (resource availability, precedence relationships, ...) and optimizing some criteria (minimizing tardiness, ...)



- Time is considered as a **continuous** dimension: domain of possible start/completion times for an activity is potentially very large
- Beside start and completion times of activities, **other types of decision variables** are often involved in real industrial scheduling problems (resource allocation, optional activities ...)

Modeling concepts (interval variables for scheduling)

- Extension of classical CSP with a new type of decision variable:
optional interval variable :

$$\text{Domain}(x) \subseteq \{\perp\} \cup \{ [s,e] \mid s, e \in \mathbb{Z}, s \leq e \}$$

Absent interval Interval of integers

- Introduction of mathematical notions such as **sequences** and **functions** to capture temporal aspects of scheduling problems

Modeling concepts (interval variables for scheduling)


- In scheduling models, interval variables usually represent an interval of time whose end-points (start/end) are decision variables of the problem
- Examples:
 - A production order, an operation in a production order
 - A sub-project in a project, a task in a sub-project
 - A batch of operations
 - The setup of a tool on a machine
 - The moving of an item by a transportation device
 - The utilization interval of a machine
 - The filling or emptying of a tank
- Idea of the model (and search) is to avoid the enumeration of start/end values

Modeling concepts (interval variables for scheduling)

- An interval variable can be **optional** meaning that it is a decision to have it present or absent in a solution.
- Examples:
 - Unperformed tasks and optional sub-projects
 - Alternative resources, modes or recipes for processing an order, each mode specifying a particular combination of operational resources
 - Operations that can be processed in different temporal modes (e.g. series or parallel), left unperformed or externalized
 - Activities that can be performed in an alternative set of batches or shifts


Example: Resource Constrained Project Scheduling Problem

- RCPSP: a very classical academical scheduling problem
 - Tasks a_i with fixed processing time P_i
 - Precedence constraints
 - Discrete resources with limited instantaneous capacity R_k
 - Tasks require some quantity of discrete resources
 - Objective is to minimize the schedule makespan



Example: Resource Constrained Project Scheduling Problem

- RCPSP: Standard time-indexed MIP formulation



Standard RCPSP (DT: Discrete Time)

$$\text{minimize} \sum_{t \in H} t x_{nt}$$

$$\sum_{t \in H} x_{it} = 1 \quad \forall i \in \mathcal{A}$$

$$\sum_{t \in H} t x_{it} + P_i \leq \sum_{t \in H} t x_{jt} \quad \forall (i, j) \in \mathcal{P}$$

$$\sum_{i \in \mathcal{A}, t \leq \tau < t + P_i} Q_{ik} x_{it} \leq R_k \quad \forall \tau \in H, \forall k \in \mathcal{R}$$

$$x_{it} \in \{0, 1\} \quad \forall i \in \mathcal{A}, \forall t \in H$$

Example: Resource Constrained Project Scheduling Problem

- Basic CP Optimizer model for RCPSP:




```
dvar interval a[i in Tasks] size i.pt;  
  
cumulFunction usage[r in Resources] =  
    sum (i in Tasks: i.qty[r]>0) pulse(a[i], i.qty[r]);  
  
minimize max(i in Tasks) endOf(a[i]);  
subject to {  
    forall (r in Resources)  
        usage[r] <= Capacity[r];  
    forall (i in Tasks, j in i.succs)  
        endBeforeStart(a[i], a[<j>]);  
}
```

Example: Resource Constrained Project Scheduling Problem

- Comparison of this time-indexed MIP formulation against the CP Optimizer model on a set of:
 - 300 classical **small** RCPSP instances (30-120 tasks) +
 - 40 slightly **more realistic** larger ones (900 tasks)
 - time-limit: 2mn, 4 threads
- Note: industrial scheduling problems are often much **larger**, typically several 1.000 tasks (we handled up to 1.000.000 tasks in an RCPSP-like scheduling application in V12.6)


Example: Resource Constrained Project Scheduling Problem

- Comparison of CP Optimizer and MIP performance on RCPSP



Example: Job-shop Scheduling Problem

- Example: Job-shop Scheduling Problem



- Minimization of makespan

Example: Job-shop Scheduling Problem


- CP Optimizer model for Job-shop:

```
dvar interval op[j in Jobs][p in Pos] size Ops[j][p].pt;
dvar sequence mchs[m in Mchs] in
  all(j in Jobs, p in Pos : Ops[j][p].mch == m) op[j][p];

minimize max(j in Jobs) endOf(op[j][nbPos-1]);
subject to {
  forall (m in Mchs)
    noOverlap(mchs[m]);
  forall (j in Jobs, p in 1..nbPos-1)
    endBeforeStart(op[j][p-1], op[j][p]);
}
```

Example: Job-shop Scheduling Problem

- Properties:
 - Complexity is **independent of the time scale**
 - CP Optimizer is able to reason **globally** over a sequence variable
 - **Avoid quadratic models** over each pair (i,j) of intervals in the sequence
- Compare:
 - Quadratic disjunctive MIP formulation with big-Ms:
$$b[i][j] \in \{0,1\} // b[i][j]=1 \text{ iff } a[i] \text{ before } a[j]$$
$$\text{end}[i] \leq \text{start}[j] + M*(1-b[i][j])$$
$$\text{end}[j] \leq \text{start}[i] + M*b[i][j]$$
 - CP Optimizer model:




Example: Job-shop Scheduling Problem

- Comparison of the CP Optimizer model vs a disjunctive MIP formulation on a set of 140 classical Job-shop instances (50-2000 tasks), time-limit: 2mn, 4 threads

Example: Job-shop Scheduling Problem

- Comparison of CP Optimizer and MIP performance on Job-Shop



Example: Flexible Job-shop Scheduling Problem


- CP Optimizer model

```
dvar interval ops [Ops];
dvar interval modes[md in Modes] optional size md.pt;
dvar sequence mchs [m in Mchs] in
  all(md in Modes: md.mch == m) modes[md];

minimize max(o in Ops) endOf(ops[o]);
subject to {
  forall (j in Jobs, o1,o2 in JobOps[j]: o2.pos==1+o1.pos)
    endBeforeStart(ops[o1],ops[o2]);
  forall (o in Ops)
    alternative(ops[o], all(md in Modes: md.opId==o.id) modes[md]);
  forall (m in Mchs)
    noOverlap(mchs[m]);
}
```

CP Optimizer modeling concepts

- CP Optimizer has mathematical concepts that naturally map to features invariably found in industrial scheduling problems




Automatic Search

- Search algorithm is **Complete**
- Core CP techniques used as a building block:
 - Tree search (Depth First)
 - Constraint propagation
- But also:
 - Deterministic multicore parallelism
 - Model presolve
 - Algorithms portfolios
 - Machine learning
 - Restarting techniques
 - Large Neighborhood Search
 - No-good learning
 - Impact-based branching
 - Opportunistic probing
 - Dominance rules
 - LP-assisted heuristics
 - Randomization
 - Evolutionary algorithms



Automatic Search – Constraint propagation

- Dedicated Propagation Algorithms for Scheduling
 - Edge-Finding
 - Not-First/Not-Last
 - Detectable Precedences
 - Timetable
 - Timetable Edge-Finding
 - Max Energy Filtering
 - etc.





Automatic Search – Large Neighborhood Search (LNS)

- LNS is able to converge very quickly to quality solutions
 - 1) Start with an existing solution





Automatic Search – Large Neighborhood Search (LNS)

- LNS is able to converge very quickly to quality solutions
 - 1) Start with an existing solution
 - 2) Take part the solution and relax it, fix structure of the rest (but not start/end times)





Automatic Search – Large Neighborhood Search (LNS)

- LNS is able to converge very quickly to quality solutions
 - 1) Start with an existing solution
 - 2) Take part the solution and relax it, fix structure of the rest (but not start/end times)




Automatic Search – Large Neighborhood Search (LNS)

- LNS is able to converge very quickly to quality solutions
 - 1) Start with an existing solution
 - 2) Take part the solution and relax it, fix structure of the rest (but not start/end times)
 - 3) Find (improved) solution



Automatic Search – LP-assisted heuristics

- Traditionally, early/tardy problems are challenging for CP-based scheduling tools as they miss a good global view of the cost



- Approach:
 - Automatically use CPLEX's LP solver to provide a solution to a relaxed version of the problem
 - Use the LP solution to guide heuristics. Start an operation as close as possible to the time proposed by CPLEX
- What is linearized?
 - Precedences
 - Execution / non-execution costs, logical constraints on optional intervals
 - "Alternative" and "Span" constraints
 - Cost function terms which are functions of start/end times

Automatic Search: some results

MISTA 2007

Problem type	Benchmark	Problem size	Reference UB	MRD	# Imp. UBS / # Instances
Trolley	[41]	230-460	[19]	-11.8%	15/15
Hybrid flow-shop	[35]	200-1000	[35]	-8.8%	19/20
Job-shop w/ E/T	[3]	30-200	[3]	-5.6%	32/48
Air traffic management	[19]	2000	[19]	-4.0%	1/1
Flow-shop w/ E/T	[27]	30-400	[14]	-3.0%	4/12
Max. quality RCPSP	[33]	30	[33]	-2.3%	NA/3600
Cumulative job-shop	[28]	150-675	[17]	-0.3%	27/86
Single proc. tardiness	[20]	200-500	[20]	0.2%	0/20
Semiconductor testing	[30]	400	[30]	0.2%	7/18
RCPSP w/ E/T	[42]	30-50	[42]	0.4%	15/60
Open-shop	[9, 40, 18]	64-400	[15, 7, 25]	0.9%	0/28
RCPSP	[23]	120	Best PSPLIB	1.6%	0/600 ³
Shop w/ setup times	[10]	50-200	[2]	2.3%	0/15
Parallel machine w/ E/T	[29]	8-200	[4]	2.6%	2/52
Job-shop	[1, 39, 43, 40]	100-500	Best OR-Lib	2.8%	0/33
Air land	[5]	10-50	[5]	3.4%	0/8
Flow-shop w/ buffers	[40]	100-500	[8]	3.6%	12/30
Flow-shop	[40]	100-500	Best OR-Lib	5.9%	0/22
Aircraft assembly	[16]	575	[13]	8.7%	0/1
Single machine w/ E/T	[11, 37, 29]	8-500	[38]	9.8%	1/100
Common due-date	[6]	100-200	[36]	14.7%	0/20

Table 1: Results of SA-LNS on 21 scheduling benchmarks

Automatic Search: some results

CP-AI-OR 2009

- Problem #1: Flow-shop with earliness/tardiness cost
- Problem #2: Oversubscribed Satellite Scheduling problem
- Problem #3: Personal tasks scheduling

	OPL Model size	CPO Automatic search (no parameter tuning) vs. state-of-the-art
#1	20 lines	Competitive with state of the art (GA, LNS)
#2	15 lines	Number of unscheduled tasks decreased by 5%
#3	42 lines	Finds solution to more instances Solution quality increased by 12.5%

Automatic Search: some results

CP-AI-OR 2015

Benchmark set	Number of instances	Lower bound improvements	Upper bound improvements	Closed instances
JobShop	48	40	3	15
JobShopOperators	222	107	215	208
FlexibleJobShop	107	67	39	74
RCPS	472	52	1	0
RCPSMax	58	51	23	1
MultiModeRCPS (j30)	552	No reference	3	535
MultiModeRCPSMax	85	84	77	85

Table 1. Results summary

Automatic Search: some results


CP 2015

- Industrial Modelling Competition at CP 2015
- <http://booleconferences.ucc.ie/indmodellingcomp>
- CP Optimizer outperformed all the other competitors on all the instances of the challenge

Automatic Search

- Some CP Optimizer industrial applications:
 - Port Management: Yantian International Container Terminals, Navis
 - Manufacturing: Ajover S.A. (plastics), TAL Group (textiles)
 - Aviation: Dassault Aviation, another large jet manufacturer
 - Workforce scheduling: A world leading IFS (Integrated Facility Services) company

Development tools



CPLEX and CP Optimizer are in the same ecosystem

		CPLEX	CP Optimizer
Interfaces		OPL, C++, Java, .NET, C, Python	OPL, C++, Java, .NET, Python
Model	Decision variables	int, float	int, interval
	Expressions	linear, quadratic	arithmetic, log, pow, ... relational, a[x], count,...
	Constraints	range	relational, logical, specialized, scheduling
Search	Search parameters	✓	✓
	Warm start	✓	✓
	Multi-core //	✓	✓
Tools	Search log	✓	✓
	I/O format	.lp, .mps,cpo
	Conflict refiner	✓	✓