

Constraints
Manuscript Draft

Manuscript Number: CONS93

Title: Evaluating the Impact of AND/OR Search on 0-1 Integer Linear Programming

Article Type: Original Research

Section/Category:

Keywords: search; AND/OR search spaces; constraint optimization; integer programming

Manuscript Region of Origin:

Abstract: AND/OR search spaces accommodate advanced algorithmic schemes for graphical models which can exploit the structure of the model. The paper extends and evaluates the depth-first and best-first AND/OR search algorithms to solving 0-1 Integer Linear Programs (0-1 ILP) within this framework. We also include a class of dynamic variable ordering heuristics while exploring an AND/OR search tree for 0-1 ILPs. We demonstrate the effectiveness of these search algorithms on a variety of benchmarks, including real-world combinatorial auctions, random uncapacitated warehouse location problems and MAX-SAT instances.

Evaluating the Impact of AND/OR Search on 0-1 Integer Linear Programming

Radu Marinescu, Rina Dechter

*Donald Bren School of Information and Computer Science
University of California, Irvine, CA 92697, USA*

Abstract

AND/OR search spaces accommodate advanced algorithmic schemes for graphical models which can exploit the structure of the model. The paper extends and evaluates the *depth-first* and *best-first* AND/OR search algorithms to solving 0-1 Integer Linear Programs (0-1 ILP) within this framework. We also include a class of dynamic variable ordering heuristics while exploring an AND/OR search tree for 0-1 ILPs. We demonstrate the effectiveness of these search algorithms on a variety of benchmarks, including real-world combinatorial auctions, random uncapacitated warehouse location problems and MAX-SAT instances.

1 Introduction

A *constraint optimization problem* is the minimization (or maximization) of an objective function subject to a set of constraints on the possible values of a set of independent decision variables. An important class of optimization problems in operations research and computer science are the 0-1 Integer Linear Programming problems (0-1 ILP) [1] where the objective is to optimize a linear function of bi-valued integer decision variables, subject to a set of linear equality or inequality constraints defined on subsets of variables. The classical approach to solving 0-1 ILPs is the *Branch-and-Bound* method [2] which maintains the best solution found so far, while discarding partial solutions which cannot improve on the best.

The AND/OR search space for graphical models [3] is a relatively new framework for search that is sensitive to the independencies in the model, often resulting in

Email addresses: radum@ics.uci.edu (Radu Marinescu),
dechter@ics.uci.edu (Rina Dechter).

1
2
3
4 substantially reduced complexities. It is based on a pseudo tree that captures conditional
5 independencies in the graphical model, resulting in a search tree exponential
6 in the depth of the pseudo tree, rather than in the number of variables.
7

8
9 The AND/OR Branch-and-Bound search (AOBB) introduced in [4,5] is a Branch-
10 and-Bound algorithm that explores a search tree in a depth-first manner for solving
11 optimization tasks in graphical models. The AND/OR Branch-and-Bound search
12 with caching algorithm (AOBB-C) [6,7] improves AOBB by allowing the algorithm
13 to save previously computed results and retrieve them when the same subproblems
14 are encountered again. The algorithm explores the context minimal search graph
15 [6,7]. A *best-first* AND/OR search algorithm (AOBF-C) that traverses the search
16 graph was introduced subsequently [8,9,7]. The algorithms were initially restricted
17 to a static variable ordering determined by the underlying pseudo tree, but sub-
18 sequent extensions to dynamic variable ordering heuristics were also introduced
19 [10,8,5]. Two such extensions, AND/OR Branch-and-Bound with Partial Variable
20 Ordering (AOBB+PVO) and best-first AND/OR search with Partial Variable Order-
21 ing (AOBF+PVO) were shown to have significant impact on several domains. The
22 above results were summarized in 2 papers which are currently and concurrently
23 under review with this paper.
24

25
26 In this paper we extend the general principles of solving combinatorial optimization
27 problems using AND/OR search with context-based caching to the class of 0-1
28 ILPs. We explore both depth-first and best-first control strategies. Under conditions
29 of admissibility and monotonicity of the guiding heuristic function, best-first search
30 is known to expand the minimal number of nodes, at the expense of using additional
31 memory [11]. We also extend dynamic variable ordering heuristics for AND/OR
32 search and explore their impact on 0-1 ILPs.
33

34
35 We demonstrate empirically the benefit of the AND/OR search approach on several
36 benchmarks for 0-1 ILP problems, including combinatorial auctions, random un-
37 capacitated warehouse location problems and MAX-SAT problem instances. Our
38 results show conclusively that these new algorithms improve dramatically over the
39 traditional OR search on this domain, in some cases with several orders of magni-
40 tude of improved performance. We illustrate the tremendous gain obtained by ex-
41 ploring problem's decomposition (using AND nodes), equivalence (by caching),
42 branching strategy (via dynamic variable ordering heuristics) and control strategy.
43 We also show that the AND/OR algorithms are competitive and in some cases out-
44 perform significantly commercial ILP solvers such as CPLEX.
45

46
47 The paper is organized as follows. Sections 2 and 3 provide background on 0-1
48 ILP and AND/OR search spaces, respectively. In Sections 4 and 5 we present the
49 depth-first AND/OR Branch-and-Bound and the best-first AND/OR search algo-
50 rithms for 0-1 ILP. Section 6 describes the AND/OR search approach that incorpo-
51 rates dynamic variable ordering heuristics. Section 7 is dedicated to our empirical
52 evaluation, Section 8 overviews related work, while Section 9 provides a summary,
53

1
2
3
4 concluding remarks and directions of future research.
5
6
7 The work we present here is based in part on 2 conference submissions [12,13].
8
9
10
11
12

13 2 Background

14
15
16

17 **Notations.** The following notations will be used throughout the paper. We denote
18 variables by uppercase letters (*e.g.*, X, Y, Z, \dots), subsets of variables by bold faced
19 uppercase letters (*e.g.*, $\mathbf{X}, \mathbf{Y}, \mathbf{Z}, \dots$) and values of variables by lower case letters
20 (*e.g.*, x, y, z, \dots). An assignment $(X_1 = x_1, \dots, X_n = x_n)$ can be abbreviated as
21 $x = (\langle X_1, x_1 \rangle, \dots, \langle X_n, x_n \rangle)$ or $x = (x_1, \dots, x_n)$. For a subset of variables \mathbf{Y} , $D_{\mathbf{Y}}$
22 denotes the Cartesian product of the domains of variables in \mathbf{Y} . $x_{\mathbf{Y}}$ and $x[\mathbf{Y}]$ are
23 both used as the projection of $x = (x_1, \dots, x_n)$ over a subset \mathbf{Y} . We denote functions
24 by letters f, h, g etc., and the scope (set of arguments) of a function f by $\text{scope}(f)$.
25
26

27
28 **DEFINITION 1 (constraint optimization problem)** *A finite constraint optimization*
29 *problem (COP) is a four-tuple $\langle \mathbf{X}, \mathbf{D}, \mathbf{F}, z \rangle$, where $\mathbf{X} = \{X_1, \dots, X_n\}$ is a set of*
30 *variables, $\mathbf{D} = \{D_1, \dots, D_n\}$ is a set of finite domains, $\mathbf{F} = \{F_1, \dots, F_r\}$ is a set*
31 *of constraints on the variables and $z(\mathbf{X})$ is a global cost function defined over*
32 *\mathbf{X} (also called objective function) to be optimized (i.e., minimized or maximized).*
33 *The scope of a constraint F_i , denoted $\text{scope}(F_i) \subseteq \mathbf{X}$, is the set of arguments of*
34 *F_i . Constraints can be expressed extensionally, through relations, or intentionally,*
35 *by a mathematical formula (e.g., equality or inequality). An optimal solution to a*
36 *COP is a complete value assignment to all the variables such that every constraint*
37 *is satisfied and the objective function is minimized or maximized.*

38
39 With every COP instance we can associate a *constraint graph* G which has a node
40 for each variable and connects any two nodes whose variables appear in the scope
41 of the same constraint. The *induced graph* [14] of G relative to an ordering d of its
42 variables, denoted $G^*(d)$, is obtained by processing the nodes in reverse order of d .
43 For each node all its earlier neighbors are connected, including neighbors connected
44 by previously added edges. Given a graph and an ordering of its nodes, the *width*
45 of a node is the number of edges connecting it to nodes lower in the ordering. The
46 *induced width* (or *treewidth*) of a graph, denoted $w^*(d)$, is the maximum width of
47 nodes in the induced graph.
48
49

50 **DEFINITION 2 (linear program)** *A linear program (LP) consists of a set of n con-*
51 *tinuous variables $\mathbf{X} = \{X_1, \dots, X_n\}$ and a set of m linear constraints (equalities*
52 *or inequalities) $\mathbf{F} = \{F_1, \dots, F_m\}$ defined on subsets of variables. The goal is to*
53 *minimize a global linear cost function, denoted $z(\mathbf{X})$, subject to the constraints.*
54 *One of the standard forms of a linear program is:*

$$1 \quad \min z(\mathbf{X}) = \sum_{i=1}^n c_i \cdot X_i \quad (1)$$

$$2 \quad \text{s.t. } \sum_{i=1}^n a_{ij} \cdot X_i \leq b_j, \quad \forall 1 \leq j \leq m \quad (2)$$

$$3 \quad X_i \geq 0, \quad \forall 0 \leq i \leq n \quad (3)$$

4 where (1) represents the linear objective function, and (2) defines the set of linear
 5 constraints. In addition, (3) ensures that all variables are positive.

6 A linear program can also be expressed in a matrix notation, as follows:

$$7 \quad \min\{\mathbf{c}^\top \mathbf{X} \mid \mathbf{A} \cdot \mathbf{X} \leq \mathbf{b}, \mathbf{X} \geq 0\} \quad (4)$$

8 where $\mathbf{c} \in \mathbb{R}^n$, $\mathbf{b} \in \mathbb{R}^m$, $\mathbf{A} \in \mathbb{R}^{m \times n}$ and $\mathbf{X} \in \mathbb{R}_+^n$. Namely, \mathbf{c} represents the cost
 9 vector and \mathbf{X} is the vector of decision variables. The vector \mathbf{b} and the matrix \mathbf{A}
 10 define the m linear constraints.

11 One of the most important constraint optimization problems in operations research
 12 and computer science is *integer programming*. Applications of integer program-
 13 ming include scheduling, routing, VLSI circuit design, combinatorial auctions, and
 14 facility location [1]. Formally:

15 **DEFINITION 3 (0-1 integer linear programming)** A 0-1 Integer Linear Program-
 16 ming (0-1 ILP) problem is a linear program where all the decision variables are
 17 constrained to have integer values 0 or 1 at the optimal solution. Formally,

$$18 \quad \min z(\mathbf{X}) = \sum_{i=1}^n c_i \cdot X_i \quad (5)$$

$$19 \quad \text{s.t. } \sum_{i=1}^n a_{ij} \cdot X_i \leq b_j, \quad \forall 1 \leq j \leq m \quad (6)$$

$$20 \quad X_i \in \{0, 1\} \quad \forall 0 \leq i \leq n \quad (7)$$

21 **Example 1** Figure 1(a) shows a 0-1 ILP instance with 6 binary decision variables
 22 (A, B, C, D, E, F) and 4 linear constraints $F_1(A, B, C), F_2(B, C, D), F_3(A, B, E),$
 23 $F_4(A, E, F)$. The objective function to be minimized is defined by $z = 7A + B -$
 24 $2C + 5D - 6E + 8F$. Figure 1(b) displays the constraint graph associated with
 25 this 0-1 ILP, where nodes correspond to the decision variables and there is an edge
 26 between any two nodes whose variables appear in the scope of the same constraint.


```

1
2
3
4
5   minimize :  $z = 7A + 3B - 2C + 5D - 6E + 8F$ 
6   subject to :
7      $3A - 12B + C \leq 3$ 
8      $-2B + 5C - 3D \leq -2$ 
9      $2A + B - 4E \leq 2$ 
10     $A - 3E + F \leq 1$ 
11     $A, B, C, D, E, F \in \{0,1\}$ 
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65

```

minimize : $z = 7A + 3B - 2C + 5D - 6E + 8F$
 subject to :
 $3A - 12B + C \leq 3$
 $-2B + 5C - 3D \leq -2$
 $2A + B - 4E \leq 2$
 $A - 3E + F \leq 1$
 $A, B, C, D, E, F \in \{0,1\}$

(a) 0-1 Integer Linear Program



(b) Constraint graph

Fig. 1. Example of a 0-1 Integer Linear Program.

If some variables are constrained to be integers (not necessarily binary), then the problem is simply called *Integer Programming*. If not all variables are constrained to be integral (they can be real), then the problem is called *Mixed Integer Programming* (MIP). Otherwise, the problem is called *0-1 Integer Programming*.

While 0-1 integer programming, and thus integer programming and MIP are all NP-hard [15], there are many sophisticated techniques that can solve very large instances in practice. We next briefly review the existing search techniques upon which we build our methods.

In **Branch-and-Bound** search, the best solution found so far (the *incumbent*) is kept in memory. Once a node in the search tree is generated, a lower bound (also known as a heuristic evaluation function) on the solution value is computed by solving a relaxed version of the problem, while honoring the commitments made on the search path so far. The most common method for doing this is to solve the problem while relaxing only the integrality constraints of all undecided variables. The resulting *linear program* (LP) can be solved fast in practice, for example using the *simplex* algorithm [16] (and in polynomial worst-case time using integer-point methods [1]). A path terminates when the lower bound is at least the value of the incumbent, or when the subproblem is infeasible or yields an integer solution. Once all paths have terminated, the incumbent is a provably optimal solution.

There are several ways to decide which leaf node of the search tree to expand next. For example, in *depth-first* Branch-and-Bound, the most recent node is expanded next. In *best-first search* (*i.e.*, A^* search [17]), the leaf with the lowest lower bound is expanded next. A^* search is desirable because for any fixed branching variable ordering, no tree search algorithm that finds a provably optimal solution can guarantee expanding fewer nodes [11]. Therefore, of the known node-selection strategies, A^* seems to be best suited when the goal is to find a provably optimal solution. A variant of a best-first node-selection strategy, called *best-bound search*, is often used in MIP [18]. While in general A^* the children are evaluated when they are generated, in best-bound search the children are queued for expansion based on their parents' values and the LP of each child is only solved if the child comes up for expansion from the queue. Thus best-bound search needs to continue until

each node on the queue has value no better than the incumbent. Best-bound search generates more nodes, but may require fewer (or more) LPs to be solved.

Branch-and-Cut Search for Integer Programming. A modern algorithm for solving MIPs is *Branch-and-Cut*, which first achieved success in solving large instances of the traveling salesman problem [19,20], and is now the core of the fastest commercial general-purpose integer programming packages. It is a *Branch-and-Bound* that uses the idea of *cutting planes* [1]. These are linear constraints that are deduced during search and, when added to the subproblem at a search node, may result in a smaller feasible space for the LP and thus a higher lower bound. Higher lower bounds can cause earlier termination of the search path, and thus yielding smaller search trees.

Software Packages. CPLEX¹ is a leading commercial software product for solving MIPs. It uses Branch-and-Cut, and it can be configured to support many different branching algorithms (*i.e.*, variable ordering heuristics). It also makes available low-level interfaces (*i.e.*, APIs) for controlling the search, as well as other components such as the pre-solver, the cutting plane engine and the LP solver.

`lp_solve`² is an open source linear (integer) programming solver based on the *simplex* and the Branch-and-Bound methods. We chose to develop our AND/OR search algorithms in the framework of `lp_solve`, because we could have access to the source code. Unlike CPLEX, `lp_solve` does not provide a cutting plane engine nor a best-bound control strategy.

3 Extending AND/OR Search Spaces to 0-1 Integer Linear Programs

As mentioned earlier, the common way of solving 0-1 ILPs is by search, namely to instantiate variables one at a time following a static/dynamic variable ordering. In the simplest case, this process defines an OR search tree, whose nodes represent states in the space of partial assignments. However, this search space does not capture independencies that appear in the structure of the problem. To remedy this problem the idea of AND/OR search spaces [21] was recently introduced to general graphical models [3]. The AND/OR search space for a graphical model is defined using a backbone *pseudo tree* [22,23].


DEFINITION 4 (pseudo tree, extended graph) Given an undirected graph $G = (V, E)$, a directed rooted tree $T = (V, E')$ defined on all its nodes is called pseudo tree if any arc of G which is not included in E' is a back-arc, namely it connects a node to an ancestor in T . Given a pseudo tree T of G , the extended graph of G

¹ <http://www.ilog.com/cplex/>

² <http://lpsolve.sourceforge.net/5.5/>

1
2
3
4
5 minimize : $z = 7A + 3B - 2C + 5D - 6E + 8F$
6 subject to :

$$\begin{aligned}3A - 12B + C &\leq 3 \\-2B + 5C - 3D &\leq -2 \\2A + B - 4E &\leq 2 \\A - 3E + F &\leq 1 \\A, B, C, D, E, F &\in \{0,1\}\end{aligned}$$



(a) 0-1 Integer Linear Program

(b) Pseudo tree




Fig. 2. AND/OR search tree for a 0-1 Integer Linear Program instance.

relative to \mathcal{T} is defined as $G^{\mathcal{T}} = (\mathbf{V}, \mathbf{E} \cup \mathbf{E}')$.

We will next specialize the AND/OR search space for a 0-1 ILPs.

3.1 AND/OR Search Trees for 0-1 Integer Linear Programs

Given a 0-1 ILP instance, its constraint graph G and a pseudo tree \mathcal{T} of G , the associated AND/OR search tree $S_{\mathcal{T}}$ has alternating levels of OR nodes and AND nodes. The OR nodes are labeled by X_i and correspond to the variables. The AND nodes are labeled by $\langle X_i, x_i \rangle$ (or simply x_i) and correspond to value assignments in the domains of the variables that are consistent relative to the constraints. The structure

of the AND/OR tree is based on the underlying pseudo tree \mathcal{T} of G . The root of the AND/OR search tree is an OR node, labeled with the root of \mathcal{T} . The children of an OR node X_i are AND nodes labeled with assignments $\langle X_i, x_i \rangle$, consistent along the path from the root. The children of an AND node $\langle X_i, x_i \rangle$ are OR nodes labeled with the children of variable X_i in \mathcal{T} . A path from the root of the search tree to a node n is denoted by π_n . If n is labeled X_i or x_i the path will be denoted by $\pi_n(X_i)$ or $\pi_n(x_i)$. The assignment sequence along π_n , denoted $asgn(\pi_n)$, is the set of assignments associated with the AND nodes along π_n .

Semantically, the OR states represent alternative ways of solving the problem, whereas the AND states represent problem decomposition into independent subproblems, all of which need be solved. When the pseudo tree is a chain, the AND/OR search tree coincides with the regular OR search tree.

As usual [21,3], a *solution tree* T of an AND/OR search tree $S_{\mathcal{T}}$ is an AND/OR subtree such that: (i) it contains the root of $S_{\mathcal{T}}$, s ; (ii) if a non-terminal AND node $n \in S_{\mathcal{T}}$ is in T then all of its children are in T ; (iii) if a non-terminal OR node $n \in S_{\mathcal{T}}$ is in T then exactly one of its children is in T ; (iv) all its terminal leaf nodes (full assignments) are consistent relative to the constraints of the 0-1 ILP.

Example 2 Consider the 0-1 ILP instance from Figure 2(a). A pseudo tree of the constraint graph, together with the back-arcs (dotted lines) are given in Figure 2(b). Figure 2(c) shows the corresponding AND/OR search tree. Notice that the partial assignment $(A = 0, B = 0, C = 0, D = 0)$ which is represented by the path $\{A, \langle A, 0 \rangle, B, \langle B, 0 \rangle, C, \langle C, 0 \rangle, D, \langle D, 0 \rangle\}$ in the AND/OR search tree, is inconsistent because the constraint $-2B + 5C - 3D \leq -2$ is violated. Similarly, the partial assignment $(A = 0, B = 0, C = 1)$ is also inconsistent due to the violation of the same constraint for any value assignment of variable D .

It was shown that:

THEOREM 1 (size of AND/OR search trees [3]) Given a 0-1 ILP instance and a backbone pseudo tree \mathcal{T} , its AND/OR search tree $S_{\mathcal{T}}$ contains all consistent solutions, and its size is $O(l \cdot 2^m)$ where m is the depth of the pseudo tree and l bounds its number of leaves. If the 0-1 ILP instance has induced width w^* , then there is a pseudo tree whose associated AND/OR search tree is $O(n \cdot 2^{w^* \cdot \log n})$.

The arcs in the AND/OR search tree are associated with *weights* that are obtained from the objective function of the given 0-1 ILP instance.

DEFINITION 5 (weights) Given a 0-1 ILP with objective function $\sum_{i=1}^n c_i \cdot X_i$ and an AND/OR search tree $S_{\mathcal{T}}$ relative to a pseudo tree \mathcal{T} , the weight $w(n, m)$ of the arc from the OR node n , labeled X_i to the AND node m , labeled $\langle X_i, x_i \rangle$, is defined as $w(n, m) = c_i \cdot x_i$.

Note that the arc-weights in general COPs are a more involved function of the input

$$\text{minimize : } z = 7A + 3B - 2C + 5D - 6E + 8F$$




Fig. 3. Weighted AND/OR search tree for the 0-1 ILP instance from Figure 2.
specification (see also [4,5] for additional details).

DEFINITION 6 (cost of a solution tree) Given a weighted AND/OR search tree S_T of a 0-1 ILP, and given a solution tree T having OR-to-AND set of arcs $\text{arcs}(T)$, the cost of T , $f(T)$, is defined by $f(T) = \sum_{e \in \text{arcs}(T)} w(e)$.

With each node n of the search tree we can associate a value $v(n)$ which stands for the optimal solution cost of the subproblem below n , conditioned on the assignment on the path leading to it [3]. $v(n)$ was shown to obey the following recursive definition:

DEFINITION 7 (node value) The value of a node n in the AND/OR search tree of a 0-1 ILP instance is defined recursively by :

$$v(n) = \begin{cases} 0 & , \text{ if } n = \langle X, x \rangle \text{ is a terminal AND node} \\ \infty & , \text{ if } n = X \text{ is a terminal OR node} \\ \sum_{m \in \text{succ}(n)} v(m) & , \text{ if } n = \langle X, x \rangle \text{ is an AND node} \\ \min_{m \in \text{succ}(n)} (w(n, m) + v(m)) & , \text{ if } n = X \text{ is an OR node} \end{cases}$$

where $\text{succ}(n)$ denotes the children of n in the AND/OR tree.

Example 3 Figure 3 shows the weighted AND/OR search tree associated with the 0-1 ILP instance from Figure 2. The numbers on the OR-to-AND arcs are the weights corresponding to the objective function. For example, the weights associated with the OR-to-AND arcs $(A, \langle A, 0 \rangle)$ and $(A, \langle A, 1 \rangle)$ are 0 and 7, respectively.




Fig. 4. Context minimal AND/OR search graph for the 0-1 ILP instance from Figure 2.

An optimal solution tree that corresponds to the assignment ($A = 0, B = 1, C = 0, D = 0, E = 1, F = 0$) with cost -3 is highlighted. Note that inconsistent portions of the tree are pruned.

Clearly, the value of the root node s is the minimal cost solution to the initial problem, namely $v(s) = \min_{\mathbf{X}} \sum_{i=1}^n c_i \cdot X_i$.

Therefore, search algorithms that traverse the AND/OR search tree and compute the value of the root node yield the answer to the problem. Consequently, depth-first search algorithms traversing the weighted AND/OR search tree are guaranteed to have time complexity bounded exponentially in the depth of the pseudo tree and can operate in linear space only.

3.2 AND/OR Search Graphs for 0-1 Integer Linear Programs

Often different nodes in the search tree root identical subtrees, and correspond to identical subproblems. Any two such nodes can be *merged*, reducing the size of the search space and converting it into a graph. Some of these mergeable nodes can be identified based on *contexts*, as described in [3] and as we briefly outline below.

Given a pseudo tree \mathcal{T} of an AND/OR search space, the *context* of an AND node

$\langle X_k, x_k \rangle$, denoted by $\text{context}(X_k) = [X_1, \dots, X_k]$, is the set of ancestors of X_k in \mathcal{T} , including X_k , ordered descendingly, that are connected in the primal graph with descendants of X_k . The context of $\langle X_k, x_k \rangle$ separates the subproblem below $\langle X_k, x_k \rangle$ from the rest of the network. The *context minimal AND/OR graph* [3] is obtained from the AND/OR search tree by merging all context mergeable nodes. It can be shown that:

THEOREM 2 (size of AND/OR graphs [3]) *Given a 0-1 ILP instance, its constraint graph G , and a pseudo tree \mathcal{T} having induced width $w^* = w_{\mathcal{T}}(G)$, the size of the context minimal AND/OR search graph based on \mathcal{T} , $\mathcal{G}_{\mathcal{T}}$, is $O(n \cdot 2^{w^*})$.*

Example 4 Figure 4(b) shows the context minimal AND/OR search graph corresponding to the 0-1 ILP from Figure 2, relative to the pseudo tree given in Figure 4(a). The square brackets indicate the AND contexts of the variables, as follows: $\text{context}(A) = \{A\}$, $\text{context}(B) = \{A, B\}$, $\text{context}(C) = \{B, C\}$, $\text{context}(D) = \{D\}$, $\text{context}(E) = \{A, E\}$ and $\text{context}(F) = \{F\}$. Note that the cache tables of A and B will not get cache hits during a depth-first traversal, because the corresponding AND levels have only one incoming arc. It can be determined from the pseudo tree inspection that variables A and B generate dead-caches [3] and their cache tables need not be stored.

4 Depth-First AND/OR Branch-and-Bound Search

Traversing AND/OR search spaces by best-first algorithms or depth-first Branch-and-Bound was described as early as [21,24,25]. In a series of papers [4,6,10,8,9] we introduced a new generation of depth-first Branch-and-Bound and best-first AND/OR search algorithms for solving constraint optimization tasks in graphical models. Our extensive empirical evaluations on a variety of probabilistic and deterministic graphical models demonstrated the power of these new algorithms over competitive approaches exploring traditional OR search spaces. In this section we revisit the notions of partial solution trees [21] to represent sets of solution trees, and heuristic evaluation function of a partial solution tree [4]. We will then recap the depth-first Branch-and-Bound algorithm for searching the AND/OR graphs, focusing on the specific properties for 0-1 ILPs.

DEFINITION 8 (partial solution tree) A partial solution tree T' of a context minimal AND/OR search graph $\mathcal{G}_{\mathcal{T}}$ is a subtree which: (1) contains the root node s of $\mathcal{G}_{\mathcal{T}}$; (2) if n is an OR node in T' then it contains one of its AND child nodes in $\mathcal{G}_{\mathcal{T}}$, and if n is an AND node it contains all its OR children in $\mathcal{G}_{\mathcal{T}}$. A node of T' is a tip node if it has no children in T' . A tip node of T' is either a terminal node (if it has no children in $\mathcal{G}_{\mathcal{T}}$), or a non-terminal node (if it has children in $\mathcal{G}_{\mathcal{T}}$).

A partial solution tree represents $\text{extension}(T')$, the set of all full solution trees which can extend it. Clearly, a partial solution tree whose all tip nodes are terminal in \mathcal{G}_T is a solution tree.

Branch-and-Bound algorithms for 0-1 ILP are guided by the LP relaxation based lower bound heuristic function. The extension of heuristic evaluation functions to subtrees in an AND/OR search space was elaborated in [4,5]. We briefly introduce here the main elements and refer the reader for further details to [4,5].

Heuristic Lower Bounds on Partial Solution Trees. We present next the notions of exact evaluation as well as heuristic evaluation functions of a partial solution tree [4,5], that are used to guide the AND/OR Branch-and-Bound.

The *exact evaluation function* $f^*(T')$ of a partial solution tree T' is the minimum of the costs of all solution trees extending T' , namely: $f^*(T') = \min\{f(T) \mid T \in \text{extension}(T')\}$. If $f^*(T'_n)$ is the exact evaluation function of a partial solution tree rooted at node n , then $f^*(T'_n)$ can be computed recursively, as follows:

1. If T'_n consists of a single node n then $f^*(T'_n) = v(n)$.
2. If n is an OR node having the AND child m in T'_n , then $f^*(T'_n) = w(n, m) + f^*(T'_m)$.
3. If n is an AND node having OR children m_1, \dots, m_k in T'_n , then $f^*(T'_n) = \sum_{i=1}^k f^*(T'_{m_i})$.

If each non-terminal tip node m of T' is assigned a heuristic lower bound estimate $h(m)$ of $v(m)$, then it induces a heuristic evaluation function on the minimal cost extension of T' . Given a partial solution tree T'_n rooted at n in the AND/OR graph \mathcal{G}_T , the *tree-based heuristic evaluation function* $f(T'_n)$, is defined recursively by:

1. If T'_n consists of a single node n , then $f(T'_n) = h(n)$.
2. If n is an OR node having the AND child m in T'_n , then $f(T'_n) = w(n, m) + f(T'_m)$.
3. If n is an AND node having OR children m_1, \dots, m_k in T'_n , then $f(T'_n) = \sum_{i=1}^k f(T'_{m_i})$.

Clearly, by definition, $f(T'_n) \leq f^*(T'_n)$. If n is the root of the context minimal AND/OR search graph, then $f(T') \leq f^*(T')$ [4,5].

During search we maintain both an upper bound $ub(s)$ on the optimal solution $v(s)$ as well as the heuristic evaluation function $f(T')$ of the current partial solution tree T' being explored, and whenever $f(T') \geq ub(s)$, then searching below the current tip node t of T' is guaranteed not to yield a better solution cost than $ub(s)$ and therefore, search below t can be halted.

In [4,5] we also showed that the pruning test can be sped up if we associate upper bounds with internal nodes as well. Specifically, if m is an OR ancestor of t in T'




Fig. 5. Illustration of the pruning mechanism.

and T'_m is the subtree of T' rooted at m , then it is also safe to prune the search tree below t , if $f(T'_m) \geq ub(m)$.

Example 5 Consider the partially explored weighted AND/OR search tree in Figure 5 (the weights and node values are given for illustration only). The current partial solution tree T' is highlighted. It contains the following nodes: A , $\langle A, 1 \rangle$, B , $\langle B, 1 \rangle$, C , $\langle C, 0 \rangle$, D , $\langle D, 1 \rangle$ and F . The nodes labeled by $\langle D, 1 \rangle$ and by F are non-terminal tip nodes and their corresponding heuristic estimates are $h(\langle D, 1 \rangle) = 2$ and $h(F) = 9$, respectively. The subtrees rooted at the AND nodes labeled $\langle A, 0 \rangle$, $\langle B, 0 \rangle$ and $\langle D, 0 \rangle$ are fully evaluated, and therefore the current upper bounds of the OR nodes labeled A , B and D , along the active path, are $ub(A) = 12$, $ub(B) = 10$ and $ub(D) = 0$, respectively. Moreover, the heuristic evaluation functions of the partial solution subtrees rooted at the OR nodes along the current path can be computed recursively, namely $f(T'_A) = 13$, $f(T'_B) = 12$ and $f(T'_D) = -1$, respectively. Notice that while we could prune below $\langle D, 1 \rangle$ because $f(T'_A) > ub(A)$, we could discover this pruning earlier by looking at node B only, because $f(T'_B) > ub(B)$. Therefore, the partial solution tree T'_A need not be consulted in this case.

The **Depth-First AND/OR Branch-and-Bound** search algorithm, AOBB-C-ILP, that traverses the context minimal AND/OR graph via full caching is described by Algorithm 1 and shown here for completeness. It specializes the Branch-and-Bound algorithm introduced in [6,7] to 0-1 ILPs. If the caching mechanism is disabled then the algorithm uses linear space only and traverses an AND/OR search tree [4,5].

As we showed in [6,7], the context based caching is done using tables. For each variable X_i , a table is reserved in memory for each possible assignment to its context. Initially, each entry has a predefined value, in our case NULL. The fringe of the search is maintained by a stack called OPEN. The current node is denoted by n , its parent by p , and the current path by π_n . The children of the current node are denoted by $succ(n)$. The flag caching is used to enable the caching mechanism.

Algorithm 1: AOBB-C-ILP: AND/OR Branch-and-Bound Search for 0-1 ILP

Input: A 0-1 ILP instance with objective function $\sum_{i=1}^n c_i X_i$, pseudo tree T rooted at X_1 , AND contexts pas_i for every variable X_i , caching set to *true* or *false*.

Output: Minimal cost solution.

```

1   $v(s) \leftarrow \infty; OPEN \leftarrow \{s\}$                                 // Initialize search stack
2  if caching == true then
3      Initialize cache tables with entries "NULL"                         // Initialize cache tables
4  while OPEN ≠  $\emptyset$  do
5       $n \leftarrow top(OPEN)$ ; remove  $n$  from OPEN;  $succ(n) \leftarrow \emptyset$            // EXPAND
6      if  $n$  is marked INFEASIBLE or INTEGER then
7           $v(n) \leftarrow \infty$  (if INFEASIBLE) or  $v(n) \leftarrow h(n)$  (if INTEGER)
8      else if  $n$  is an OR node, labeled  $X_i$  then
9          foreach  $x_i \in D_i$  do
10         create an AND node  $n'$ , labeled  $\langle X_i, x_i \rangle$ 
11          $v(n') \leftarrow 0; h(n') \leftarrow LP(P_{n'})$                                 // Solve the LP relaxation
12          $w(n, n') \leftarrow c_i \cdot x_i$                                          // Compute the arc weight
13         mark  $n'$  as INFEASIBLE or INTEGER if the LP relaxation is infeasible or has an integer solution
14          $succ(n) \leftarrow succ(n) \cup \{n'\}$ 
15     else if  $n$  is an AND node, labeled  $\langle X_i, x_i \rangle$  then
16          $cached \leftarrow false; deadend \leftarrow false$ 
17         if caching == true and Cache(asgn( $\pi_n$ )[ $pas_i$ ]) ≠ NULL then
18              $v(n) \leftarrow Cache(asgn(\pi_n)[pas_i])$                                 // Retrieve value
19              $cached \leftarrow true$                                               // No need to expand below
20         foreach OR ancestor  $m$  of  $n$  do
21              $lb \leftarrow evalPartialSolutionTree(T'_m)$ 
22             if  $lb \geq v(m)$  then
23                  $deadend \leftarrow true$                                            // Pruning
24                 break
25         if  $deadend == false$  and cached == false then
26             foreach  $X_j \in children_T(X_i)$  do
27                 create an OR node  $n'$  labeled  $X_j$ 
28                  $v(n') \leftarrow \infty; h(n') \leftarrow LP(P_{n'})$                       // Solve the LP relaxation
29                 mark  $n'$  as INFEASIBLE or INTEGER if the LP relaxation is infeasible or has an integer solution
30                  $succ(n) \leftarrow succ(n) \cup \{n'\}$ 
31     else if  $deadend == true$  then
32          $succ(p) \leftarrow succ(p) - \{n\}$ 
33     Add  $succ(n)$  on top of OPEN                                         // PROPAGATE
34     while  $succ(n) == \emptyset$  do
35         if  $n$  is an OR node, labeled  $X_i$  then
36             if  $X_i == X_1$  then
37                 return  $v(n)$                                                  // Search is complete
38              $v(p) \leftarrow v(p) + v(n)$                                          // Update AND node value (summation)
39         else if  $n$  is an AND node, labeled  $\langle X_i, x_i \rangle$  then
40             if caching == true and  $v(n) \neq \infty$  then
41                 Cache(asgn( $\pi_n$ )[ $pas_i$ ])  $\leftarrow v(n)$                          // Save AND node value in cache
42             if  $v(p) > (w(p, n) + v(n))$  then
43                  $v(p) \leftarrow w(p, n) + v(n)$                                          // Update OR node value (minimization)
44         remove  $n$  from succ(p)
45          $n \leftarrow p$ 

```

Each node n in the search graph maintains its current value $v(n)$, which is updated based on the values of its children. For OR nodes, the current $v(n)$ is an upper bound on the optimal solution cost below n . Initially, $v(n)$ is set to ∞ if n is OR, and 0 if n is AND, respectively. The heuristic function $h(n)$ in the search graph is computed by solving the LP relaxation of the subproblem rooted at n , conditioned

Algorithm 2: Recursive computation of the heuristic evaluation function.

```

1   function evalPartialSolutionTree( $T'_n$ )
2     Input: Partial solution subtree  $T'_n$  rooted at node  $n$ .
3     Output: Heuristic evaluation function  $f(T'_n)$ .
4
5     1 if  $\text{succ}(n) == \emptyset$  then
6       2 return  $h(n)$ 
7     3 else
8       4 if  $n$  is an AND node then
9         5 let  $m_1, \dots, m_k$  be the OR children of  $n$  in  $T'_n$ 
10        6 return  $\sum_{i=1}^k \text{evalPartialSolutionTree}(T'_{m_i})$ 
11      7 else if  $n$  is an OR node then
12        8 let  $m$  be the AND child of  $n$  in  $T'_n$ 
13        9 return  $w(n, m) + \text{evalPartialSolutionTree}(T'_m)$ 

```

on the current partial assignment along π_n (i.e., $\text{asgn}(\pi_n)$) (lines 11 and 28, respectively). Notice that if the LP relaxation is infeasible, we assign $h(n) = \infty$ and in this case $v(n) = \infty$, denoting inconsistency. Similarly, if the LP has an integer solution, then $h(n)$ equals $v(n)$. In both cases, $\text{succ}(n)$ is set to the empty set, thus avoiding n 's expansion (lines 6–7).

Before expanding the current AND node n , its cache table is checked (line 18). If the same context was encountered before, it is retrieved from the cache, and $\text{succ}(n)$ is set to the empty set, which will trigger the PROPAGATE step. Otherwise, the node is expanded in the usual way, depending on whether it is an AND or OR node (lines 8–32). The algorithm also computes the heuristic evaluation function for every partial solution subtree rooted at the OR ancestors of n along the path from the root (lines 20–24). The search below n is terminated if, for some OR ancestor m , $f(T'_m) \geq v(m)$, where $v(m)$ is the current upper bound on the optimal cost below m . The recursive computation of $f(T'_m)$ is described in Algorithm 2.

The node values are updated by the PROPAGATE step (lines 34–45). It is triggered when a node has an empty set of descendants (note that as each successor is evaluated, it is removed from the set of successors in line 44). This means that all its children have been evaluated, and their final values are already determined. If the current node is the root, then the search terminates with its value (line 37). If n is an OR node, then its parent p is an AND node, and p updates its current value $v(p)$ by summation with the value of n (line 38). An AND node n propagates its value to its parent p in a similar way, by minimization (lines 42–43). Finally, the current node n is set to its parent p (line 45), because n was completely evaluated. Search continues either with a *propagation* step (if conditions are met) or with an *expansion* step.

AOBB-C-ILP is restricted to a static variable ordering determined by the underlying pseudo tree and explores the context minimal AND/OR search graph via *full caching*. However, if the memory requirements are prohibitive, rather than using full caching, AOBB-C-ILP can be modified to use a memory bounded caching scheme that saves only those nodes whose context size can fit in the available memory, as shown in [6,7].

1
2
3
4 **5 Best-First AND/OR Search**
5
6
7

8 We now direct our attention to a *best-first* rather than depth-first control strategy for
9 traversing the context minimal AND/OR graph and present a best-first AND/OR
10 search algorithm for 0-1 ILP. The algorithm uses similar amounts of memory as
11 the depth-first AND/OR Branch-and-Bound with full caching. The algorithm was
12 described in details in [8,9,7] and evaluated for general constraint optimization
13 problems. By specializing it to 0-1 ILP using the LP relaxation for h , we get
14 AOBF-C-ILP. For completeness sake, we describe the algorithm again including
15 minor modifications for the 0-1 ILP case.
16
17

18 The algorithm, denoted by AOBF-C-ILP (Algorithm 3), specializes Nilsson’s
19 AO* algorithm [21] to AND/OR search spaces for 0-1 ILPs. It interleaves for-
20 ward expansion of the best partial solution tree (`EXPAND`) with a cost revision step
21 (`REVISE`) that updates node values, as detailed in [21]. The explicated AND/OR
22 search graph is maintained by a data structure called \mathcal{G}'_T , the current node is n , s is
23 the root of the search graph and the current best partial solution subtree is denoted
24 by T' . The children of a node n are denoted by $\text{succ}(n)$.
25
26

27 First, a top-down, graph-growing operation finds the best partial solution tree by
28 tracing down through the marked arcs of the explicit AND/OR search graph \mathcal{G}'_T
29 (lines 4–9). These previously computed marks indicate the current best partial sol-
30 ution tree from each node in \mathcal{G}'_T . Before the algorithm terminates, the best partial
31 solution tree, T' , does not yet have all of its leaf nodes terminal. One of its non-
32 terminal leaf nodes n is then expanded by generating its successors, depending on
33 whether it is an OR or an AND node. Notice that when expanding an OR node, the
34 algorithm does not generate AND children that are already present in the explicit
35 search graph \mathcal{G}'_T (lines 13–15). All these identical AND nodes in \mathcal{G}'_T are easily
36 recognized based on their contexts. Upon node’s n expansion, a heuristic under-
37 estimate $h(n')$ of $v(n')$ is assigned to each of n ’s successors $n' \in \text{succ}(n)$ (lines
38 12–25). Again, $h(n')$ is obtained by solving the LP relaxation of the subproblem
39 rooted at n' , conditioned on the current partial assignment of the path to the root. As
40 before, AOBF-C-ILP avoids expanding those nodes for which the corresponding
41 LP relaxation is infeasible or yields an integer solution (lines 18–22 and 28–32).
42
43

44 The second operation in AOBF-C-ILP is a bottom-up, cost revision, arc marking,
45 SOLVE-labeling procedure (lines 26–40). Starting with the node just expanded n ,
46 the procedure revises its value $v(n)$, using the newly computed values of its suc-
47 cessors, and marks the outgoing arcs on the estimated best path to terminal nodes.
48 This revised value is then propagated upwards in the graph. The revised value $v(n)$
49 is an updated lower bound estimate of the cost of an optimal solution to the sub-
50 problem rooted at n . During the bottom-up step, AOBF-C-ILP labels an AND
51 node as **SOLVED** if all of its OR child nodes are solved, and labels an OR node as
52 **SOLVED** if its marked AND child is also solved. The algorithm terminates with the
53
54

Algorithm 3: AOBF-C-ILP: Best-First AND/OR Search for 0-1 ILP

Input: A 0-1 ILP instance with objective function $\sum_{i=1}^n c_i X_i$, pseudo tree T rooted at X_1 , AND contexts pas_i for every variable X_i

Output: Minimal cost solution.

```

1   $v(s) \leftarrow h(s); \mathcal{G}'_T \leftarrow \{s\}$                                 // Initialize
2  while  $s$  is not labeled SOLVED do
3       $S \leftarrow \{s\}; T' \leftarrow \emptyset;$                                      // Create the marked partial solution tree
4      while  $S \neq \emptyset$  do
5           $n \leftarrow \text{top}(S)$ ; remove  $n$  from  $S$ 
6           $T' \leftarrow T' \cup \{n\}$ 
7          let  $L$  be the set of marked successors of  $n$ 
8          if  $L \neq \emptyset$  then
9               $L$  add  $L$  on top of  $S$ 

10     let  $n$  be any nonterminal tip node of the marked  $T'$  (rooted at  $s$ )           // EXPAND
11     if  $n$  is an OR node, labeled  $X_i$  then
12         foreach  $x_i \in D_i$  do
13             let  $n'$  be the AND node in  $\mathcal{G}'_T$  having context equal to  $pas_i$ 
14             if  $n' == \text{NULL}$  then
15                 create an AND node  $n'$  labeled  $\langle X_i, x_i \rangle$ 
16                  $h(n') \leftarrow LP(P_{n'})$ ;  $v(n') \leftarrow h(n')$                   // Solve the LP relaxation
17                  $w(n, n') \leftarrow c_i \cdot x_i$                                          // Compute the arc weight
18                 label  $n'$  as INFEASIBLE or INTEGER if the LP relaxation is infeasible or has an integer solution
19                 if  $n'$  is INTEGER or TERMINAL then
20                      $n'$  label  $n'$  as SOLVED
21                 else if  $n'$  is INFEASIBLE then
22                      $v(n') \leftarrow \infty$ 
23
24              $succ(n) \leftarrow succ(n) \cup \{n'\}$ 

25     else if  $n$  is an AND node, labeled  $\langle X_i, x_i \rangle$  then
26         foreach  $X_j \in \text{children}_T(X_i)$  do
27             create an OR node  $n'$  labeled  $X_j$ 
28              $h(n') \leftarrow LP(P_{n'})$ ;  $v(n') \leftarrow h(n')$                   // Solve the LP relaxation
29             label  $n'$  as INFEASIBLE or INTEGER if the LP relaxation is infeasible or has an integer solution
30             if  $n'$  is INTEGER then
31                 mark  $n'$  as SOLVED
32             else if  $n'$  is INFEASIBLE then
33                  $v(n') \leftarrow \infty$ 
34
35              $succ(n) \leftarrow succ(n) \cup \{n'\}$ 

36      $\mathcal{G}'_T \leftarrow \mathcal{G}'_T \cup succ(n)$ 
37      $S \leftarrow \{n\}$                                                        // REVISE
38     while  $S \neq \emptyset$  do
39         let  $m$  be a node in  $S$  such that  $m$  has no descendants in  $\mathcal{G}'_T$  still in  $S$ ; remove  $m$  from  $S$ 
40         if  $m$  is an AND node, labeled  $\langle X_i, x_i \rangle$  then
41              $v(m) \leftarrow \sum_{m' \in succ(m)} v(m')$ 
42             mark all arcs to the successors
43             label  $m$  as SOLVED if all its children are labeled SOLVED
44
45         else if  $m$  is an OR node, labeled  $X_i$  then
46              $v(m) = \min_{m' \in succ(m)} (w(m, m') + v(m'))$ 
47             mark the arc through which this minimum is achieved
48             label  $m$  as SOLVED if the marked successor is labeled SOLVED
49
50         if  $m$  changes its value or  $m$  is labeled SOLVED then
51              $m$  add to  $S$  all those parents of  $m$  such that  $m$  is one of their successors through a marked arc.
52
53
54
55     return  $v(s)$                                                        // Search terminates

```

optimal solution when the root node s is labeled SOLVED. We next summarize the complexity of both depth-first and best-first AND/OR graph search [3,6,8,7]:

1
2
3
4 **THEOREM 3 (complexity)** *Depth-first AND/OR Branch-and-Bound and best-first*
5 *AND/OR search algorithms guided by a pseudo tree T are sound and complete for*
6 *solving 0-1 ILPs. Their time and space complexity is $O(n \cdot 2^{w^*})$, where w^* is the*
7 *induced width of the pseudo tree.*

10 **AOBB versus AOBF.** Best-first search AOBF with the same heuristic function as
11 depth-first Branch-and-Bound AOOB is likely to expand the smallest number of
12 nodes [11], but empirically this depends on how quickly AOOB will find an optimal solution.
13 Second, AOOB can use far less memory by avoiding dead-caches for example (e.g., when the context minimal search graph is a tree), while AOBF has
14 to keep the explicated search graph in memory. Third, AOOB can be used as an
15 anytime scheme, namely whenever interrupted, the algorithm outputs the best solution found so far, unlike AOBF which outputs a solution upon termination only.
16 All the above points show that the relative merit of best-first versus depth-first over
17 context minimal AND/OR search spaces cannot be determined by theory [11] and
18 empirical evaluation is necessary.
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65

6 Dynamic Variable Orderings

The depth-first and best-first AND/OR search algorithms presented in the previous sections assumed a static variable ordering determined by the underlying pseudo tree of the constraint graph. However, the mechanism of identifying unifiable AND nodes based solely on their contexts is hard to extend when variables are instantiated in a different order than that dictated by the pseudo tree. In this section we discuss a strategy that allows dynamic variable orderings in depth-first and best-first AND/OR search, when both algorithms traverse an AND/OR search tree. The approach called *Partial Variable Ordering (PVO)*, which combines the static AND/OR decomposition principle with a dynamic variable ordering heuristic, was described and tested also for general constraint optimization over graphical models in [10,5]. For completeness sake, we review it briefly next.

Variable Orderings for Integer Programming. At every node in the search tree, the search algorithm has to decide what variable to instantiate next. One common method in operations research is to select next the *most fractional variable*, i.e., variable whose LP value is furthest from being integral [26]. Finding a candidate variable under this rule is fast and the method yields small search trees on many problem instances.

A more sophisticated approach, which is better suited for certain hard problems is *strong branching* [27]. This method performs a one-step lookahead for each variable that is non-integral in the LP at the node. The one-step lookahead computation solve the LP relaxations for each of the children of the candidate variable, and a score is computed based on the LP values of the children. The next variable to

1
2
3
4 instantiate is selected as the one with the highest score among the candidates.
5
6

7 **Partial Variable Ordering (PVO).** *AND/OR Branch-and-Bound with Partial Vari-*
8 *able Ordering* (resp. *Best-First AND/OR Search with Partial Variable Ordering*),
9 denoted by AOBB+PVO-ILP (resp. AOBF+PVO-ILP), uses the static graph-based
10 decomposition given by a pseudo tree with a dynamic semantic ordering heuristic
11 applied over chain portions of the pseudo tree. For simplicity and without loss of
12 generality we consider the *most fractional variable* as our semantic variable order-
13 ing heuristic. Clearly, it can be replaced by any other heuristic.
14
15

16 Consider the pseudo tree from Figure 2 inducing the following variable groups
17 (or chains): $\{A, B\}$, $\{C, D\}$ and $\{E, F\}$, respectively. This implies that variables
18 $\{A, B\}$ should be considered before $\{C, D\}$ and $\{E, F\}$. The variables in each
19 group can be dynamically ordered based on the semantic ordering heuristic.
20
21

22 AOBB+PVO-ILP (resp. AOBF+PVO-ILP) can be derived from Algorithm 1 (resp.
23 Algorithm 3) with some simple modifications. As usual, the algorithm traverses an
24 AND/OR search tree in a depth-first (resp. best-first) manner, guided by a pre-
25 computed pseudo tree \mathcal{T} . When the current AND node n , labeled $\langle X_i, x_i \rangle$, is ex-
26 panded in the forward step, the algorithm generates its OR successor m , labeled X_j ,
27 based on the semantic ordering heuristic. Specifically, m corresponds to the most
28 fractional variable in the current pseudo tree chain. If there are no uninstantiated
29 variables left in the current chain, namely variable X_i was instantiated last, then the
30 OR successors of n are labeled by the most fractional variables from the variable
31 groups rooted by X_i in \mathcal{T} .
32
33
34
35

36 37 38 39 40 41 42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60 61 62 63 64 65 7 Experimental Results

We evaluated the performance of the depth-first and best-first AND/OR search algorithms on 0-1 ILP problem classes such as combinatorial auction, uncapacitated warehouse location problems and MAX-SAT problem instances. We implemented our algorithms in C++ and carried out all experiments on a 2.4GHz Pentium IV with 2GB of RAM, running Windows XP.

Algorithms. The detailed outline of the experimental evaluation is given in Table 1. We evaluated the following 6 classes of AND/OR search algorithms:

- 1 Depth-first and best-first search algorithms using a static variable ordering and exploring the AND/OR tree, denoted by AOBB-ILP and AOBF-ILP, respectively.
- 2 Depth-first and best-first search algorithms using dynamic partial variable orderings and exploring the AND/OR tree, denoted by AOBB+PVO-ILP and AOBF+PVO-ILP, respectively.

Table 1
Detailed outline of the experimental evaluation for 0-1 ILP.

Problem classes	Tree search		Graph search	ILP solvers
	AOBB-ILP	AOBB+PVO-ILP	AOBB-C-ILP	BB (lp_solve)
	AOBF-ILP	AOBF+PVO-ILP	AOBF-C-ILP	CPLEX 11.0
Combinatorial Auctions	✓	✓	✓	✓
Warehouse Location Problems	✓	✓	✓	✓
MAX-SAT Instances	✓	✓	✓	✓

3 Depth-first and best-first search algorithms with caching that explore the context minimal AND/OR graph and use static variable orderings, denoted by AOBB-C-ILP and AOBF-C-ILP, respectively.

All of these AND/OR algorithms use a *simplex* implementation based on the open-source `lp_solve` library to compute the guiding LP relaxation. For this reason, we compare them against the OR Branch-and-Bound algorithm available from the `lp_solve` library, denoted by BB. The pseudo tree used by the AND/OR algorithms was constructed using the hypergraph partitioning heuristic described in [10,5] and outlined briefly below. BB, AOBB+PVO-ILP and AOBF+PVO-ILP used a dynamic variable ordering heuristic based on *reduced costs* (*i.e.*, dual values) [1]. Specifically, the next fractional variable to instantiate has the smallest reduced cost in the solution of the LP relaxation. Ties are broken lexicographically.

We note however that the AOBB-ILP and AOBB-C-ILP algorithms support a restricted form of dynamic variable and value ordering. Namely, there is a dynamic internal ordering of the successors of the node just expanded, before placing them onto the search stack. Specifically, in line 33 of Algorithm 1, if the current node n is AND, then the independent subproblems rooted by its OR children can be solved in decreasing order of their corresponding heuristic estimates (variable ordering). Alternatively, if n is OR, then its AND children corresponding to domain values can also be sorted in decreasing order of their heuristic estimates (value ordering).

For reference, we also ran the ILOG CPLEX version 11.0 solver (with default settings), which uses a best-first control strategy, dynamic variable ordering heuristic based on strong branching, as well as cutting planes to tighten the LP relaxation. It explores however an OR search tree.

In the MAX-SAT domain we ran, in addition, three specialized solvers:

- 1 MaxSolver [28], a DPLL-based algorithm that uses a 0-1 non-linear integer formulation of the MAX-SAT problem,
- 2 toolbar [29], a classic OR Branch-and-Bound algorithm that solves MAX-SAT as a Weighted CSP problem [30], and
- 3 PBS [31], a DPLL-based solver capable of propagating and learning pseudo-boolean constraints as well as clauses.

1
2
3
4 MaxSolver and toolbar were shown to perform very well on random MAX-
5 SAT instances with high graph connectivity [29], whereas PBS exhibits better per-
6 formance on relatively sparse MAX-SAT instances [28]. These algorithms explore
7 an OR search space.
8

9
10 Throughout our empirical evaluation we will address the following questions that
11 govern the performance of the proposed algorithms:
12

- 13
14 1 The impact of AND/OR versus OR search.
15 2 The impact of best-first versus depth-first AND/OR search.
16 3 The impact of caching.
17 4 The impact of dynamic variable orderings.
18

19
20 **Constructing the Pseudo Tree.** Our heuristic for generating a low height balanced
21 pseudo tree is based on the recursive decomposition of the dual hypergraph associ-
22 ated with the 0-1 ILP instance. The dual hypergraph of a 0-1 ILP with \mathbf{X} variables
23 and \mathbf{F} constraints is a pair (\mathbf{V}, \mathbf{E}) where each constraint in \mathbf{F} is a vertex $v_i \in \mathbf{V}$ and
24 each variable in \mathbf{X} is a hyperedge $e_j \in \mathbf{E}$ connecting all the constraints (vertices)
25 in which it appears.
26

27
28 Generating heuristically good hypergraph separators can be done using a package
29 called hMetIS³, which we used following [32]. The vertices of the hypergraph
30 are partitioned into two balanced (roughly equal-sized) parts, denoted by \mathcal{H}_{left} and
31 \mathcal{H}_{right} respectively, while minimizing the number of hyperedges across. A small
32 number of crossing edges translates into a small number of variables shared
33 between the two sets of constraints. \mathcal{H}_{left} and \mathcal{H}_{right} are then each recursively par-
34 titioned in the same fashion, until they contain a single vertex. The result of this
35 process is a tree of hypergraph separators which can be shown to also be a pseudo
36 tree of the original model where each separator corresponds to a subset of variables
37 chained together [10,5].
38

39
40 **Measures of Performance.** We report CPU time (in seconds) and number of nodes
41 visited. We also specify the number of variables (n), the number of constraints (c),
42 the depth of the pseudo trees (h) and the induced width of the graphs (w^*) obtained
43 for each problem instances. The best performance points are highlighted. In each
44 table, “-” denotes that the respective algorithm exceeded the time limit. Similarly,
45 “out” stands for exceeding the 2GB memory limit.
46

47
48
49
50
51
52
53

7.1 Combinatorial Auctions

54
55

56
57 In **combinatorial auctions** (CA), an auctioneer has a set of goods, $M = \{1, 2, \dots, m\}$
58 to sell and the buyers submit a set of bids, $\mathcal{B} = \{B_1, B_2, \dots, B_n\}$. A bid is a tuple
59

60
61
62
63
64
65

³ Available at: <http://www-users.cs.umn.edu/karypis/metis/hmetis>

$B_j = (S_j, p_j)$, where $S_j \subseteq M$ is a set of goods and $p_j \geq 0$ is a price. The winner determination problem is to label the bids as winning or loosing so as to maximize the sum of the accepted bid prices under the constraint that each good is allocated to at most one bid. The problem can be formulated as a 0-1 ILP, as follows:

$$\begin{aligned}
 & \max \sum_{j=1}^n p_j x_j \\
 \text{s.t. } & \sum_{j|i \in S_j} x_j \leq 1 \quad i \in \{1..m\} \\
 & x_j \in \{0, 1\} \quad j \in \{1..n\}
 \end{aligned} \tag{8}$$

Combinatorial auctions can also be formulated as binary Weighted CSPs [30], as described in [14]. Therefore, in addition to the 0-1 ILP solvers, we also ran `toolbar` which is a specialized OR Branch-and-Bound algorithm that maintains a level of local consistency called *existential directional arc-consistency* [33].

Figures 6 and 7 display the results for experiments with combinatorial auctions drawn from the *regions-upv* (Figure 6) and *arbitrary-upv* (Figure 7) distributions of CATS 2.0 test suite [34]. The *regions-upv* problem instances simulate the auction of radio spectrum in which a government sells the right to use specific segments of spectrum in different geographical areas. The *arbitrary-upv* problem instances simulate the auction of various electronic components. The suffix *upv* indicates that the bid prices were drawn from a *uniform* distribution. We looked at moderate size auctions having 100 goods and increasing number of bids. The number of bids is also the number of variables in the 0-1 ILP model. Each data point represents an average over 10 instances drawn uniformly at random from the respective distribution. The header of each plot in Figures 6 and 7 shows the average induced width and depth of the pseudo trees.

AND/OR vs. OR search. When comparing the AND/OR versus OR search regimes, we observe that both depth-first and best-first AND/OR search algorithms improve considerably over the OR search algorithm, BB, especially when the number of bids increases and the problem instances become more difficult. In particular, the depth-first and best-first AND/OR search algorithm using partial variable orderings, AOBB+PVO-ILP and AOBF+PVO-ILP, are the winners on this domain, among the `lp_solve` based solvers. For example, on the *regions-upv* auctions with 400 bids (Figure 6), AOBF+PVO-ILP is on average about 8 times faster than BB. Similarly, on the *arbitrary-upv* auctions with 280 bids (Figure 7), the difference in running time between AOBB+PVO-ILP and BB is about 1 order of magnitude. Notice that on the *regions-upv* dataset, `toolbar` is outperformed significantly by BB as well as the AND/OR algorithms. On the *arbitrary-upv* dataset, `toolbar` outperforms dramatically the `lp_solve` based solvers. However, the size of the search space explored by `toolbar` is significantly larger than the ones explored




Fig. 6. Comparing depth-first and best-first AND/OR search algorithms with static and dynamic variable orderings. CPU time in seconds (top) and number of nodes visited (bottom) for solving combinatorial auctions from the **regions-upv** distribution with 100 goods and increasing number of bids. Time limit 3 hours.





Fig. 7. Comparing depth-first and best-first AND/OR search algorithms with static and dynamic variable orderings. CPU time in seconds (top) and number of nodes visited (bottom) for solving combinatorial auctions from the *arbitrary-upv* distribution with 100 goods and increasing number of bids. Time limit 3 hours.



55 Fig. 8. Comparison with CPLEX. CPU time in seconds (top) and number of nodes (bottom)
56 between toolbar and three other solvers for solving combinatorial auctions from the **regions-upv** distribution with 100 goods
57 and increasing number of bids. Time limit 3 hours.




Fig. 9. Comparison with CPLEX. CPU time in seconds (top) and number of nodes visited (bottom) for solving combinatorial auctions from the **arbitrary-upv** distribution with 100 goods and increasing number of bids. Time limit 3 hours.

1
2
3
4 by the AND/OR algorithms. Therefore, toolbar's better performance in this case
5 can be explained by the far smaller computational overhead of the arc-consistency
6 based heuristic used, compared with the LP relaxation based heuristic.
7

8
9 **AOBB vs. AOBF.** When comparing further best-first versus depth-first AND/OR
10 search, we see that AOBF-ILP (resp. AOBF+PVO-ILP) improve considerably
11 over AOBB-ILP (resp. AOBB+PVO-ILP), especially on the *regions-upv* dataset.
12 The gain observed when moving from depth-first AND/OR Branch-and-Bound to
13 best-first AND/OR search is primarily due to the optimal cost, which bounds the
14 horizon of best-first more effectively than for depth-first search.
15

16
17 **Impact of caching.** When looking at the impact of caching on AND/OR search,
18 we notice that the graph search algorithms AOBB-C-ILP and AOBF-C-ILP ex-
19 panded the same number of nodes as the tree search algorithms AOBB-ILP and
20 AOBF-ILP, respectively (see Figures 6 and 7). This indicates that, for this do-
21 main, the context minimal AND/OR search graph explored is a tree. Or, the LP
22 relaxation is very accurate in this case and the AND/OR algorithms only explore
23 a small part of the pseudo tree, for which the corresponding context-based cache
24 entries are actually dead-caches.
25

26
27 **Impact of dynamic variable orderings.** We can see that using dynamic variable
28 ordering heuristics improves the performance of best-first AND/OR search only.
29 For depth-first AND/OR search, the performance deteriorated sometimes (see for
30 example AOBB-ILP vs. AOBB+PVO-ILP on *regions-upv* auctions in Figure 6).
31

32
33 **Comparison with CPLEX.** In Figures 8 and 9 we contrast the results obtained with
34 CPLEX, toolbar, BB, AOBB+PVO-ILP and AOBF+PVO-ILP on the *regions-
35 upv* (Figure 8) and *arbitrary-npv* (Figure 9) distributions, respectively. Clearly, we
36 can see that CPLEX is the best performing solver on these datasets. In particular, it
37 is several orders of magnitude faster than the `lp_solve` based solvers, especially
38 the baseline BB solver. Its excellent performance is leveraged by the powerful cut-
39 ting planes engine as well as the proprietary variable ordering heuristic used. Note
40 that on the *arbitrary-upv* dataset, toolbar is competitive with CPLEX only for
41 relatively small number of bids.
42

43
44 Section A in the appendix provides additional experimental results on combina-
45 torial auctions derived from the *regions-npv* and *arbitrary-npv* distributions. The
46 suffix *nrv* indicates that the bid prices were drawn from a *normal* distribution.
47

48
49

7.2 Uncapacitated Warehouse Location Problems

50

51
52 In the **uncapacitated warehouse location problem** (UWLP) a company considers
53 opening m warehouses at some candidate locations in order to supply its n exist-
54 ing stores. The objective is to determine which warehouse to open, and which of
55

these warehouses should supply the various stores, such that the sum of the maintenance and supply costs is minimized. Each store must be supplied by exactly one warehouse. The typical 0-1 ILP formulation of the problem is as follows:

$$\min \sum_{j=1}^n \sum_{i=1}^m c_{ij} x_{ij} + \sum_{i=1}^m f_i y_i \quad (9)$$

$$\begin{aligned} \text{s.t. } & \sum_{i=1}^m x_{ij} = 1 \quad \forall j \in \{1..n\} \\ & x_{ij} \leq y_i \quad \forall j \in \{1..n\}, \forall i \in \{1..m\} \\ & x_{ij} \in \{0, 1\} \quad \forall j \in \{1..n\}, \forall i \in \{1..m\} \\ & y_i \in \{0, 1\} \quad \forall i \in \{1..m\} \end{aligned}$$

where f_i is the cost of opening a warehouse at location i and c_{ij} is the cost of supplying store j from the warehouse at location i .

Table 2 display the results obtained for 16 randomly generated UWLP instances⁴ with 50 warehouses, 200 and 400 stores, respectively. The warehouse opening and store supply costs were chosen uniformly randomly between 0 and 1000. These are large problems with 10,050 variables and 10,500 constraints for the *uwlp-50-200* class, and 20,050 variables and 20,400 constraints for the *uwlp-50-400* class, respectively, having induced widths of 50 and pseudo trees with depths of 123.

AND/OR vs. OR search. When looking at AND/OR versus OR search, we can see that in almost all test cases the AND/OR algorithms dominate BB. On the *uwlp-50-200-013* instance, for example, AOBF+PVO-ILP causes a speed-up of 186 over BB, exploring a search tree 1,142 times smaller. Similarly, on the *uwlp-50-400-001* instance, AOBB+PVO-ILP outperforms BB by almost 2 orders of magnitude in terms of running time and size of the search space explored. On this domain, the best performing algorithm among the lp_solve based solvers is best-first AOBF+PVO-ILP.

AOBB vs. AOBF. When comparing best-first against depth-first AND/OR Branch-and-Bound search we observe only minor savings in running time in favor of best-first search. This can be explained by the already small enough search space traversed by the algorithms, which does not leave room for additional improvements due to the optimal cost bound exploited by best-first search.

Impact of caching. When looking at the impact of caching we see again that AOBB-C-ILP and AOBF-C-ILP visited the same number of nodes as AOBB-ILP and AOBF-ILP, respectively (see columns 3 and 5 in Table 2). This shows again that the context minimal AND/OR search graph explored by the AOBB-C-ILP

⁴ Problem generator from <http://www.mpi-sb.mpg.de/units/ag1/projects/benchmarks/UflLib/>

Table 2

CPU time in seconds and number of nodes visited for solving Uncapacitated Warehouse Location Problems with 50 warehouses 200 (top part) and 400 (bottom part) stores, respectively. No time limit.

uwlp	BB (lp_solve) CPLEX		AOBB-ILP AOBF-ILP		AOBB+PVO-ILP AOBF+PVO-ILP		AOBB-C-ILP AOBF-C-ILP	
	time	nodes	time	nodes	time	nodes	time	nodes
50 warehouses 200 locations: (n=10,050, c=10,500), (w*=50, h=123)								
uwlp-50-200-004	61.08 0.80	142 0	46.39 37.58	46 24	17.47 15.49	10 3	46.42 36.27	46 24
uwlp-50-200-005	1591.89 9.91	1,692 81	404.94 287.64	233 97	125.81 145.53	50 37	405.72 270.99	233 97
uwlp-50-200-011	256.19 7.97	358 37	233.96 88.22	246 41	78.74 75.83	39 22	233.21 83.75	246 41
uwlp-50-200-013	13693.76 8.94	14,846 37	116.19 111.28	44 26	78.86 74.53	24 13	116.25 105.72	44 26
uwlp-50-200-017	711.04 2.15	998 3	123.14 48.06	118 21	18.17 16.84	9 2	124.70 47.77	118 21
uwlp-50-200-018	1477.74 5.74	2,666 8	161.03 54.58	146 21	59.52 32.33	37 8	161.05 52.41	146 21
uwlp-50-200-020	2179.39 7.47	3,668 28	190.77 87.58	138 33	68.91 48.33	36 10	190.81 83.70	138 33
uwlp-50-200-021	3252.60 6.66	5,774 25	609.74 80.55	580 30	37.63 46.80	9 7	608.24 92.08	580 30
50 warehouses 400 locations: (n=20,050, c=20,400), (w*=50, h=123)								
uwlp-50-400-001	13638.55 10.76	12,548 12	743.75 130.03	374 20	106.63 81.63	29 8	743.68 126.39	374 20
uwlp-50-400-004	820.89 6.52	942 6	1114.47 126.97	794 25	55.10 51.85	10 3	1117.55 123.19	794 25
uwlp-50-400-005	57532.67 30.55	32,626 58	2719.09 331.87	617 36	247.03 131.58	50 8	2722.26 313.09	617 36
uwlp-50-400-006	365.93 3.59	632 0	48.41 51.62	11 8	32.31 32.65	1 1	48.44 51.95	11 8
uwlp-50-400-008	599.49 3.40	560 0	175.60 119.28	49 13	96.66 60.27	21 3	175.67 116.42	49 13
uwlp-50-400-009	17608.98 9.02	17,262 6	281.02 132.27	76 14	97.00 78.05	9 2	281.30 128.58	76 14
uwlp-50-400-011	22727.61 8.07	22,324 7	193.91 93.11	77 12	64.28 64.58	5 4	193.89 92.06	77 12
uwlp-50-400-012	5468.30 4.49	4,174 0	671.90 164.64	307 32	52.22 52.95	4 2	671.77 159.28	307 32

1
2
3
4 and AOBF-C-ILP algorithms was a tree and therefore all cache entries were dead-
5 caches.
6

7
8 **Impact of dynamic variable orderings.** We also observe that the dynamic variable
9 ordering had a significant impact on performance in this case, especially for depth-
10 first search. For example, on the *uwlp-50-200-021* instance, AOBB+PVO-ILP is
11 16 times faster than AOBB-ILP and expands 64 times fewer nodes. However, the
12 difference in running time between the best-first search algorithms, AOBF-ILP
13 and AOBF+PVO-ILP, is smaller compared to what we see for depth-first AND/OR
14 search. This is because the search space explored by AOBF-ILP is already small
15 enough and the savings in number of nodes caused by dynamic variable orderings
16 cause only minor time savings.
17

18
19 **Comparison with CPLEX.** When looking at the results obtained with CPLEX
20 (column 2 in Table 2), we notice again its excellent performance in terms of both
21 running time and size of the search space explored. However, we see that in some
22 cases AOBF+PVO-ILP actually explored fewer nodes than CPLEX (e.g., *uwlp-50-*
23 *200-021*). This is important because it shows that the relative worse performance of
24 AOBF+PVO-ILP versus CPLEX is due mainly to the much slower *simplex* imple-
25 mentation of the former, lack of cutting planes engine as well as the naive dynamic
26 variable ordering heuristic used.
27

28
29
30
31
32
33 **7.3 MAX-SAT Instances**
34
35

36 Given a set of Boolean variables the goal of **maximum satisfiability** (MAX-SAT)
37 is to find a truth assignment to the variables that violates the least number of clauses.
38 We experimented with problem classes *pret* and *dubois* from the SATLIB⁵ library,
39 which were previously shown to be difficult for 0-1 ILP solvers [29].
40

41 MAX-SAT can be formulated as a 0-1 ILP [35] or pseudo-Boolean formula [36,37].
42 In the 0-1 ILP model, a Boolean variable v is mapped to an integer variable x that
43 takes value 1 when v is *True* or 0 when it is *False*. Similarly, $\neg v$ is mapped to
44 $1 - x$. With these mappings, a clause can be formulated as a linear inequality. For
45 example, the clause $(v_1 \vee \neg v_2 \vee v_3)$ can be mapped to $x_1 + (1 - x_2) + x_3 \geq 1$.
46 Here, the inequality means that the clause must be satisfied in order for the left side
47 of the inequality to have a value no less than one.
48

49 However, a clause in a MAX-SAT may not be satisfied, so that the corresponding
50 inequality may be violated. To address this issue, an auxiliary integer variable y is
51 introduced to the left side of a mapped inequality. Variable $y = 1$ if the correspond-
52 ing clause is unsatisfied, making the inequality valid; otherwise, $y = 0$. Since the
53 objective is to minimize the number of violated clauses, it is equivalent to mini-
54

55
56
57
58
59
60 ⁵ <http://www.satlib.org/>
61
62
63
64
65

Table 3

CPU time in seconds and number of nodes visited for solving *pret* MAX-SAT instances.

Time limit 10 hours.

pret (w*, h)	BB (lp_solve)		MaxSolver time	toolbar PBS		AOBB-ILP		AOBB+PVO-ILP		AOBB-C-ILP		
	CPLEX			PBS		AOBF-ILP		AOBF+PVO-ILP		AOBF-C-ILP		
	time	nodes		time	nodes	time	nodes	time	nodes	time	nodes	
pret60-40 (6, 13)	-	-	9.47	53.89 0.004	7,297,773 565	7.88	1,255	8.41	1,216	7.38	1,216	
pret60-60 (6, 13)	-	-	9.48	53.66 0.004	7,297,773 495	8.56	1,259	8.70	1,247	7.30	1,140	
pret60-75 (6, 13)	-	-	9.37	53.52 0.003	7,297,773 543	6.97	1,124	6.80	1,089	6.34	1,067	
pret150-40 (6, 15)	-	-	-	- 0.02	- 2,592	95.11	6,625	108.84	7,152	75.19	5,625	
pret150-60 (6, 15)	-	-	-	- 0.01	- 2,873	98.88	6,851	112.64	7,347	78.25	5,813	
pret150-75 (6, 15)	-	-	-	- 0.02	- 2,898	108.14	7,311	115.16	7,452	84.97	6,114	
	out					98.95	6,282	103.03	6,394	20.95	1,430	

mize the sum of the auxiliary variables that are forced to take value 1. For example, $(v_1 \vee \neg v_2 \vee v_3), (v_2 \vee v_4)$ can be written as an 0-1 ILP of minimizing $z = y_1 + y_2$, subject to the constraints of $x_1 + (1 - x_2) + x_3 + y_1 \geq 1$ and $x_2 + (1 - x_4) + y_2 \geq 1$.

7.3.1 pret Instances

Table 3 shows the results for experiments with 6 *pret* instances. These are unsatisfiable instances of graph 2-coloring with parity constraints. The size of these problems is relatively small (60 variables with 160 clauses for *pret60* and 150 variables with 400 clauses for *pret150*, respectively).

AND/OR vs. OR search. When comparing AND/OR versus OR search we see again that the AND/OR algorithms improved dramatically over BB. For instance, on the *pret150-75* network, AOBB-ILP finds the optimal solution in less than 2 minutes, whereas BB exceeds the 10 hour time limit. Similarly, MaxSolver and toolbar could not solve the instance within the time limit. Overall, PBS offers the best performance on this dataset.

AOBB vs. AOBF. The best-first AND/OR search algorithms improve sometimes considerably over the depth-first ones, especially when exploring an AND/OR graph (e.g., see AOBF-C-ILP versus AOBB-C-ILP in the leftmost column of Table 3). Moreover, the search space explored by AOBF-C-ILP appears to be the smallest. This indicates that the computational overhead of AOBF-C-ILP is mainly due to evaluating its guiding lower bounding heuristic evaluation function.

Impact of caching. When looking at the depth-first AND/OR Branch-and-Bound

graph search algorithm we only observe minor improvements due to caching. This is probably because most of the cache entries were actually dead-caches. On the other hand, best-first AOBF-C-ILP exploits the relatively small size of the context-minimal AND/OR graph (*i.e.*, in this case the problem structure is captured by a very small context with size 6 and a shallow pseudo tree with depth 13 or 15) and achieves the best performance among the ILP solvers.

Impact of dynamic variable orderings. We also see that the dynamic variable ordering did not have an impact on search performance for both depth-first and best-first algorithms.

Comparison with CPLEX. Both depth-first and best-first AND/OR search algorithms outperformed dramatically CPLEX on this dataset. On the *pret60-40* instance, for example, AOBF-C-ILP is 2 orders of magnitude faster than CPLEX. Similarly, on *pret150-40*, CPLEX exceeded the memory limit.

7.3.2 *dubois Instances*

Figure 10 displays the results for experiments with random *dubois* instances with increasing number of variables. These are unsatisfiable 3-SAT instances with $3 \times \text{degree}$ variables and $8 \times \text{degree}$ clauses, each of them having 3 literals. As in the previous test case, the *dubois* instances have very small contexts of size 6 and shallow pseudo trees with depths ranging from 10 to 20.

AND/OR vs. OR search. As before, we see that the AND/OR algorithms are far superior to BB, which could not solve any of the test instances within the 3 hour time limit. PBS is again the overall best performing algorithm, however it failed to solve 4 test instances: on instance *dubois130*, for which $\text{degree} = 130$, it exceeded the 3 hour time limit, whereas on instances *dubois180*, *dubois200* and *dubois260* the clause/pseudo-boolean constraint learning mechanism caused the solver to run out of memory. We note that MaxSolver and toolbar were not able to solve any of the test instances within the time limit.

AOBB vs. AOBF. Best-first search outperforms again depth-first search, especially when exploring the AND/OR graph. However, the depth-first tree search algorithms AOBB-ILP and AOBB+PVO-ILP were better than the best-first tree search counterparts in this case. This was probably caused by the internal dynamic variable ordering used by AOBB-ILP and AOBB+PVO-ILP to solve independent subproblems rooted at the AND nodes in the search tree.

Impact of caching. We can see that AOBF-C-ILP takes full advantage of the relatively small context minimal AND/OR search graph and, on some of the larger instances, it outperforms its ILP competitors with up to one order of magnitude in terms of both running time and number of nodes expanded. On this dataset as well AOBF-C-ILP explores the smallest search space, but its computational overhead




Fig. 10. Comparing depth-first and best-first AND/OR search algorithms with static and dynamic variable orderings. CPU time in seconds (top) and number of nodes visited (bottom) for solving ***dubois*** MAX-SAT instances. Time limit 3 hours. CPLEX, BB, toolbar and MaxSolver were not able to solve any of the test instances within the time limit.

1
2
3
4 does not pay off in terms of running time when compared with PBS. The impact of
5 caching on AND/OR Branch-and-Bound is not that pronounced as for best-first.
6
7
8

9 **Impact of dynamic variable orderings.** The dynamic variable ordering had a mi-
10 nor impact on depth-first AND/OR search only (*e.g.*, see AOBB+PVO-ILP versus
11 AOBB-ILP in Figure 10).

12 **Comparison with CPLEX.** The performance of CPLEX was quite poor on this
13 dataset and could not solve any of the test instances within the time limit.
14
15
16
17
18
19

20 8 Related Work

21
22

23 The idea of exploiting structural properties of the problem in order to enhance
24 the performance of search algorithms in constraint satisfaction is not new. Freuder
25 and Quinn [22] introduced the concept of pseudo tree arrangement of a constraint
26 graph as a way of capturing independencies between subsets of variables. Subse-
27 quently, *pseudo tree search* is conducted over a pseudo tree arrangement of the
28 problem which allows the detection of independent subproblems that are solved
29 separately. More recently, [38] extended pseudo tree search [22] to optimization
30 tasks in order to boost the Russian Doll search [39] for solving Weighted CSPs.
31 Our depth-first AND Branch-and-Bound and best-first AND/OR search algorithms
32 for 0-1 ILPs are also related to the Branch-and-Bound method proposed by [25]
33 for acyclic AND/OR graphs and game trees, as well as the depth-first and best-
34 first AND/OR search algorithms for general constraint optimization over graphical
35 models introduced in [4,6,10,8,9,5,7].
36
37
38
39

40 Dechter’s graph-based back-jumping algorithm [40] uses a depth-first (DFS) span-
41 ning tree to extract knowledge about dependencies in the graph. The notion of
42 DFS-based search was also used by [41] for a distributed constraint satisfaction al-
43 gorithm. Bayardo and Miranker [23] reformulated the pseudo tree search algorithm
44 in terms of back-jumping and showed that the depth of a pseudo tree arrangement
45 is always within a logarithmic factor off the induced width of the graph.
46
47

48 *Recursive Conditioning* (RC) [32] is based on the divide and conquer paradigm.
49 Rather than instantiating variables to obtain a tree structured network like the cycle
50 cutset scheme, RC instantiates variables with the purpose of breaking the network
51 into independent subproblems, on which it can recurse using the same technique.
52 The computation is driven by a data-structure called *dtree*, which is a full binary
53 tree, the leaves of which correspond to the network functions. It can be shown that
54 RC explores an AND/OR space [3]. A pseudo tree can be generated from the static
55 ordering of RC dictated by the *dtree*. This ensures that whenever RC splits the
56 problem into independent subproblems, the same happens in the AND/OR space.
57
58
59
60
61
62
63
64

1
2
3
4 *Backtracking with Tree-Decomposition* (BTD) [42] is a memory intensive method
5 for solving constraint satisfaction and optimization problems which combines search
6 techniques with the notion of tree decomposition. This mixed approach can in fact
7 be viewed as searching an AND/OR search space whose backbone pseudo tree is
8 defined by and structured along the tree decomposition [3]. What is defined in [42]
9 as structural goods, that is parts of the search space that would not be visited again
10 as soon as their consistency (or optimal value) is known, corresponds precisely to
11 the decomposition of the AND/OR space at the level of AND nodes, which root in-
12 dependent subproblems. The BTD algorithm is not linear space, it uses substantial
13 caching and may be exponential in the treewidth.
14
15

16
17 *Value Elimination* [43] is a recently developed algorithm for Bayesian inference. It
18 was already explained in [43] that, under static variable ordering, there is a strong
19 relation between Value Elimination and Variable Elimination. Given a static order-
20 ing d for Value Elimination, it can be shown that it actually traverses an AND/OR
21 space [3]. The pseudo tree underlying the AND/OR search graph traversal by Value
22 Elimination can be constructed as the bucket tree in reversed d . However, the traver-
23 sal of the AND/OR space will be controlled by d , advancing the frontier in a hybrid
24 depth or breadth first manner.
25
26
27

31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60 5 **9 Summary and Conclusion**

30
31 The paper investigates the impact of the AND/OR search spaces perspective to
32 solving optimization problems from the class of 0-1 Integer Linear Programs. In
33 earlier papers [4,6,8,9,5,7] we showed that the AND/OR search paradigm can im-
34 prove general constraint optimization algorithms. Here, we demonstrate empiri-
35 cally the benefit of AND/OR search to 0-1 Integer Linear Programs.
36
37

38 Specifically, we extended and evaluated the depth-first and best-first AND/OR search
39 algorithm traversing the AND/OR search tree or context minimal AND/OR graph
40 to solving 0-1 ILPs. We also extended the algorithms with dynamic variable or-
41 dering strategies. Our empirical evaluation demonstrated on a variety of 0-1 ILP
42 benchmark problems that the AND/OR search algorithms outperform the clas-
43 sic depth-first OR Branch-and-Bound sometimes by several orders of magnitude.
44 We summarize next the most important factors influencing performance, including
45 dynamic variable orderings, caching, as well as the search control strategy (*e.g.*,
46 depth-first versus the best-first).
47
48

- 49
50
51
52
53
54
55 • **Depth-first versus best-first search.** Our results showed conclusively that the
56 AND/OR search algorithms using a best-first control strategy and traversing ei-
57 ther an AND/OR search tree or graph were able, in many cases, to improve con-
58 siderably over the depth-first search ones (*e.g.*, combinatorial auctions from Fig-
59 ures 6 and A.1, *dubois* MAX-SAT instances from Figure 10).

- **Impact of caching.** For problems with relatively small contexts (treewidth), the memory intensive best-first AND/OR search algorithms were shown to outperform dramatically the corresponding tree search algorithms (*e.g.*, *dubois* MAX-SAT instances from Figure 10). The impact of caching on the depth-first AND/OR Branch-and-Bound search algorithms was less prominent on these types of problems (*e.g.*, *pret* and *dubois* MAX-SAT instances from Table 3 and Figure 10, respectively) probably because most of the cache entries were dead-caches. For problems with very large contexts (*e.g.*, combinatorial auctions from Figures 6 and A.1, UWLP instances from Table 2) the context minimal AND/OR graph explored was a tree, and therefore caching had no impact.
- **Impact of dynamic variable orderings.** The AND/OR search approach was already shown to be powerful when used in conjunction with dynamic variable ordering schemes [10,5]. Here, for 0-1 ILPs we also show that the AND/OR Branch-and-Bound with partial variable orderings sometimes outperformed the AND/OR Brach-and-Bound restricted to a static variable ordering by one order of magnitude (*e.g.*, UWLP instances from Table 2). Similarly, best-first AND/OR search with partial variable orderings improved considerably over its counterpart using a static ordering (*e.g.*, combinatorial auctions from Figures 6 and A.1).
- **AND/OR solvers versus CPLEX.** Our current implementation of the depth-first and best-first AND/OR search is far from being fully optimized with respect to commercial 0-1 ILP solvers such as CPLEX, as it relies on an open source implementation of the *simplex* algorithm, as well as a naive dynamic variable ordering heuristic. Nevertheless, we demonstrated that on selected classes of 0-1 ILPs the AND/OR algorithms outperformed CPLEX in terms of both the number of nodes explored (*e.g.*, UWLP instances from Table 2) and CPU time (*e.g.*, *pret* MAX-SAT instances from Table 3).

Future Work. Our depth-first and best-first AND/OR search approach can accommodate all known enhancement schemes. In particular, it can be modified to incorporate *cutting planes* to tighten the linear relaxation of the current subproblem. The space required by the best-first AND/OR search can be enormous, due to the fact that all the nodes generated by the algorithm have to be saved prior to termination. Therefore, the algorithm can be extended to incorporate a memory bounding scheme [21,44,45]. We can also incorporate good initial upper bound techniques (using incomplete schemes), which in some cases can allow a best-first performance using depth-first AND/OR Branch-and-Bound algorithms, as we demonstrated in [7].

Acknowledgments

This work was partially supported by the NSF grants IIS-0086529 and IIS-0412854, the MURI ONR award N00014-00-1-0617 and the NIH grant R01-HG004175-02.

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65

References

- [1] G. Nemhauser and L. Wolsey. *Integer and combinatorial optimization*. Wiley, 1988.
- [2] E. Lawler and D. Wood. Branch-and-bound methods: A survey. *Operations Research*, 14(4):699–719, 1966.
- [3] R. Dechter and R. Mateescu. AND/OR search spaces for graphical models. *Artificial Intelligence*, 171(2-3):73–106, 2007.
- [4] R. Marinescu and R. Dechter. And/or branch-and-bound for graphical models. In *International Joint Conference on Artificial Intelligence (IJCAI)*, pages 224–229, 2005.
- [5] R. Marinescu and R. Dechter. And/or branch-and-bound search for combinatorial optimization in graphical models. *Technical Report, University of California, Irvine (under review)*, 2008.
- [6] R. Marinescu and R. Dechter. Memory intensive branch-and-bound search for graphical models. In *National Conference on Artificial Intelligence (AAAI)*, 2006.
- [7] R. Marinescu and R. Dechter. Memory intensive and/or search for combinatorial optimization in graphical models. *Technical Report, University of California, Irvine (under review)*, 2008.
- [8] R. Marinescu and R. Dechter. Best-first and/or search for graphical models. In *National Conference on Artificial Intelligence (AAAI)*, pages 1171–1176, 2007.
- [9] R. Marinescu and R. Dechter. Best-first and/or search for most probable explanations. In *International Conference on Uncertainty in Artificial Intelligence (UAI)*, 2007.
- [10] R. Marinescu and R. Dechter. Dynamic orderings for and/or branch-and-bound search in graphical models. In *European Conference on Artificial Intelligence (ECAI)*, pages 138–142, 2006.
- [11] R. Dechter and J. Pearl. Generalized best-first search strategies and the optimality of a^* . *Journal of the ACM*, 32(3):505–536, 1985.
- [12] R. Marinescu and R. Dechter. And/or branch-and-bound search for pure 0/1 integer linear programming problems. In *International Conference on Integration of AI and OR Techniques in Constraint Programming for Combinatorial Optimization (CPAIOR)*, pages 152–166, 2006.
- [13] R. Marinescu and R. Dechter. Best-first and/or search for 0/1 integer programming. In *International Conference on Integration of AI and OR Techniques in Constraint Programming for Combinatorial Optimization (CPAIOR)*, 2007.
- [14] Rina Dechter. *Constraint Processing*. MIT Press, 2003.
- [15] R. Karp. Reducibility among combinatorial problems. *Complexity of Computer Computations*, Plenum Press, NY, pages 85–103, 1972.

- [16] G.B. Dantzig. Maximization of a linear function of variables subject to linear inequalities. *Activity Analysis of Production and Allocation*, 1951.
- [17] N. Nilsson P. Hart and B. Raphael. A formal basis for the heuristic determination of minimum cost paths. *IEEE Transactions on Systems Science and Cybernetics*, 8(2):100–107, 1968.
- [18] J. Wolsey. *Integer Programming*. John Wiley & Sons, 1998.
- [19] M. Padberg and G. Rinaldi. Optimization of a 532-city symmetric traveling salesman problem by branch-and-cut. *Operations Research Letters*, 6(1):1–7, 1987.
- [20] M. Padberg and G. Rinaldi. A branch-and-cut algorithm for the resolution of large scale symmetric traveling salesman problems. *SIAM Review*, 33(1):60–100, 1991.
- [21] Nils J. Nilsson. *Principles of Artificial Intelligence*. Tioga, 1980.
- [22] E. Freuder and M. Quinn. Taking advantage of stable sets of variables in constraint satisfaction problems. In *International Joint Conference on Artificial Intelligence (IJCAI)*, pages 1076–1078, 1985.
- [23] R. Bayardo and D. Miranker. On the space-time trade-off in solving constraint satisfaction problems. In *International Joint Conference on Artificial Intelligence (IJCAI)*, pages 558–562, 1995.
- [24] J. Pearl. *Heuristics: Intelligent search strategies for computer problem solving*. Addison-Wesley, 1984.
- [25] L. Kanal and V. Kumar. *Search in artificial intelligence*. Springer-Verlag., 1988.
- [26] Laurence Wosley. *Integer Programming*. John Wiley & Sons, 1998.
- [27] V. Chvatal D. Applegate, R. Bixby and W. Cook. Finding cuts in the tsp (a preliminary report). *Technical Report 95-05, DIMACS, Rutgers University*, 1995.
- [28] Z. Xing and W. Zhang. Efficient strategies for (weighted) maximum satisfiability. In *Principles and Practice of Constraint Programming (CP)*, pages 660–705, 2004.
- [29] S. de Givry, J. Larrosa, and T. Schiex. Solving max-sat as weighted csp. In *Principles and Practice of Constraint Programming (CP)*, 2003.
- [30] S. Bistarelli, U. Montanari, and F. Rossi. Semiring based constraint solving and optimization. *Journal of the ACM*, 44(2):309–315, 1997.
- [31] F. Aloul, A. Ramani, I. Markov, and K. Sakallah. Pbs: A backtrack search pseudo-boolean solver. In *Symposium on the Theory and Applications of Satisfiability Testing (SAT)*, 2002.
- [32] A. Darwiche. Recursive conditioning. *Artificial Intelligence*, 126(1-2):5–41, 2001.
- [33] S. de Givry, F. Heras, J. Larrosa, and M. Zytnicki. Existential arc consistency: getting closer to full arc consistency in weighted csps. In *International Joint Conference in Artificial Intelligence (IJCAI)*, 2005.

- 1
2
3
4 [34] K. Leyton-Brown, M. Pearson, and Y. Shoham. Towards a universal test suite for
5 combinatorial auction algorithms. In *ACM Electronic Commerce*, pages 66–76, 2000.
6
7 [35] S. Joy, J. Mitchell, and B. Borchers. A branch and cut algorithm for max-sat and
8 weighted max-sat. In *Satisfiability Problem: Theory and Applications*, pages 519–
9 536, 1997.
10
11 [36] J.P. Walser. *Integer Optimization Local Search*. Springer, 1999.
12
13 [37] H. Dixon and M.L. Ginsberg. Inference methods for a pseudo-boolean satisfiability
14 solver. In *National Conference on Artificial Intelligence (AAAI)*, pages 635–640, 2006.
15
16 [38] J. Larrosa, P. Meseguer, and M. Sanchez. Pseudo-tree search with soft constraints. In
17 *European Conference on Artificial Intelligence (ECAI)*, pages 131–135, 2002.
18
19 [39] M. Lemaitre G. Verfaillie and T. Schiex. Russian doll search for solving constraint
20 optimization problems. In *National Conference on Artificial Intelligence (AAAI)*,
21 pages 298–304, 1996.
22
23
24 [40] R. Dechter. Enhancement schemes for constraint processing: Backjumping, learning
25 and cutset decomposition. *Artificial Intelligence*, 41(3):273–312, 1990.
26
27 [41] Z. Collin, R. Dechter, and S. Katz. On the feasibility of distributed constraint
28 satisfaction. In *International Joint Conference on Artificial Intelligence (IJCAI)*, pages
29 318–324, 1991.
30
31
32 [42] P. Jegou and C. Terrioux. Decomposition and good recording for solving max-csp.
33 In *European Conference on Artificial Intelligence (ECAI)*, pages 196–200, 2004.
34
35 [43] F. Bacchus, S. Dalmao, and T. Pittaci. Value elimination: Bayesian inference via
36 backtracking search. In *Uncertainty in Artificial Intelligence (UAI)*, pages 20–28,
37 2003.
38
39 [44] P. Chakrabati, S. Ghose, A. Acharya, and S. de Sarkar. Heuristic search in restricted
40 memory. *Artificial Intelligence*, 41(3):197–221, 1989.
41
42 [45] R. Korf. Linear-space best-first search. *Artificial Intelligence*, 62(1):41–78, 1993.
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64

1
2
3
4 **A Additional Results on Combinatorial Auctions**
5
6
7
8
9
10
11
12
13
14

Figures A.1 and A.2 show the results for experiments with combinatorial auctions generated from the *regions-npv* (Figure A.1) and *arbitrary-npv* (Figure A.2) distributions of the CATS 2.0 suite. As mentioned earlier, the bid prices of these auctions were drawn from a *normal* rather than the uniform distribution used in Section 7.1. As before, each data point represents an average over 10 random instances.

The spectrum of results is similar to what we observed in Section 7.1. The AND/OR algorithms outperformed BB by a significant margin. Caching had no impact on these datasets as well, namely the context minimal AND/OR graph explored was a tree (in Figures A.1 and A.2 the curves corresponding to graph search AOBB-C-ILP and AOBF-C-ILP overlap with those corresponding to tree search AOBB-ILP and AOBF-ILP). On the *arbitrary-npv* dataset, `toolbar` outperformed again the `lp_solve` based solvers, indicating that in this case the EDAC heuristic had a far smaller overhead than the LP based one.

Figures A.3 and A.4 show the results obtained with CPLEX on the *regions-npv* (Figure A.3) and *arbitrary-npv* (Figure A.4) distributions, respectively. Clearly, we can see that CPLEX is the best performing solver on these datasets. It is several orders of magnitude faster than all other ILP solvers. On the *arbitrary-npv* auctions, `toolbar` is competitive with CPLEX only for relatively small number of bids.




Fig. A.1. Comparing depth-first and best-first AND/OR search algorithms with static and dynamic variable orderings. CPU time in seconds (top) and number of nodes visited (bottom) for solving combinatorial auctions from the **regions-npv** distribution with 100 goods and increasing number of bids. Time limit 3 hours.




Fig. A.2. Comparing depth-first and best-first AND/OR search algorithms with static and dynamic variable orderings. CPU time in seconds (top) and number of nodes visited (bottom) for solving combinatorial auctions from the *arbitrary-npv* distribution with 100 goods and increasing number of bids. Time limit 3 hours.




Fig. A.3. Comparison with CPLEX. CPU time in seconds (top) and number of nodes visited (bottom) for solving combinatorial auctions from the **regions-npv** distribution with 100 goods and increasing number of bids. Time limit 3 hours.




Fig. A.4. Comparison with CPLEX. CPU time in seconds (top) and number of nodes visited (bottom) for solving combinatorial auctions from the **arbitrary-npv** distribution with 100 goods and increasing number of bids. Time limit 3 hours.