

Continuously Relaxing Over-constrained Conditional Temporal Problems through Generalized Conflict Learning and Resolution

Keywords:

Search in Planning and Scheduling,
Constraints, Satisfiability, and Search::Applications,
Constraints, Satisfiability, and Search::Solvers and Tools.

Abstract

Over-constrained temporal problems are commonly encountered while operating autonomous and decision support systems. An intelligent system must learn the humans' preferences over the problem in order to generate preferred resolutions that minimize perturbations. We present the Best-first Conflict-Directed Relaxation (BCDR) algorithm for enumerating the best continuous relaxation for an over-constrained Conditional Temporal Problem (CTP) with controllable choices. BCDR reformulates such a problem by replacing its relaxable constraints with fuzzy temporal constraints and solves the problem using a conflict-directed approach. It extends the traditional Conflict-directed A* (CD-A*) algorithm to conditional temporal problems, by first generalizing the conflict learning process to include all discrete variable assignments and continuous temporal constraints, and then by guiding the forward search away from known infeasible regions using conflict resolutions. When evaluated empirically on a range of coordinated car sharing network problems, BCDR demonstrates a substantial improvement in performance and solution quality compared to previous conflict-direct approaches.

1 Introduction

Temporal constraint networks [Dechter *et al.*, 1991] have been widely used to model planning and scheduling problems in daily life. They have been used to describe and reason over conditional and uncertain situations with multiple alternative plans. However, a solution to a temporal problem does not always exist. For example, in a car sharing network scenario, a user needs four hours to complete his shopping trip but only has three hours of reservation time. It is not enough for a scheduling program to just signal a failure. Instead, it should explain the situation and propose alternative plans for the user so that he can make a more informed decision, either to extend the reservation or to drop goals. Such a scenario is usually framed as an over-constrained temporal problem, and the goal is to find one or a set of preferred relaxations to the temporal constraints in the problem so that a consistent schedule can be found.

Prior work on over-constrained temporal problems starts with [Beaumont *et al.*, 2001], in which Partial Constraint Satisfaction techniques [Freuder and Wallace, 1992] are implemented to find the subset of temporal constraints that can be satisfied. Later, disjunctions and optimality were added in the context of over-constrained Disjunctive Temporal Problems with Preferences (DTPPs) [Peintner *et al.*, 2005]. In a DTPP, the disjuncts of every constraint are assigned a preference function that maps the temporal constraint to a cost value. The optimal partial solution is obtained by enumerating consistent subproblems using Branch & Bound, as well as optimization techniques introduced in [Khatib *et al.*, 2001]. However, most of the prior work has focused on restoring consistency through complete suspension of constraints, while in real-world scenarios, the user often wants to preserve as much of the schedule as possible and minimize the perturbations.

In this paper, we present our continuous relaxation approach, the Best-first Conflict-Directed Relaxation algorithm (BCDR), to address this issue. BCDR efficiently resolves over-constrained Conditional Temporal Problems with controllable variables. It reformulates an over-constrained temporal problem by replacing its relaxable constraints with fuzzy temporal constraints: their bounds can be relaxed to restore consistency. BCDR uses a conflict-directed strategy similar to Conflict-directed A* [Williams and Ragno, 2002] to enumerate continuous relaxations in best-first order: it learns conflicts between constraints and variable assignments, and uses the resolutions to these conflicts to guide the search away from infeasible regions.

Note that this paper is not concerned about the dynamic or weak consistency of Conditional Temporal Problems with uncontrollable discrete variables (CTPs and CTPPs, [Tsamardinos, 2003; Falda *et al.*, 2007]). We are only concerned about controllable variables that are not dependent on observation events. Solving such a problem is simpler than determining the dynamic/weak consistency of a CTP in that those tasks may require the enumeration of all possible scenarios.

2 Example

To motivate the need for continuously relaxing over-constrained temporal problems, we describe an example in the domain of a coordinated car sharing network, such as Zipcar. Such a network provides an hourly rental service to its members: a rental car may be used by multiple members in a

day. Each member must time their usage well so that the car can be returned on time. Otherwise, the next reservation will be affected and a penalty fee will be applied.

Consider the following example on John's trip for grocery shopping and lunch. He has reserved a car from 11am to 2pm, and is planning to go to one of the two grocery stores nearby: A or B. John has different preferences over them and their shopping times vary from 35 minutes to 50 minutes. After grocery shopping, John would like to have lunch at a restaurant, either X, Y or Z, before going home. Lunch takes different amount of time in each restaurant. Finally, driving times to these locations are different, and John has to return his car back home in three hours (11am to 2pm) so that the next person can start her trip on time.

We model John's trip with a Controllable Conditional Temporal Problem (CCTP), and use it to determine the best strategy for him that includes: which grocery store to visit, which restaurant to dine at, how much time he can spend at each location and whether he should extend his reservation. We start by defining two variables for the decisions he needs to make: GS (Grocery store) and RT (Restaurant). GS has two options in its domain: A (40) and B (100). Each option is associated with a positive reward value that represents John's preferences towards it, the larger the better. The other variable RT has three options: X (70), Y (80) and Z (30).

Next, we define twelve events with time points for the problem (Table 1): a reference point in time (S_T) that represents the beginning of the trip at 11am; a time point that indicates the end of the trip (R_T); and time points representing the arrival and departure of each locations (store A and B, restaurant X, Y and Z).

Events with Time Points		
Trip starts	S_T	Store A Arrive/Leave
Trip ends	R_T	Store B Arrive/Leave
Restaurant X Arrive/Leave		X_A, X_L
Restaurant Y Arrive/Leave		Y_A, Y_L
Restaurant Z Arrive/Leave		Z_A, Z_L

Table 1: Events in John's trip

Constraints (in minutes)		
$C_1: A_L - A_A \geq 40$	$C_6: A_A - S_T \in [35, 50]$	$GS \leftarrow A$
$C_2: B_L - B_A \geq 35$	$C_7: B_A - S_T \in [35, 40]$	$GS \leftarrow B$
$C_3: X_L - X_A \geq 50$	$C_8: R_T - X_L \in [45, 50]$	$RT \leftarrow X$
$C_4: Y_L - Y_A \geq 75$	$C_9: R_T - Y_L \in [40, 50]$	$RT \leftarrow Y$
$C_5: Z_L - Z_A \geq 100$	$C_{10}: R_T - Z_L \in [50, 60]$	$RT \leftarrow Z$
C_{11}	$X_A - A_L [30, 40]$	$GS \leftarrow A \text{ and } RT \leftarrow X$
C_{12}	$Y_A - A_L [25, 30]$	$GS \leftarrow A \text{ and } RT \leftarrow Y$
C_{13}	$Z_A - A_L [20, 25]$	$GS \leftarrow A \text{ and } RT \leftarrow Z$
C_{14}	$X_A - B_L [35, 40]$	$GS \leftarrow B \text{ and } RT \leftarrow X$
C_{15}	$Y_A - B_L [25, 40]$	$GS \leftarrow B \text{ and } RT \leftarrow Y$
C_{16}	$Z_A - B_L [30, 35]$	$GS \leftarrow B \text{ and } RT \leftarrow Z$
C_{17}	$R_T - S_T \in [0, 180]$	

Table 2: Conditional Temporal Constraints in the CCTP

(Table 2) shows all the conditional temporal constraints in the CCTP that encode the temporal relaxations between events. Constraints C_1 through C_5 are linear constraints that represent John's desired length of stay at five locations. For example, $B_L - B_A \geq 35$ indicates that John would like to spend at least 35 minutes at store B. These constraints are labeled by the assignments made to the decision variables: a

constraint is activated only if its label assignment is made. For example, C_2 will be considered only if John chooses to shop at B, as shown in the right side of (Table 2). Constraints C_6 through C_{16} are simple temporal constraints that encode the driving time required between locations. They are conditioned on assignments made to either GS or RT, or both (C_{11} through C_{16}). Finally, C_{17} constrains the duration of John's trip to three hours.

Some of the constraints highlighted in bold (C_1 through C_5 and C_{17}) are fuzzy temporal constraints. They can be relaxed in order to restore the consistency of the problem, if necessary. Each fuzzy constraint comes with one or two cost functions that describe John's preferences towards the relaxations for the upper and lower bounds. These functions map the relaxation from LB to LB' , or from UB to UB' , to a positive cost value, as seen in (Figure 1). If the upper bound of C_{17} is relaxed from 180 minutes to 200 minutes, meaning that John delays his return by 20 minutes, the cost will be 40. On the other hand, if he shortens his lunch time by relaxing the lower bound of C_3 to 30, the cost would be 100. In this example, we assume that all other fuzzy constraints have linear cost functions with gradient 1.




Figure 1: Preference functions for C_3 and C_{17}

Relaxation 1	Relaxation 2	Relaxation 3
$GS \leftarrow B$	$GS \leftarrow B$	$GS \leftarrow B$
$RT \leftarrow Y$	$RT \leftarrow X$	$RT \leftarrow Y$
$C_4 \text{ to } \geq 50$	$C_3 \text{ to } \geq 48$	$C_4 \text{ to } \geq 55$
$C_{17} \text{ to } \in [0, 185]$	$C_2 \text{ to } \geq 17$	$C_2 \text{ to } \geq 25$
Utility: 152.5	Utility: 151	Utility: 150
$\searrow \text{Do not relax } C_{17} \swarrow$	$\searrow C_2 \text{ is at least 25} \swarrow$	

Table 3: Three preferred relaxations to the CCTP

Before relaxing any constraints, there is no consistent solution to the problem. The cause of failure is that three hours is not enough for John to complete both shopping and dining tasks: driving to the nearest grocery store and restaurant will consume at least 100 minutes, which brings the minimum trip duration to 200 minutes. Therefore, one or more temporal constraints need to be relaxed. (Table 3) shows three consistent relaxations for the CCTP ranked in best-first order. Relaxation 1 is first presented to John that decides to shop at B and have lunch in Y. The lunch time should be reduced to 50 minutes and the reservation be extended by 5 minutes. The utility of the relaxation is 152.5, which is computed by summing up the reward of two assignments and subtracting the cost of relaxing C_4 and C_{17} . If John changes his mind and decides not to relax C_{17} , Relaxation 2 will be generated which incorporates this new requirement. It takes John to X for lunch, shortens the lunch time to 48 minutes and reduces the shopping time to 17 minutes. If John is still unsatisfied, he may add an additional requirement that shopping time should be no less than 25 minutes. BCDR will continue the search

and present Relaxation 3, which respects both newly added requirements.

This example demonstrates the advantage of continuous relaxations: minimizing perturbations. Compared to discrete relaxations, which may suggest that John not shop or have lunch, continuous relaxations minimally relax the constraints to restore consistency. In addition, the reactive nature of BCDR makes it possible to adapt to newly added constraints, enumerate relaxations accordingly and enable the user to adjust his or her requirements on-line.

3 Problem Statement

Temporal problems with choices are usually modeled using Conditional Temporal Problems (CTPs,[Tsamardinos, 2003]). It is a generalization of the restricted problem class of Simple Temporal Problems (STPs,[Dechter *et al.*, 1991]) by adding uncontrollable discrete choices and by conditioning the occurrence of events and simple temporal constraints on the outcomes of these choices. CTP is capable of modeling conditional plans and uncertainty during executions.

Definition 1. A CTP is a 6-tuple $\langle V, E, L, OV, O, P \rangle$ where:

- P is a set of Boolean atomic propositions;
- V is a set of events representing designated time points;
- E is a set of simple temporal constraints that restricts the time points in V , such as $l_{ij} \leq v_i - v_j \leq u_{ij}, l_{ij}, u_{ij} \in \mathcal{R}$;
- $L : V \rightarrow Q$ is a function that attaches conjunctions of literals of P , $q_i \in Q$, to each event $v_i \in V$;
- $OV \subseteq V$ is a set of observation events that provides the truth value for $p_i \in P$ through function $O : P \rightarrow OV$.

In a CTP, each event is associated with a conjunctive set of literals, called a *label*. If the label of an event is evaluated to be true, the event is said to be activated and needs to be scheduled. Otherwise, the event and its associated temporal constraints can be ignored.

The solution to a CTP is a schedule that assigns a time point to each event in the CTP and is consistent with the temporal constraints. There are three notions of CTP consistency: Strong, Dynamic and Weak consistency, depending on the assumptions made over the outcomes of observation variables [Tsamardinos, 2003]. Conditional Temporal Problems with Preferences (CTPPs,[Falda *et al.*, 2007]) extend CTPs by allowing fuzzy temporal constraints and fuzzy atomic propositions. This allows the user to specify preferences over the execution time of each event $v_i \in V$, and compare two schedules T_1 and T_2 using a preference function that maps a schedule to a value function $f : T \rightarrow \mathcal{R}^+$ [Khatib *et al.*, 2001].

The problem that BCDR addresses, Controllable Conditional Temporal Problems (CCTPs), shares most of the notation with CTPs and CTPPs. However, there are two important differences. First, we assume that all variables are controllable. To determine the consistency of a CCTP, it is sufficient to find one consistent set of discrete variable assignments. Second, CCTP extends the domains of discrete variables from binary to any finite domains, and allows the discrete variables to be conditioned on assignments to other variables.

Definition 2. A CCTP is a 8-tuple $\langle V, E, RE, L_v, L_p, P, f_v, f_e \rangle$ where:

- P is a set of controllable finite domain discrete variables;
- V is a set of events representing designated time points;
- E is a set of temporal constraints between pairs of events $v_i \in V$;
- $RE \in E$ is a set of fuzzy temporal constraints whose bounds can be relaxed.
- $L_v : V \rightarrow Q$ is a function that attaches conjunctions of assignments to P , $q_i \in Q$, to some events $v_i \in V$;
- $L_p : P \rightarrow Q$ is a function that attaches conjunctions of assignments to P , $q_i \in Q$, to some variables $p_i \in P$;
- $f_p : Q \rightarrow \mathcal{R}^+$ is a function that maps each assignment to every controllable discrete variables, $q_{ij} : p_i \leftarrow \text{value}_j$, to a positive reward value.
- $f_e : (e_i, e'_i) \rightarrow \mathcal{R}^+$ is a function that maps the relaxation to one fuzzy temporal constraint $e_i \in RE$, from e_i to e'_i , to a positive cost value.

To allow the relaxation for an over-constrained temporal problem, we include fuzzy temporal constraints in the definition of CCTP, similar to the soft constraints in a Simple Temporal Problem with Preferences (STPP,[Rossi *et al.*, 2002]). We do not use a disjunctive set of temporal bounds for soft constraints. Instead, the constraint is soft in that its lower or upper bounds can be relaxed at the price of increasing cost. The cost is defined over the degree of relaxation made to the lower and upper bounds.

There are two preference functions, f_p and f_e . f_p is a reward function over the assignments to controllable discrete variables $p_i \in P$. Each assignment is mapped to a positive reward value, such as $RT \leftarrow X : 50$. The larger the number is, the more preferred the choice will be. f_e is a positive cost function defined over fuzzy constraints. The cost of relaxing an upper bound constraint $E_{ij} : v_i - v_j \leq u_{ij}$ from u_{ij} to u'_{ij} is $f_{eij}(u'_{ij} - u_{ij})$. (Figure 1(b)) shows an example function defined over $u'_{ij} - u_{ij}$.

For temporal constraints restricting the lower bounds between two events, the cost function is $f_{eij}(l_{ij} - l'_{ij})$, such as (Figure 1(a)). We assume that the user always prefers smaller relaxations, therefore all f_e functions must be monotonically increasing, and equal to 0 when there is no relaxation. Therefore, f_e can be viewed as a semi-convex [Khatib *et al.*, 2001] function with a segment of zero cost when there is no relaxation. This assumption simplifies our relaxation process, as the tightest relaxation will always result in the lowest cost. For fuzzy simple temporal constraints, two separate cost functions are required for the lower and upper bounds.

We define the solution to a CCTP as a pair $\langle A, R \rangle$, where:

- A is a complete set of assignments to some discrete variables in P that leaves no variable unassigned.
- R is a set of relaxed bounds of some fuzzy constraints in RE .

such that the CCTP is temporally consistent. The utility of a relaxation is computed by subtracting the relaxation cost from the assignment reward: $\sum_{p_i} f_{p_i}(p_i \leftarrow \text{value}_i) - \sum_{e_i} f_{e_i}(e_i \rightarrow e'_i)$. The most preferred relaxation to a CCTP is the one with the highest utility value according to f_p and f_e .

Note that CCTP is similar, though different in notations, to the Optimal Conditional Simple Temporal Problem (OCSTP) formulation introduced by [Effinger, 2006]. OCSTP uses a conditional CSP-like formulation in which conditional temporal constraints are encoded as the domain values of the discrete variables. A depth-first strategy is used to find a set of variable assignments that enables a consistent set of temporal constraints. However, it is difficult to encode the fuzzy temporal constraints using an OCSTP formulation. We will compare the performance of its search strategy to BCDR in Section 5.

4 Approach

In this section, we present the Best-first Conflict-Directed Relaxation algorithm that enumerates the relaxations to a CCTP in best-first order. This can be viewed as an extension to the Conflict-Directed A* algorithm [Williams and Ragno, 2002] by generalizing the conflicts learning and resolution capability. CD-A* enumerates likely solutions to discrete domain CSPs with conflicts learned from inconsistent sets of assignments. Once detected, a conflict is used to prune the search space by extending each partial candidate with its resolutions. To resolve a CCTP using the conflict-directed strategy, we have to first generalize the conflicts to include conditional and temporal constraints, then generate both discrete and continuous constituent relaxations to the conflict. We will first give an overview of the BCDR algorithm, and then discuss the conflict learning and resolution in detail.

4.1 The BCDR algorithm

BCDR takes an A* search strategy by evaluating each partial candidate using an admissible heuristic function and expanding the search tree in best-first order. The first relaxation found is guaranteed to be the best one. It uses two types of expansions to explore the search space, *Expand on an unassigned variable* and *Expand on a new conflict*, which differentiates BCDR from previous relaxation algorithms. The pseudo code of BCDR is given in (Algorithm 1).

BCDR starts with an empty candidate in the queue (Line 1). A candidate is a 4-tuple $\langle A, R, C_r, C_{cont} \rangle$ with assignments A , relaxations R , resolved conflicts C_r and continuously resolved conflicts C_{cont} , all being empty lists in the first candidate. BCDR continues looping until the first relaxation is found that makes the CCTP consistent (Line 11). If BCDR does not find a consistent relaxation until the queue is exhausted, it returns null indicating that no relaxation exists for the input CCTP (Line 24).

Within each loop, BCDR first dequeues the best partial candidate (Line 6). It checks if $Cand$ resolves all known conflicts (Line 3). If not, an unresolved conflict $currCFT$ will be returned by function `RESOLVEKNOWNCONFLICTS?`, which compares the resolved conflicts C_r in $Cand$ with all known conflicts C . $currCFT$ is then used for expanding $Cand$ by function `EXPANDONCONFLICT` (Line 21). The child candidates of $Cand$ will then be enqueued.

If $Cand$ resolves all known conflicts, BCDR then checks if it is complete by comparing its assignments and all unassigned variables in the CCTP (Line 9). If $Cand$ is incomplete, BCDR will expand it using the assignments to

Input: A CCTP $T = \langle V, E, RE, L_v, L_p, P, f_v, f_e \rangle$.

Output: A relaxation $\langle A, R \rangle$ that maximizes $f_v - f_e$.

Initialization:

```

1  $Cand \leftarrow \langle A, R, C_r, C_{cont} \rangle$ ; the first candidate;
2  $Q \leftarrow \{Cand\}$ ; a priority queue that records candidates;
3  $C \leftarrow \{\}$ ; the set of all known conflicts;
4  $U \leftarrow V$ ; the list of unassigned controllable variables;
Algorithm:
5 while  $Q \neq \emptyset$  do
6    $Cand \leftarrow \text{Dequeue}(Q)$ ;
7    $currCFT \leftarrow \text{RESOLVEKNOWNCONFLICTS?}(Cand, C)$ ;
8   if  $currCFT == \text{null}$  then
9     if  $\text{isComplete?}(Cand, U)$  then
10        $newCFT \leftarrow \text{CONSISTENCYCHECK}(cand)$ ;
11       if  $newCFT == \text{null}$  then
12          $\mid$  return  $Cand$ ;
13       else
14          $C \leftarrow C \cup \{newCFT\}$ ;
15          $Q \leftarrow Q \cup \{Cand\}$ ;
16       endif
17     else
18        $Q \leftarrow Q \cup \text{EXPANDONVARIABLE}\{Cand, U\}$ 
19     endif
20   else
21      $Q \leftarrow Q \cup \text{EXPANDONCONFLICT}\{Cand, currCFT\}$ ;
22   endif
23 end
24 return  $null$ ;
```

Algorithm 1: The BCDR algorithm

one unassigned variable through function `EXPANDONVARIABLE` (Line 18). For example, assume that we need to expand a partial candidate $\{GS=A, RT=X\}$ with variable $FD:\{Steak, Salmon\}$, we simply create two child candidates that extends the partial candidate using two possible assignments of FD (Figure 2(a)). The expanded candidates will be added back to Q .




Figure 2: Example of expanding on candidate and conflict

If $Cand$ is complete, BCDR proceeds to check its consistency using function `CONSISTENCYCHECK` (Line 10). If no conflict is returned, $Cand$ will be returned as the best relaxation (Line 12). If a new conflict, $newCFT$, is detected by `CONSISTENCYCHECK`, BCDR will record it and put $Cand$ back to the queue for future expansions (Line 14,15).

4.2 Learning Conflicts through Negative Cycles

Given a complete candidate that assigns all active discrete variables, function `CONSISTENCYCHECK` checks the consistency of all activated temporal constraints. BCDR implements the Incremental Temporal Consistency algorithm [hsiang Shu et al., 2005] for checking temporal consistency. If the set of temporal constraints is inconsistent, ITC will return a simple negative cycle as the cause of failure. We can extract the minimal inconsistent set of temporal constraints, also called minimal conflict [Liffiton et al., 2005], using this simple negative cycle. For example, (Figure 3) shows a simple negative cycle detected in John's trip: the reservation time

is too tight for activities at B and Y .




Figure 3: A negative cycle in John’s trip

Previous approaches [Effinger and Williams, 2005; Li and Williams, 2005] only extract the discrete variable assignments as conflict, which is $\{GS = B; RT = Y\}$ in this case. Since we are looking for relaxations to temporal constraints, we should include them in the conflict as well, and use their relaxations to resolve the conflict. In addition, a temporal constraint may depend on one or more assignments: its label must be included in the conflict as well. In short, BCDR learns a conflict from a simple negative cycle. A conflict is composed of the temporal constraints involved in the cycle and the assignments required to activate them. For example, the generalized conflict we can learn from (Figure 3) is:

Assignments: $GS=B; RT=Y;$

Constraints: $RT-ST \in [0, 180]; GS=B \rightarrow BA-ST \in [35, 40]; GS=B \rightarrow BL-BA \geq 35; GS=B \wedge RT=Y \rightarrow YA-BL \in [25, 40]; RT=Y \rightarrow YL-YA \geq 75; RT=Y \rightarrow RT-YL \in [40, 50];$

4.3 Generalized Conflict Resolutions

Given a minimal conflict, we can compute their resolutions and use them to expand existing candidates so that future expansions of the candidates will not enter the infeasible region represented by this conflict again. This is the core principle behind conflict-directed search. Previous approaches generate the resolutions, which are called constituent relaxations, by either flipping the assignments to the discrete variables [Williams and Ragni, 2002; Effinger and Williams, 2005] or suspending temporal constraints [Moffitt and Pollack, 2005]. BCDR generalizes the conflict resolution to include both discrete assignments and temporal constraint relaxations: the more we can learn from a conflict, the larger infeasible region we may avoid in the forward search. In addition, we would like to relax the temporal constraints continuously to the minimal extent, instead of completely suspending them, in order to minimize the perturbations.

Function EXPANDONCONFLICT is presented in (Algorithm 2). The resolution is separated into two stages: First, we generate constituent relaxations by negating variable assignments (Line 3-12). If a variable v_i is conditioned on other assignments, in addition to flipping the assignment to v_i , we can also negate its label. This deactivates the variable and resolves the conflict. For example, for a conflict that involves assignment $FD_Y = Steak$, if we know that variable FD_Y has label $RT = Y$, we can resolve the conflict by flipping the assignment of either FD_Y or $RT:FD_Y=Salmon, RT=X$ or $RT=Z$.

In the second stage, we compute the optimal continuous relaxation to the fuzzy temporal constraints that can resolve the conflict (Line 13-19). We formulate the relaxation as an optimization problem with linear constraints (Line 13) and semi-convex objective function (Line 14). The objective function is the minimization over the sum of the relaxation costs of all fuzzy constraints. The variables in this optimization problem

Input: A candidate to expand $Cand \langle A, R, C_r, C_{cont} \rangle$ and a minimal conflict $currCFT$.

Output: A set of expanded candidates $newCands$.

Initialization:

```

1  $newCands \leftarrow \{\}$ ;
2  $CFTs \leftarrow C_{cont} \cup \{currCFT\}$ ; conflicts to be resolved
   continuously;
3 Algorithm:
4   for  $a \in A$  do
5      $A_{alter} \leftarrow A_{alter} \cup \text{GETALTERNATIVES}(a)$ ;
6      $A_{alter} \leftarrow A_{alter} \cup \text{GETALTERNATIVES}(\text{label}(a))$ ;
7   end
8   for  $a_{extend} \in A_{alter}$  do
9     if NOTCOMPETING( $A, a_{extend}$ ) then
10       $Cand_{new} \leftarrow \langle A \cup \{a_{extend}\}, R, C_r, C_{cont} \rangle$ ;
11       $newCands \leftarrow newCands \cup Cand_{new}$ ;
12    end
13  end
14   $\langle E_{fuzzy}, N_{value} \rangle \leftarrow \text{EXTRACTCONSTRAINTS}(CFTs)$ ;
15   $f_{obj} \leftarrow \sum_{e \in E_{fuzzy}} f_e(\Delta e)$ ;
16  if  $R_{new} \neq \text{null}$  then
17     $Cand_{new} \leftarrow \langle A, R_{new}, C_r, C_{cont} \rangle$ ;
18     $newCands \leftarrow newCands \cup Cand_{new}$ ;
19  end
20 return  $newCands$ ;
```

Algorithm 2: Function EXPANDONCONFLICT

are ΔLB_i s and ΔUB_i s, which are the relaxations applied to each fuzzy temporal constraint. They are non-negative and their sum must compensate for the negative value of the conflict (Line 15). The optimal relaxation will not over-relax any constraints, due to the semi-convex assumption over preference functions. It is sufficient to relax the fuzzy constraints to the extent that just eliminates the negative cycle, that is:

$$\begin{aligned} \min \sum_{i \in \text{conflict}} (f_{ei}(u'_{ij} - u_{ij}) + f_{ei}(l'_{ij} - l_{ij})) \\ \text{s.t. } \sum_{i \in \text{conflict}} (e'_{ij} - e_{ij}) = -1 \times N_{value} \end{aligned}$$

For example, the conflict in (Figure 3) involves six constraints. The negative value for this conflict is -30. Among the six constraints, three of them are fuzzy constraints whose bounds can be relaxed: $Reservation \in [0, 180]$, $BL-BA \geq 35$ and $YL-YA \geq 75$. We can define the following optimization problem for computing the continuous relaxation:

$$\begin{aligned} \min(f(\Delta(BL-BA)) + f(\Delta(YL-YA)) + f(\Delta(RT-ST))) \\ \text{s.t. } \Delta(BL-BA) + \Delta(YL-YA) + \Delta(RT-ST) = 30; \end{aligned}$$

The solution to the above optimization problem is a set of relaxed bounds of the fuzzy temporal constraints that resolves the conflict and minimizes the cost. In this case, the best relaxation is: Relax $YL-YA$ to 50 and $Reservation$ to 185. The cost is 27.5. In fact, this problem can also be viewed as a Simple Temporal Problem with Preferences, but on a much smaller scale when compared to the CCTP. [Khatib et al., 2001] demonstrates that finding the optimal solution to a STPP with semi-convex preferences is tractable. In real world applications, we may substitute different optimization algorithms, depending on the preference functions, to improve efficiency.

In total, BCDR generates four constituent relaxations: three new assignments derived from negating assignments and one continuous relaxation. They are used to extend the partial candidates so that future extensions will not run into

the same conflict again, as demonstrated in (Figure 2(b)). Note that one extension is removed due to its conflicting assignments.

5 Experimental Results

To demonstrate the effectiveness of our approach, we present empirical results to compare two BCDR implementations: BCDR-GC (generalized conflict learning and resolution) and BCDR-DC (discrete conflict resolution only). BCDR-DC implements the discrete conflict resolution technique [Li and Williams, 2005]. It computes continuous relaxations once a complete candidate is generated and learns conflicts from an inconsistent set of temporal constraints. Both implementations are set to find the best continuous relaxation. In addition, we compare BCDR-GC to DFS-GC, a depth-first implementation with the generalized conflict resolution technique. DFS-GC is a time-critical alternative to BCDR-GC, and we demonstrate their difference in run-time performance and solution quality.

5.1 Setup

We generated random CCTPs using a simulated car sharing network similar to Section 2. To make it more complex, we extended the problem to allow multiple users and cars: there is always another user waiting for the shared car following each reservation; and there are multiple cars that can be reserved in parallel. In addition, two users using different cars may want to meet during their reservations. We use the following control parameters in the problem generator to control the complexity of a test case:

- N_u : number of users per car. $1 \leq N_u \leq 10$.
- N_c : number of cars available. $1 \leq N_c \leq 12$.
- N_{act} : number of activities per reservation. $1 \leq N_{act} \leq 8$.
- N_{opt} : number of alternatives per activity. $2 \leq N_{opt} \leq 10$.

The total number of discrete variables in a test case is $N_u \times N_c \times N_{act}$, and the domain size of each variable is determined by its N_{opt} . We use the map of Boston to randomly sample the locations. The driving time is computed using an average speed between 30 and 50 mph. The duration at each location and the reservation time, T_{act} and T_{res} , are randomly sampled in $[0, 90]$ and $[0, 360]$ (minutes), respectively. These durations are fuzzy temporal constraints that can be relaxed. We define linear preference functions over these fuzzy constraints with costs sampled between 0 and 10. The reward for each variable assignment (location selection for each activity) ranges from 0 to 1000. Finally, arbitrary temporal constraints are added between cars to simulate a meeting between two users. We use LPSolve as the linear optimizer for all three algorithms [Berkelaar *et al.*,]. In total, 2400 test cases were generated with the number of constraints ranging from 50 to 10000. The time out for each test is 20 seconds, which is usually the maximum time a user would like to wait for a reservation system.

5.2 Results

The results are presented in (Figure 4). Each dot in the graph represents the averaged results computed across all test cases in that category. As can be seen in (Figure 4a), the number of candidates expanded by BCDR-GC before returning

the best relaxation is significantly less than that by BCDR-DC. The generalized conflict learning and resolution efficiently prunes the inconsistent regions in the search domain and avoids nearly 30% of unnecessary candidate expansions compared to discrete conflict resolutions. The reduced number of candidate expansions helps BCDR-GC achieve higher run-time performance compared to BCDR-DC: the average savings is around 10%-15% (Figure 4b). We believe that the run-time performance of BCDR-GC can be further improved if we implement the continuous conflict resolution in an incremental manner, since the EXPANDONCONFLICT function keeps solving optimization problems with a growing set of constraints.




Figure 4: Benchmark results of each algorithm
Next, if the user wants a quick solution, DFS-GC is a good alternative to BCDR-GC. (Figure 4a) shows that DFS-GC expands 50% less candidates than BCDR-GC when the first relaxation is returned, and cuts the run-time by half (Figure 4b). However, the faster results come at a cost of decreased solution quality: the utility of the first solution returned by DFS-GC are 70% lower when compared to that returned by BCDR-GC. If time permits, the user may continue running DFS-GC after obtaining the first relaxation with a Branch & Bound approach, or use BCDR-GC that guarantees to return the best relaxation.

6 Contributions

In this paper, we presented the Best-first Conflict-Directed Relaxation algorithm, the first approach that continuously relaxes over-constrained conditional temporal problems. Compared to previous relaxation algorithms, which restore temporal consistency by suspending constraints, BCDR minimizes the perturbations by continuously relaxing temporal constraints to the minimal extent. It reformulates these problems as Controllable Conditional Temporal Problems, in which relaxable temporal constraints are replaced by fuzzy constraints. With the implementation of generalized conflict learning and resolution, BCDR is more efficient in enumerating the best relaxations when compared to previous conflict-directed approaches, and experimental results have demonstrated its effectiveness in resolving large and highly constrained real-world problems.

References

- [Beaumont *et al.*, 2001] Matthew Beaumont, Abdul Sattar, Michael Maher, and John Thornton. Solving overconstrained temporal reasoning problems. In *Proceedings of the 14th Australian Joint Conference on Artificial Intelligence (AI-2001)*, pages 37–49, 2001.
- [Berkelaar *et al.*,] Michel Berkelaar, Kjell Eikland, and Peter Notebaert. *Ipsolve : Open source (Mixed-Integer) Linear Programming system*.
- [Dechter *et al.*, 1991] Rina Dechter, Itay Meiri, and Judea Pearl. Temporal constraint networks. *Artificial Intelligence*, 49(1-3):61–95, 1991.
- [Effinger and Williams, 2005] Robert Effinger and Brian Williams. Conflict-directed search through disjunctive temporal plan networks. *CSAIL Research Abstracts - 2005*, 2005.
- [Effinger, 2006] Robert T. Effinger. Optimal temporal planning at reactive time scales via dynamic backtracking branch and bond, 2006.
- [Falda *et al.*, 2007] Marco Falda, Francesca Rossi, and K. Brent Venable. Dynamic consistency in fuzzy conditional temporal problems. In *PROC. COPLAS 2007, CP 07 WORKSHOP*, 2007.
- [Freuder and Wallace, 1992] Eugene C. Freuder and Richard J. Wallace. Partial constraint satisfaction. *Artificial Intelligence*, 58(1-3):21–70, 1992.
- [hsiang Shu *et al.*, 2005] I hsiang Shu, Robert Effinger, and Brian C. Williams. Enabling fast flexible planning through incremental temporal reasoning with conflict extraction. In *Proceedings of the 15th International Conference on Automated Planning and Scheduling (ICAPS 05)*, pages 252–261, 2005.
- [Khatib *et al.*, 2001] Lina Khatib, Robert Morris, Robert Morris, and Francesca Rossi. Temporal constraint reasoning with preferences. pages 322–327. Morgan Kaufmann, 2001.
- [Li and Williams, 2005] Hui Li and Brian Williams. Generalized conflict learning for hybrid discrete/linear optimization. In *In CP-2005*, pages 415–429, 2005.
- [Liffiton *et al.*, 2005] M.H. Liffiton, M.D. Moffitt, M.E. Pollack, and K.A. Sakallah. Identifying conflicts in overconstrained temporal problems. In *Proceedings of the 19th International Joint Conference on Artificial Intelligence (IJCAI05)*, pages 205–211, 2005.
- [Moffitt and Pollack, 2005] Michael D. Moffitt and Martha E. Pollack. Partial constraint satisfaction of disjunctive temporal problems. In *Proceedings of the 18th International Florida Artificial Intelligence Research Society Conference (FLAIRS-2005)*, 2005.
- [Peintner *et al.*, 2005] Bart Peintner, Michael D. Moffitt, and Martha E. Pollack. Solving over-constrained disjunctive temporal problems with preferences. In *Proceedings of the 15th International Conference on Automated Planning and Scheduling (ICAPS 2005)*, 2005.
- [Rossi *et al.*, 2002] Francesca Rossi, Alessandro Sperduti, Kristen Brent Venable, Lina Khatib, Paul H. Morris, and Robert A. Morris. Learning and solving soft temporal constraints: An experimental study. In *CP*, pages 249–263, 2002.
- [Tsamardinos, 2003] T. & Pollack M. E. Tsamardinos, I.; Vidal. Ctp: A new constraint-based formalism for conditional, temporal planning. *Constraints*, 8:365–388, 2003.
- [Williams and Rago, 2002] Brian C. Williams and Robert J. Rago. Conflict-directed a* and its role in model-based embedded systems. *Journal of Discrete Applied Mathematics*, 2002.