Power of Pre-Processing: Production Scheduling with Variable Energy Pricing and Power-Saving States

Ondřej Benedikt[1,2], István Módos[1,2], and Zdeněk Hanzálek[1]

[1] Czech Institute of Informatics, Robotics and Cybernetics, Czech Technical University in Prague, Czech Republic
[2] Faculty of Electrical Engineering, Czech Technical University in Prague, Czech Republic
{ondrej.benedikt,istvan.modos,zdenek.hanzalek}@cvut.cz

Abstract. This paper addresses a single machine scheduling problem with non-preemptive jobs to minimize the total electricity cost. Two latest trends in the area of the energy-aware scheduling are considered, namely the variable energy pricing and the power-saving states of a machine. Scheduling of the jobs and the machine states are considered jointly to achieve the highest possible savings. Although this problem has been previously addressed in the literature, the reported results of the state-of-the-art method show that the optimal solutions can be found only for instances with up to 35 jobs and 209 intervals within 3 hours of computation. We propose an elegant pre-processing technique called SPACES for computing the optimal switching of the machine states with respect to the energy costs. The optimal switchings are associated with the shortest paths in an interval-state graph that describes all possible transitions between the machine states in time. This idea allows us to implement efficient integer linear programming and constraint programming models of the problem while preserving the optimality. The efficiency of the models lies in the simplification of the optimal switching representation. The results of the experiments show that our approach outperforms the existing state-of-the-art exact method. On a set of benchmark instances with varying sizes and different state transition graphs, the proposed approach finds the optimal solutions even for the large instances with up to 190 jobs and 1277 intervals within an hour of computation.

Keywords: Single machine production scheduling · Machine states · Variable energy costs · Total energy cost minimization.

1 Introduction

Energy-efficient scheduling has been attracting a considerable amount of attention lately, as reported in both [5] and [6]. The trend is most likely to continue in the future since the energy-efficient scheduling helps to achieve sustainability

of the production by both decreasing the production cost and minimizing its environmental impact. Gahm et al. in [5] identified promising approaches to the energy-aware scheduling, including, among others, (i) the optimization of the energy demand by considering the power-saving states of the machines, and (ii) the participation in demand response programs, which are used by the electric utilities to reward the energy consumers for shifting their energy consumption to *off-peak* intervals [10].

In this work, we study a single machine scheduling problem to minimize the total energy cost (TEC) of the production. We consider both the power-saving states of the machine and the time-of-use (TOU) pricing. The TOU pricing is one of the demand response programs, in which the electricity price may differ every hour. The scheduling problems with TOU pricing have been extensively addressed in the literature [4,7,8].

Considering the power-saving states of the machine, Mouzon et al. in [11] identified that a significant energy cost reduction can be attained. However, the switchings between the machine states need to be planned carefully because of their non-negligible energy costs and transition times.

The integration of the power-saving states and the TOU pricing was originally proposed by Shrouf et al. [12], who designed an integer linear programming (ILP) model for the single machine problem with the fixed order of the jobs. However, it was proven in [2] that the problem with the fixed order of the jobs can be solved in polynomial time. Aghelinejad et al. [1] improved and generalized the existing ILP model to consider even an arbitrary order of the jobs, in which case the problem is $\mathcal{NP}$-hard [2]. However, in both [1] and [12], only small instances of the problem have been solved optimally.

In this paper, we describe a novel pre-processing technique for the single machine scheduling problem, which was introduced in [12] and further studied in [1]. Our pre-processing technique pre-computes the optimal switching behavior in time w.r.t. energy costs. The pre-computed costs of the optimal switchings allow us to design efficient exact ILP and constraint programming (CP) models. In contrast, the ILP model proposed in [1] explicitly formulates the transition behavior of the machine, which needs to be optimized jointly with the scheduling of the jobs. As shown by the experiments, our approach outperforms the existing ILP model [1], which is, to the best of our knowledge, the state-of-the-art among the exact methods for this problem. Our ILP model is able to solve all the benchmark instances with up to 190 jobs and 1277 pricing intervals within the time-limit. On the other hand, the state-of-the-art ILP model from the literature scales only up to instances with 60 jobs and 316 intervals.

2   Problem Statement

Let $\mathcal{I} = \{I_1, I_2, \ldots, I_h\}$ be a set of *intervals*, which partition the scheduling horizon. The *energy costs* for the intervals are given by the vector $\boldsymbol{c} = (c_1, c_2, \ldots, c_h)$, where $c_i \in \mathbb{Z}_{\geq 0}$ is the energy (electricity) cost associated with interval I_i. It is assumed, that every interval is one time unit long, i.e., $I_1 = [0, 1)$, $I_2 = [1, 2)$,

$\ldots, I_h = [h-1, h)$. Note that the physical representation of the time unit length can be different depending on the required granularity of the scheduling horizon.

Let $\mathcal{J} = \{J_1, J_2, \ldots, J_n\}$ be a set of *jobs*, which must be scheduled on a single machine, that is available throughout the whole scheduling horizon; we assume that $n \geq 1$. Each job J_j is characterized by its *processing time* $p_j \in \mathbb{Z}_{>0}$, given in the number of intervals. Scheduling of the jobs is non-preemptive, and the machine can process at most one job at the time. All the jobs are available at the beginning of the scheduling horizon.

During each interval, the machine is operating in one of its *states* $s \in \mathcal{S}$ or transits from one state to another. Let us denote the *transition time function* by $T : \mathcal{S} \times \mathcal{S} \to \mathbb{Z}_{\geq 0} \cup \{\infty\}$, and the *transition power function* by $P : \mathcal{S} \times \mathcal{S} \to \mathbb{Z}_{\geq 0} \cup \{\infty\}$. The transition from state s to state s' lasts $T(s, s')$ intervals and has power consumption $P(s, s')$, which is the constant rate of the consumed energy at every time unit. The value ∞ means that the direct transition does not exist. We assume that $P(s, s)$ denotes the power consumption of the machine while staying in state s for the duration of one interval.

Note that the transition time/power functions are general enough to represent many kinds of machines, e.g., those studied in [1,3,11,12].

During the first and the last interval, the machine is assumed to be in off state $\texttt{off} \in \mathcal{S}$. Besides, the machine has a single processing state, $\texttt{proc} \in \mathcal{S}$, which must be active during the processing of the jobs. Due to the transition from/to the initial/last $\texttt{off}$ state, the machine cannot be in $\texttt{proc}$ state during the early/late intervals. Hence, we denote the *earliest* and the *latest* interval during which the machine can be in $\texttt{proc}$ state by I_{earl} and I_{late}, respectively.

A *solution* is a pair $(\boldsymbol{\sigma}, \boldsymbol{\Omega})$, where $\boldsymbol{\sigma} = (\sigma_1, \sigma_2, \ldots, \sigma_n) \in \mathbb{Z}_{\geq 0}^n$ is the vector denoting the start time of the jobs, and $\boldsymbol{\Omega} = (\Omega_1, \Omega_2, \ldots, \Omega_h) \in (\mathcal{S} \times \mathcal{S})^h$ represents the active state or transition in each interval. The solution is feasible if the following four conditions are satisfied.

1. the machine processes at most one job at a time;
2. the jobs are processed when the machine is in $\texttt{proc}$ state, i.e.,
   $$\forall J_j \in \mathcal{J} \; \forall i \in \{\sigma_j + 1, \ldots, \sigma_j + p_j\} : \Omega_i = (\texttt{proc}, \texttt{proc}),$$
   where $\{a, \ldots, b\}$ is $\{a, a+1, \ldots, b\}$;
3. the machine is in $\texttt{off}$ state during the first and the last interval, i.e.,
   $\Omega_1 = (\texttt{off}, \texttt{off})$, and $\Omega_h = (\texttt{off}, \texttt{off})$;
4. all transitions are valid with respect to the transition time function.

The total energy cost (TEC) of solution $(\boldsymbol{\sigma}, \boldsymbol{\Omega})$ is

$$\sum_{I_i \in \mathcal{I}} c_i \cdot P(\Omega_i), \tag{1}$$

where $P(\Omega_i)$ represents $P(s, s')$ for $\Omega_i = (s, s')$. The goal of the scheduling problem is to find a feasible solution minimizing the total energy cost (1).

The above-defined problem was introduced in [12] and is denoted in standard Graham's notation as 1, TOU|states|TEC. The problem was shown to be strongly $\mathcal{NP}$-hard, see [2].

Example: Here, we present a small example to illustrate the proposed notation. Let us consider a scheduling horizon consisting of 16 intervals, $\mathcal{I} = \{I_1, \ldots, I_{16}\}$, and the associated energy costs $\boldsymbol{c} = (2, 1, 2, 1, 8, 16, 14, 3, 2, 5, 3, 10, 3, 2, 1, 2)$. Let us have three jobs, $\mathcal{J} = \{J_1, J_2, J_3\}$ with processing times $p_1 = 2$, $p_2 = 1$, and $p_3 = 2$. Considering the machine states, we assume $\mathcal{S} = \{\texttt{proc}, \texttt{off}, \texttt{idle}\}$. The values of the transition time function and the transition power function are given in Fig. 1. For the given transition time function, we have $I_{\mathrm{earl}} = I_4$ and $I_{\mathrm{late}} = I_{14}$. Note that the same machine states and transitions were originally proposed in [12].


Fig. 1. Parameters of the transition power function $P(s, s')$ and transition time function $T(s, s')$, and the corresponding transition graph, where every edge from s to s' is labeled by $T(s, s')/P(s, s')$.

The optimal solution to the given instance is depicted in Fig. 2, where

$$\boldsymbol{\sigma} = (9, 3, 12), \text{ and}$$

$$\begin{aligned}
\boldsymbol{\Omega} = (&(\texttt{off}, \texttt{off}), (\texttt{off}, \texttt{proc}), (\texttt{off}, \texttt{proc}), (\texttt{proc}, \texttt{proc}), (\texttt{proc}, \texttt{off}), (\texttt{off}, \texttt{off}), \\
&(\texttt{off}, \texttt{off}), (\texttt{off}, \texttt{proc}), (\texttt{off}, \texttt{proc}), (\texttt{proc}, \texttt{proc}), (\texttt{proc}, \texttt{proc}), (\texttt{idle}, \texttt{idle}), \\
&(\texttt{proc}, \texttt{proc}), (\texttt{proc}, \texttt{proc}), (\texttt{proc}, \texttt{off}), (\texttt{off}, \texttt{off})).
\end{aligned}$$

The TEC of the optimal solution is equal to 177.

3   Solution Approach

In this section, we first describe how to pre-compute the optimal switching behavior of the machine and the corresponding costs. Afterward, we design efficient ILP and CP models (called ILP-SPACES and CP-SPACES) that integrate the pre-computed optimal switching costs.

3.1   Instance Pre-processing: Computation of the Optimal Switching

Given two states s, s' in which the machine is during two intervals $I_i, I_{i'}$ such that $i < i'$, the pre-processing computes the optimal transitions from (s, I_i) to

Interval I_i	I_1	I_2	I_3	I_4	I_5	I_6	I_7	I_8	I_9	I_{10}	I_{11}	I_{12}	I_{13}	I_{14}	I_{15}	I_{16}
Energy cost c_i	2	1	2	1	8	16	14	3	2	5	3	10	3	2	1	2
Processing				6						30	18		18	12		
Idle												20				
Off	0					0	0									0
Turn-off ($\searrow$)					8										1	
Turn-on ($\nearrow$)		8	16					24	16							
Schedule	off	$\nearrow$	J_2	$\searrow$	off		$\nearrow$		J_1		idle	J_3			$\searrow$	off

Fig. 2. The optimal schedule for the example instance. Each cell, corresponding to interval I_i and a state/transition, contains the value $c_i \cdot P(\Omega_i)$. The sum over all these values gives the TEC equal to 177.

$(s', I_{i'})$ over all possible states w.r.t. the energy cost. Formally, the pre-processing solves the following optimization problem

$$\min_{\Omega_{i+1}, \Omega_{i+2}, \ldots, \Omega_{i'-1}} \sum_{i''=i+1}^{i'-1} c_{i''} \cdot P(\Omega_{i''}). \tag{2}$$

such that $((s,s), \Omega_{i+1}, \Omega_{i+2}, \ldots, \Omega_{i'-1}, (s', s'))$ are valid transitions w.r.t. to the transition time function. We call this an *optimal switching problem*. As an illustration, the cost of the optimal switching in Fig. 2 from $(\texttt{proc}, I_4)$ to $(\texttt{proc}, I_{10})$ equals 48. Interestingly, the optimal switching problem can be solved in polynomial time by finding the shortest path in an *interval-state* graph, which is explained in the rest of this section.

The interval-state graph is defined by a triplet (V, E, w), where V is the set of *vertices*, E is the set of *edges* and $w : E \to \mathbb{Z}_{\geq 0}$ are the *weights* of the edges. The set of the vertices and edges of this graph are defined as follows:

$$V = \{v_{1,\texttt{off}}\} \cup \{v_{i,s} : I_i \in \mathcal{I} \setminus \{I_1\}, s \in \mathcal{S}\} \cup \{v_{h+1,\texttt{off}}\}, \tag{3}$$

$$\begin{aligned} E = &\{(v_{1,\texttt{off}}, v_{2,\texttt{off}})\} \\ &\cup \{(v_{i,s}, v_{i+T(s,s'),s'}) : s, s' \in \mathcal{S}, I_i \in \mathcal{I} \setminus \{I_1\}, \\ &\qquad\qquad T(s,s') \neq \infty, (i-1) + T(s,s') \leq h-1\} \\ &\cup \{(v_{h,\texttt{off}}, v_{h+1,\texttt{off}})\}. \end{aligned} \tag{4}$$

Informally, each vertex $v_{i,s} \in V$ represents that at the beginning of interval I_i the machine is in state s. Each edge $(v_{i,s}, v_{i',s'}) \in E$ corresponds to the direct transition from state s to state s' that lasts $T(s,s') = (i' - i)$ intervals. The condition $(i-1) + T(s,s') \leq h-1$ ensures, that only transitions completing at most at the beginning of interval I_h are present in the interval-state graph.

The edges are weighted by the total energy cost of the corresponding transition w.r.t. the costs of energy in intervals, i.e., *weight* of edge $(v_{i,s}, v_{i',s'}) \in E$ is

defined as

$$w(v_{i,s}, v_{i',s'}) = \sum_{i''=i}^{i'-1} c_{i''} \cdot P(s, s') \,. \tag{5}$$

Returning to the optimal switching problem (2), the optimal transitions from (s, I_i) to $(s', I_{i'})$ w.r.t. the energy cost can be obtained by finding the shortest path from $v_{i+1,s}$ to $v_{i',s'}$ in the interval-state graph. We denote the cost of the optimal switching by function $l : V \times V \to \mathbb{Z}_{\geq 0}$. The values of l can be computed using the Floyd-Warshall algorithm in $\mathcal{O}(h^3 \cdot |\mathcal{S}|^3)$ time.

However, for the scheduling decisions, only some of the switchings are interesting. Since all the jobs need to be scheduled in the proc state of the machine, the optimal switchings need to be resolved only in the 'space', i.e., the sequence of intervals: (i) between two consecutive intervals with proc; (ii) between the first off and the first proc; and (iii) the last proc and the last off. The cost of the switchings between $s, s' \in \{\text{off}, \text{proc}\}^2$ are recorded by function $c^\star : \mathcal{I}^2 \to \mathbb{Z}_{\geq 0}$ defined as

$$c^\star(i, i') = \begin{cases} l(v_{i+1,\text{proc}}, v_{i',\text{proc}}) & i > 1, i' < h \qquad \text{case (i)} \\ l(v_{2,\text{off}}, v_{i',\text{proc}}) & i = 1, i' < h \qquad \text{case (ii)} \\ l(v_{i+1,\text{proc}}, v_{h,\text{off}}) & i > 1, i' = h \qquad \text{case (iii)} \end{cases} \tag{6}$$

for each $i < i'$. The vector of states corresponding to $c^\star(i, i')$, i.e., the *optimal switching behavior* of the machine between i and i', is denoted by $\boldsymbol{\Omega}^\star(i, i')$. As an example, see the Fig. 2, where intervals $\{I_5, I_6, \ldots, I_9\}$ represent the space between two consecutive jobs J_2, J_1 with cost $c^\star(4, 10) = 48$.

Values of $c^\star$ can be computed efficiently using an algorithm that we call the *Shortest Path Algorithm for Cost Efficient Switchings* (SPACES). In every iteration $I_i \in \mathcal{I} \setminus \{I_h\}$, SPACES computes all values $c^\star(i, i+1), c^\star(i, i+2), \ldots, c^\star(i, h)$ by finding the shortest paths from $v_{i+1,\text{proc}}$ (or $v_{2,\text{off}}$ if $i = 1$) to all other vertices in the interval-state graph. The shortest paths are obtained with Dijkstra algorithm that runs in $\mathcal{O}(|E| + |V| \cdot \log |V|)$ if implemented using the priority queues. Since the Dijkstra algorithm is started h times, the complexity of SPACES is

$$\mathcal{O}(h \cdot (|E| + |V| \cdot \log |V|)) = \mathcal{O}(h^2 \cdot |\mathcal{S}| \cdot (|\mathcal{S}| + \log h \cdot |\mathcal{S}|)) \,. \tag{7}$$

Moreover, to increase the performance further, iterations i can be computed in parallel since they are independent to each other.

Example (continued): Continuing the Example, Fig. 3 shows the whole interval-state graph for the given instance. The green dashed path shows the optimal switching behavior of the machine assuming that the machine is in proc state during intervals I_4 and I_{10}; at first the machine is turned off (during I_5), then it remains off (during intervals I_6 and I_7), and is turned on afterward (intervals I_8, I_9). The optimal switching cost $c^\star(4, 10)$ is, in this case, 48. The optimal switching behavior is

$$\boldsymbol{\Omega}^\star(4, 10) = ((\text{proc}, \text{off}), (\text{off}, \text{off}), (\text{off}, \text{off}), (\text{off}, \text{proc}), (\text{off}, \text{proc})) \,. \tag{8}$$
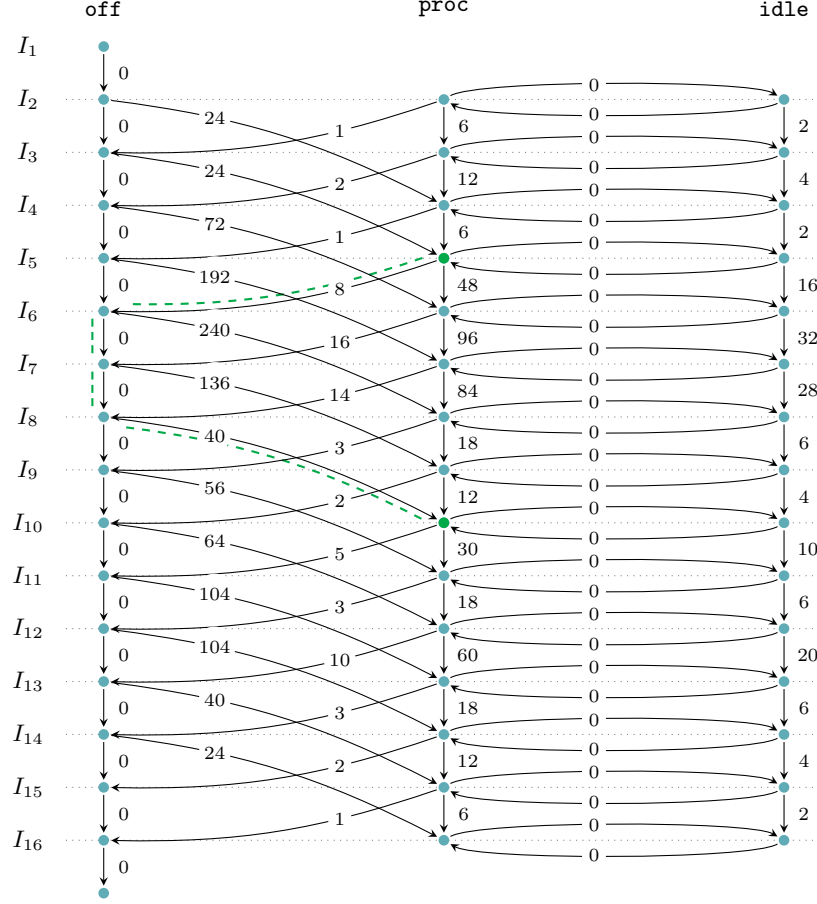
Fig. 3. Interval-state graph for the Example instance from Section 2 with highlighted optimal switching behavior from $(\mathtt{proc}, I_4)$ to $(\mathtt{proc}, I_{10})$.

3.2   Integer Linear Programming Model ILP-SPACES

In the ILP model proposed in [1], the state transition functions are explicitly encoded. In contrast, our ILP-SPACES model works only with the optimal switching costs pre-computed by the SPACES algorithm, thus encoding the transitions implicitly without sacrificing the optimality. The only task of the ILP solver is then to schedule the jobs, and select appropriate spaces in between, such that the TEC is minimized. Thus, the structure of our model is greatly simplified, with positive impact on its performance.

Formally, the variables used in the ILP-SPACES model are

- job start time $s_{j,i} \in \{0, 1\}$: equals 1 if job J_j starts at the beginning of interval I_i, otherwise 0;

– space activation $x_{i,i'} \in \{0,1\}$: equals 1 if the machine undergoes the optimal switching defined by $\boldsymbol{\Omega}^\star(i,i')$, otherwise 0.

The complete model follows.

$$\min \sum_{\substack{I_i,I_{i'} \in \mathcal{I} \\ i<i'}} x_{i,i'} \cdot c^\star(i,i') + \sum_{\substack{J_j \in \mathcal{J} \\ I_i \in \mathcal{I}}} s_{j,i} \cdot c_{j,i}^{(\mathrm{job})}, \tag{9}$$

$$\sum_{I_i \in \mathcal{I}} s_{j,i} = 1, \ \forall J_j \in \mathcal{J}, \tag{10}$$

$$s_{j,i} = 0, \ \forall J_j \in \mathcal{J}, \forall i \in \{1,\dots,\mathrm{earl}-1\} \cup \{\mathrm{late}-p_j+2,\dots,h\}, \tag{11}$$

$$\sum_{J_j \in \mathcal{J}} \sum_{i'=\max\{2,i-p_j+1\}}^{i} s_{j,i'} + \sum_{i'=1}^{i-1} \sum_{i''=i+1}^{h} x_{i',i''} = 1, \ \forall I_i \in \{I_2, I_3, \dots, I_{h-1}\}. \tag{12}$$

The objective (9) minimizes the total energy cost, consisting of the optimal switching cost of active spaces, and the cost of jobs processing, where

$$c_{j,i}^{(\mathrm{job})} = \sum_{i'=i}^{i+p_j-1} c_{i'} \cdot P(\texttt{proc},\texttt{proc}) \tag{13}$$

for job $J_j \in \mathcal{J}$ and $i \in \{\mathrm{earl},\dots,\mathrm{late}-p_j+1\}$.

Constraint (10) forces every job to be scheduled exactly once, and constraint (11) forbids the job to be scheduled before I_{earl} and after I_{late}. Finally, the last constraints (12) force the machine to be processing a job or to be undergoing some transition during every interval and forbid overlaps between them.

Search Space Reduction Various methods can be employed to reduce the search space without sacrificing the optimality. One of such methods is pruning of the spaces variables that lead to infeasible solutions if activated.

The pruning works as follows. For each $I_i, I_{i'}$ such that $i < i'$, the available time for processing the jobs is computed for both left (before I_i) and right (after $I_{i'}$) part of the scheduling horizon, i.e., $i - \mathrm{earl} + 1$ and $\mathrm{late} - i' + 1$, respectively. Then, activating the switching behavior $\boldsymbol{\Omega}^\star(i,i')$ leads to an infeasible solution if one of the following *pruning conditions* holds

PC.1: The largest job can be fitted in neither part, i.e.,

$$\max_{J_j \in \mathcal{J}} p_j > i - \mathrm{earl} + 1 \quad \wedge \quad \max_{J_j \in \mathcal{J}} p_j > \mathrm{late} - i' + 1. \tag{14}$$

PC.2: The total available time for processing is less than the sum of all the processing times, i.e.,

$$(i - \mathrm{earl} + 1) + (\mathrm{late} - i' + 1) < \sum_{J_j \in \mathcal{J}} p_j. \tag{15}$$

If any of these conditions holds, the corresponding space variable $x_{i,i'}$ is not created in ILP-SPACES.

3.3 Constraint Programming Model CP-SPACES

The idea of the CP-SPACES model is similar to the ILP-SPACES, with the exception that the spaces are not fixed – they are allowed to 'float' within the scheduling horizon. In consequence, the spaces do not have fixed costs because the cost depends on the position of the space in the horizon and its length. In our CP-SPACES model, costs are formulated with an *Element* expression, which is integrated in the objective. To describe the CP-SPACES model, we use the IBM CP formalism [9].

Variables: Three types of interval variables are used in the CP model.

To represent the jobs, we use optional interval variables $s_{j,i}$, $\forall J_j \in \mathcal{J}, i \in \{\text{earl}, \ldots, \text{late} - p_j + 1\}$, which model whether job J_j starts at the beginning of the interval I_i. Only one such variable, represented by interval variable s_j, will be present in the schedule for each job. Length of $s_{j,i}$ is fixed to p_j, and its start is fixed to $i - 1$.

Finally, the optional interval variables $x_{\ell,k}$ represent the 'floating' spaces of fixed length. For each possible *length* $\ell \in \{1, 2, \ldots, h - 2 - \sum_{J_j \in \mathcal{J}} p_j\}$, we create $K(\ell) = \left\lfloor \frac{h - 2 - \sum_{J_j \in \mathcal{J}} p_j}{\ell} \right\rfloor$ variables that are indexed by $k \in \{1, 2 \ldots, K(\ell)\}$. Note that the number $K(\ell)$ gives the upper bound on the number of the spaces of length ℓ that may appear in a feasible schedule, while $\ell_{\max} = h - 2 - \sum_{J_j \in \mathcal{J}} p_j$ gives an upper bound on the space length.

Constraints: Since the machine is assumed to be in `off` state during I_1 and I_h, the earliest and the latest interval during which a switching might occur is I_2 and I_{h-1}, respectively. Hence, starts (ends) of the spaces are restricted by

$$\left.\begin{array}{l} \text{StartOf}(x_{\ell,k}) \geq 1 \\ \text{EndOf}(x_{\ell,k}) \leq h - 1 \end{array}\right\} \ \forall \ell \in \{1, \ldots, \ell_{\max}\}, k \in \{1, \ldots, K(\ell)\}. \qquad (16)$$

As mentioned previously, the spaces have fixed lengths, i.e.,

$$\text{LengthOf}(x_{\ell,k}) = \ell, \ \forall \ell \in \{1, \ldots, \ell_{\max}\}, k \in \{1, \ldots, K(\ell)\}. \qquad (17)$$

Relationship between variables s_j and $s_{j,i}$ is given by

$$\text{Alternative}(s_j, \{s_{j,i} : i \in \{\text{earl}, \ldots, \text{late} - p_j + 1\}\}), \ \forall J_j \in \mathcal{J}, \qquad (18)$$

To ensure that jobs and spaces are not overlapping, we use the *NoOverlap* constraint,

$$\text{NoOverlap}(\{x_{\ell,k} : \ell \in \{1, \ldots, \ell_{\max}\}, k \in \{1, \ldots, K(\ell)\}\} \cup \{s_j : J_j \in \mathcal{J}\}). \qquad (19)$$

The lengths of the spaces are constrained by

$$\sum_{\ell=1}^{\ell_{\max}} \sum_{k=1}^{K(\ell)} \text{LengthOf}(x_{\ell,k}) = \ell_{\max}, \qquad (20)$$

to ensure that the whole scheduling horizon is filled.

Finally, to eliminate some symmetries in the model, space variables $x_{\ell,k}$ are constrained such that space with index k can be present in the solution only if all the spaces $x_{\ell,k'}$ of the same length with $k' < k$ are present, i.e.,

$$\text{PresenceOf}(x_{\ell,k}) \geq \text{PresenceOf}(x_{\ell,k+1}), \ \forall \ell \in \{1, \ldots, \ell_{\max}\}, k \in \{1, \ldots, K(\ell)-1\}. \tag{21}$$

Objective: The objective is to minimize the TEC, here expressed as

$$\sum_{\ell=1}^{\ell_{\max}} \sum_{k=1}^{K(\ell)} \text{Element}(\boldsymbol{c}_\ell^\star, \text{StartOf}(x_{\ell,k})) + \sum_{J_j \in \mathcal{J}} \sum_{i=\text{earl}}^{\text{late}-p_j+1} \text{PresenceOf}(s_{j,i}) \cdot c_{j,i}^{(\text{job})}, \tag{22}$$

where the first part corresponds to the cost for optimal switchings between the job processings, and the second part corresponds to the cost for job processing. To compute the cost of the present spaces, vector

$$\boldsymbol{c}_\ell^\star = (c^\star(1, 1+\ell+1), c^\star(2, 2+\ell+1), \ldots, c^\star(h-\ell-1, h)) \tag{23}$$

is used to represent the optimal switching costs for the given ℓ addressed by the start of space $x_{\ell,k}$ (indexed from 1).

Remark 1. Note that there are many different ways of implementing the CP model. For example, each possible space might be fixed in time, similarly as in ILP-SPACES. Then, the objective would simplify to a sum of spaces presences multiplied by their costs. Another alternative would be to have 'floating' spaces with variable lengths. In that case, the number of the interval variables needed to represent the spaces would decrease to $n-1$, but the element expression in the objective would need to be indexed by both the start time and the length of each space. Also, the element expressions could be replaced by overlap expressions, etc. We have tried multiple different alternatives; however, the performances of the models on preliminary benchmark instances were more or less similar. The described CP model was slightly better than the others; therefore, we use it for the experiments in Section 4.

Remark 2. We can use a similar idea to reduce the search space as applied to ILP-SPACES. However, in the case of CP-SPACES, the spaces variables cannot be pruned since they are not fixed in time. Instead, we can enforce the corresponding cost of this space to a large number, which effectively deactivates it.

4   Experiments

This section evaluates how ILP-SPACES and CP-SPACES models perform in comparison to the ILP-REF model proposed in [1]. The comparison is made on

a set of randomly generated instances with varying sizes and different machine transition graphs; see Section 4.1 for the description of the generated dataset. The results are presented in Section 4.2.

All experiments were executed on 2x Intel(R) Xeon(R) Silver 4110 CPU 2.10 GHz with 188 GB of RAM (16 cores in total). For solving the ILP and CP models, we used Gurobi 8 and IBM CP Optimizer 12.9, respectively. Except for the time-limit and search phases in CP-SPACES, which branched on jobs first, all the solver parameters were set to default values.

The generated instances and their solutions are publicly available at
https://github.com/CTU-IIG/EnergyStatesAndCostsSchedulingDatasets.

4.1 Instances

The instances in the dataset can be divided according to

1. a number of jobs:
    (a) MEDIUM: medium instances with $n \in \{30, 60, 90\}$;
    (b) LARGE: large instances with $n \in \{150, 170, 190\}$;
2. a machine transition graph:
    (a) NOSBY: a simple model with no standby state used by [12,1], see Fig. 1 for its description;
    (b) TWOSBY: a model with two standby states shown in Fig. 4.

For fixed n and a machine transition graph, 12 random instances are generated in the following way (48 instances in the whole dataset). The processing times of the jobs are randomly sampled from discrete uniform distribution $\mathcal{U}\{1,5\}$. The number of intervals in each instance is obtained as a multiple of the total processing time plus the required number of intervals to turn the machine on and off, where the multiple is taken from set the $\{1.3, 1.6, 1.9, 2.2\}$. The energy cost in each interval is randomly sampled from $\mathcal{U}\{1, 10\}$. For instances differing only in the number of intervals, the energy costs are sampled gradually, i.e., the energy costs of all the intervals in an instance with a shorter horizon are the same as for the corresponding intervals in an instance with a longer horizon.

Note that the distributions for sampling the processing times and the energy costs of the intervals are the same as proposed in [12,1].

4.2 Results

All the presented Tables 1, 2, 3, and 4 have the same structure: each row represents one instance characterized by the number of the jobs n and the number of intervals h. The objective value ub of the found feasible solution, lower bound lb and the running time t are given for each tested model. If the objective value or the lower bound is in bold font, the corresponding value is known to be optimal. Therefore, if both objective and the lower bound are in bold, the solver was able to prove the solution optimality within the time-limit. If the solver reached its given time-limit on an instance without proving the optimality of a solution, the value in the corresponding cell in t column is TLR.

Fig. 4. Example of a transition graph with multiple standby states; every edge from s to s' is labeled by $T(s, s')/P(s, s')$.

Additionally, we report the pre-processing time P-P for the large instances. For medium-size instances, the pre-processing time is negligible with the average time 0.69 s and maximum time 2.93 s.

The last rows in each table shows the average running time on each model and the average optimality gap. The average time is computed over all instances; if the solver time-outed on some instance, the specified time-limit is taken as the running time on that instance. The optimality gap on each instance is defined as

$$\frac{ub - lb^{\text{best}}}{ub} \cdot 100 \ [\%] \,, \tag{24}$$

where lb^{best} is the best lower bound obtained over all models on that instance.

Results for Medium Instances The results of the experiment with medium-size instances for NOSBY and TWOSBY transition graphs are shown in Table 1 and Table 2, respectively. In these tables we can see that ILP-SPACES finds the optimal solutions and proves their optimality for all instances. On the other hand, the model ILP-REF proposed in [1] finds the optimal solution and proves the optimality only for 11 instances out of 24 within the time-limit (600 s). Moreover, some of the non-optimal solutions found by ILP-REF are far from the optimum, for example, the objective of the solution found for $n = 90, h = 621$ on TWOSBY is more than twice the objective of the optimal one found by ILP-SPACES.

Unfortunately, CP-SPACES is not able to prove the optimality of any instance within the time-limit. However, the average optimality gaps (3.40 % for NOSBY and 1.79 % for TWOSBY) reveals that it can find near-optimal solutions. The performance of both CP-SPACES and ILP-SPACES is slightly influenced by a more complex transition graph, whereas the performance of ILP-REF deteriorates significantly (average optimality gap 2.14 % for NOSBY increased to 29.59 % for TWOSBY).

Table 1. MEDIUM+NOSBY: Comparison of found upper bound ub, lower bound lb and runtime t between the models. Time-limit is 600 s and TLR stands for time-limit reached.

Instance		ILP-REF [1]			CP-SPACES			ILP-SPACES		
n	h	ub [-]	lb [-]	t [s]	ub [-]	lb [-]	t [s]	ub [-]	lb [-]	t [s]
30	104	**1 426**	**1 426**	3.7	1 448	496	TLR	**1 426**	**1 426**	1.3
30	127	**1 394**	**1 394**	4.9	1 412	488	TLR	**1 394**	**1 394**	1.7
30	150	**1 394**	**1 394**	5.7	1 414	484	TLR	**1 394**	**1 394**	2.4
30	173	**1 394**	**1 394**	7.0	1 458	484	TLR	**1 394**	**1 394**	3.1
60	258	**4 290**	**4 290**	88.5	4 379	1 724	TLR	**4 290**	**4 290**	7.7
60	316	**3 994**	**3 994**	344.7	4 117	1 584	TLR	**3 994**	**3 994**	29.0
60	374	**3 836**	3 826	TLR	3 952	1 424	TLR	**3 836**	**3 836**	29.0
60	432	3 956	3 800	TLR	3 972	1 380	TLR	**3 833**	**3 833**	46.9
90	363	6 044	5 839	TLR	6 004	2 328	TLR	**5 920**	**5 920**	7.0
90	445	5 778	5 567	TLR	5 760	2 232	TLR	**5 686**	**5 686**	166.0
90	528	5 916	4 695	TLR	5 916	2 168	TLR	**5 431**	**5 431**	64.5
90	610	5 901	4 514	TLR	5 829	1 828	TLR	**5 373**	**5 373**	147.1
Average time [s]:				337.9			>600			42.1
Average optimality gap [%]:				2.14			3.40			0.00

Table 2. MEDIUM+TWOSBY: Comparison of found upper bound ub, lower bound lb and runtime t between the models. Time-limit is 600 s and TLR stands for time-limit reached.

Instance		ILP-REF [1]			CP-SPACES			ILP-SPACES		
n	h	ub [-]	lb [-]	t [s]	ub [-]	lb [-]	t [s]	ub [-]	lb [-]	t [s]
30	106	**3 815**	**3 815**	29.4	**3 815**	1 240	TLR	**3 815**	**3 815**	1.4
30	129	**3 804**	**3 804**	30.7	3 815	1 220	TLR	**3 804**	**3 804**	2.3
30	152	**3 804**	**3 804**	42.0	3 815	1 210	TLR	**3 804**	**3 804**	7.0
30	175	**3 804**	**3 804**	61.4	3 815	1 210	TLR	**3 804**	**3 804**	9.5
60	254	**10 863**	**10 863**	588.1	11 058	4 190	TLR	**10 863**	**10 863**	2.0
60	311	10 289	10 087	TLR	10 574	3 860	TLR	**10 248**	**10 248**	43.3
60	368	**9 917**	9 696	TLR	10 163	3 470	TLR	**9 917**	**9 917**	82.1
60	426	20 346	9 133	TLR	10 055	3 340	TLR	**9 874**	**9 874**	233.9
90	370	17 179	14 818	TLR	15 470	5 900	TLR	**15 379**	**15 379**	140.2
90	454	22 808	12 951	TLR	15 156	5 680	TLR	**14 923**	**14 923**	138.6
90	538	25 992	11 868	TLR	15 107	5 500	TLR	**14 548**	**14 548**	403.8
90	621	29 558	11 406	TLR	15 152	4 620	TLR	**14 392**	**14 392**	225.8
Average time [s]:				412.6			>600			107.5
Average optimality gap [%]:				29.59			1.79			0.00

Results for Large Instances The results of the experiment with large instances for NOSBY and TWOSBY transition graphs are shown in Table 3 and Table 4, respectively. The results for CP-SPACES are not included, since we were unable to obtain solutions to all the instances from the IBM CP Optimizer. We observed that the solver used all the available RAM and swap memory (188 GB + 191 GB), which indicates internal issues of the solver. However, for three instances where the CP solver was able to find the solutions, the objective was better than for ILP-REF.

Looking at the results of ILP-SPACES, we can see that it found the optimal solutions and proved their optimality for all 24 instances. On the other hand, ILP-REF was able to find the optimal solutions for only two smallest instances, but was not able to prove their optimality within the specified time-limit (3 600 s). Comparing the average optimality gaps, ILP-REF achieved 8.39 % on NOSBY transition graph and 61.33 % on TWOSBY, whereas ILP-SPACES achieved 0 % optimality gap on both transition graphs. This shows that ILP-SPACES scales even to large instances.

Table 3. LARGE+NOSBY: Comparison of found upper bound ub, lower bound lb and runtime t between the models. Time-limit is 3 600 s and TLR stands for time-limit reached.

Instance		ILP-REF [1]			ILP-SPACES			P-P
n	h	ub [-]	lb [-]	t [s]	ub [-]	lb [-]	t [s]	t [s]
150	527	**8 582**	8 567	TLR	**8 582**	**8 582**	187	1.0
150	647	8 726	8 240	TLR	**8 409**	**8 409**	277	2.9
150	767	8 557	7 787	TLR	**8 132**	**8 132**	624	5.5
150	888	8 976	6 780	TLR	**8 078**	**8 078**	511	9.1
170	650	10 596	9 628	TLR	**10 068**	**10 068**	290	2.3
170	799	10 794	8 832	TLR	**9 820**	**9 820**	1 087	4.6
170	948	10 940	8 343	TLR	**9 637**	**9 637**	806	9.3
170	1 097	11 189	8 124	TLR	**9 620**	**9 620**	1 345	13.4
190	757	12 555	11 206	TLR	**12 008**	**12 008**	246	3.9
190	930	12 882	10 521	TLR	**11 758**	**11 758**	942	6.9
190	1 104	12 791	9 949	TLR	**11 611**	**11 611**	3 147	13.3
190	1 277	12 757	0	TLR	**11 465**	**11 465**	1 348	22.7
Average time [s]:				>3600			901	7.9
Average optimality gap [%]:				8.39			0.00	

Table 4. LARGE+TWOSBY: Comparison of found upper bound *ub*, lower bound *lb* and runtime *t* between the models. Time-limit is $3\,600\,$s and TLR stands for time-limit reached.

Instance		ILP-REF [1]			ILP-SPACES			P-P
n	h	ub [-]	lb [-]	t [s]	ub [-]	lb [-]	t [s]	t [s]
150	529	**21 910**	21 562	TLR	**21 910**	**21 910**	130	1.1
150	649	29 425	20 685	TLR	**21 821**	**21 821**	702	3.1
150	769	37 764	18 140	TLR	**21 353**	**21 353**	949	5.2
150	890	43 929	16 799	TLR	**21 266**	**21 266**	701	8.5
170	651	28 425	24 983	TLR	**25 807**	**25 807**	809	2.6
170	799	39 095	21 981	TLR	**25 518**	**25 518**	1 244	5.0
170	948	46 083	20 709	TLR	**25 279**	**25 279**	2 922	8.5
170	1 096	53 177	20 091	TLR	**25 279**	**25 279**	2 162	14.1
190	756	38 471	27 984	TLR	**30 563**	**30 563**	797	4.2
190	929	46 319	26 166	TLR	**30 224**	**30 224**	1 069	7.5
190	1 102	53 751	24 630	TLR	**30 224**	**30 224**	2 069	13.6
190	1 275	61 547	0	TLR	**30 071**	**30 071**	2 572	23.7
Average time [s]:				>3600			1344	8.1
Average optimality gap [%]:				61.33			0.00	

5 Conclusions

Continuing on the recent research of the single-machine scheduling problem with the variable energy costs and power-saving machine states, we propose a pre-processing algorithm SPACES, which pre-computes the optimal switching behavior of the machine for all possible spaces in the schedule. The pre-processing can be done in polynomial time and works well even for large instances of the problem, e.g., it takes 23 s to pre-process our largest benchmark instance with 190 jobs and 1277 intervals. The pre-computed switching costs are successfully integrated into novel CP and ILP models, which are compared to the state-of-the-art exact ILP model on a set of benchmark instances. Results show that our approach outperforms the existing methods considering all aspects – the runtime, the provided lower bounds and the upper bounds. Using our models, we obtain the optimal solutions even for the large instances with up to 190 jobs and 1277 intervals, which have been previously tackled only heuristically [1].

References

1. Aghelinejad, M., Ouazene, Y., Yalaoui, A.: Production scheduling optimisation with machine state and time-dependent energy costs. International Journal of Production Research **56**(16), 5558–5575 (2018). https://doi.org/10.1080/00207543.2017.1414969
2. Aghelinejad, M., Ouazene, Y., Yalaoui, A.: Complexity analysis of energy-efficient single machine scheduling problems. Operations Research Perspectives **6**, 100105 (2019). https://doi.org/10.1016/j.orp.2019.100105
3. Benedikt, O., Šůcha, P., Módos, I., Vlk, M., Hanzálek, Z.: Energy-aware production scheduling with power-saving modes. In: van Hoeve, W.J. (ed.) Integration of Constraint Programming, Artificial Intelligence, and Operations Research. pp. 72–81. Springer International Publishing, Cham (2018)
4. Fang, K., Uhan, N.A., Zhao, F., Sutherland, J.W.: Scheduling on a single machine under time-of-use electricity tariffs. Annals of Operations Research **238**(1), 199–227 (Mar 2016)
5. Gahm, C., Denz, F., Dirr, M., Tuma, A.: Energy-efficient scheduling in manufacturing companies: A review and research framework. European Journal of Operational Research **248**(3), 744 – 757 (2016). https://doi.org/10.1016/j.ejor.2015.07.017
6. Gao, K., Huang, Y., Sadollah, A., Wang, L.: A review of energy-efficient scheduling in intelligent production systems. Complex & Intelligent Systems (2019)
7. Gong, X., Pessemier, T.D., Martens, L., Joseph, W.: Energy- and labor-aware flexible job shop scheduling under dynamic electricity pricing: A many-objective optimization investigation. Journal of Cleaner Production **209**, 1078 – 1094 (2019). https://doi.org/10.1016/j.jclepro.2018.10.289
8. Hadera, H., Harjunkoski, I., Sand, G., Grossmann, I.E., Engell, S.: Optimization of steel production scheduling with complex time-sensitive electricity cost. Computers & Chemical Engineering **76**, 117 – 136 (2015). https://doi.org/10.1016/j.compchemeng.2015.02.004
9. Laborie, P., Rogerie, J., Shaw, P., Vilím, P.: IBM ILOG CP optimizer for scheduling. Constraints **23**(2), 210–250 (Apr 2018). https://doi.org/10.1007/s10601-018-9281-x
10. Merkert, L., Harjunkoski, I., Isaksson, A., Säynevirta, S., Saarela, A., Sand, G.: Scheduling and energy – industrial challenges and opportunities. Computers & Chemical Engineering **72**, 183 – 198 (2015). https://doi.org/10.1016/j.compchemeng.2014.05.024, a Tribute to Ignacio E. Grossmann
11. Mouzon, G., Yildirim, M.B., Twomey, J.: Operational methods for minimization of energy consumption of manufacturing equipment. International Journal of Production Research **45**(18–19), 4247–4271 (2007). https://doi.org/10.1080/00207540701450013
12. Shrouf, F., Ordieres-Meré, J., García-Sánchez, A., Ortega-Mier, M.: Optimizing the production scheduling of a single machine to minimize total energy consumption costs. Journal of Cleaner Production **67**, 197 – 207 (2014). https://doi.org/10.1016/j.jclepro.2013.12.024