

A Multivariate Complexity Analysis of the Material Consumption Scheduling Problem

#7381

Planning, Routing, and Scheduling: Scheduling

Abstract

The NP-hard MATERIAL CONSUMPTION SCHEDULING PROBLEM and closely related problems have been thoroughly studied since the 1980's. Roughly speaking, the problem deals with minimizing the makespan when scheduling jobs that consume non-renewable resources. We focus on the single-machine case without preemption: from time to time, the resources of the machine are (partially) replenished, thus allowing for meeting a necessary pre-condition for processing further jobs, each of which having individual resource demands. We initiate a systematic exploration of the parameterized (exact) complexity landscape of the problem, providing parameterized tractability as well as intractability results. Doing so, we mainly investigate how parameters related to the resource supplies influence the computational solvability. Thereby, we get a deepened understanding of the algorithmic complexity of this fundamental scheduling problem.

Keywords. Scheduling with non-renewable resources, Parameterized algorithmics and complexity, NP-hard problems, fine-grained complexity, exact algorithms, Integer Linear Programming.

Introduction

Consider the following motivating example. Every day, an agent works for a number of clients, all of equal importance. The clients, one-to-one corresponding to jobs, each time request a service with individual processing time and individual consumption of a non-renewable resource; examples for such resources include raw material, energy, and money. Since the agent wants to end the working day as soon as possible, the goal is to finish all jobs earliest possible, known as minimizing the makespan in the scheduling literature. Unfortunately, the agent only has a limited initial supply of the resource which is to be renewed (with potentially different amounts) at known points of time during the day. Since the job characteristics (resource consumption, job length) and the resource delivery characteristics (delivery amount, point of time) are known in advance, the objective thus is to find a feasible job schedule minimizing the makespan. Notably, jobs cannot be preempted and only one at a time can be executed. In Figure 1 we provide a concrete numerical example

with five jobs with varying job lengths and resource requirements.

In the scheduling literature, the described problem setting is known as minimizing the makespan on a single machine with non-renewable resources. Notably, in our example we considered the special but perhaps most prominent case of just one type of resource. More specifically, we study the single-machine variant of the NP-hard MATERIAL CONSUMPTION SCHEDULING PROBLEM. Formally, we have a set \mathcal{R} of resources and a set $\mathcal{J} = \{J_1, \dots, J_n\}$ of jobs to be scheduled on a single machine without preemption. The machine can process at most one job at a time. Each job has a processing time $p_j \in \mathbb{Z}_+$ and a resource requirement $a_{ij} \in \mathbb{Z}_+$ from resource $i \in \mathcal{R}$. We have resource supplies at q different points of time $0 = u_1 < u_2 < \dots < u_q$, where the vector $\tilde{b}_\ell \in \mathbb{Z}_+^{|\mathcal{R}|}$ represents the quantities supplied at time u_ℓ . The starting time S_j for each job J_j is specified by a *schedule* σ , which is *feasible* if (i) the jobs do not overlap in time, and (ii) at any point of time t the total supply from each resource is at least the total request of the jobs starting until t , that is,

$$\sum_{\ell: u_\ell \leq t} \tilde{b}_\ell \geq \sum_{j: S_j \leq t} a_{ij}, \quad \forall i \in \mathcal{R}.$$

Note that in case of just one resource type (as in our starting example), we simply drop the indices corresponding to the single resource. The objective is to minimize the maximum job completion time (makespan) $C_{\max} := \max_{j \in \mathcal{J}} C_j$, where C_j is the completion time of job J_j . In the rest of the paper, we make the following simplifying assumptions guaranteeing sanity of the instances and filtering out trivial cases.

Assumption 1. *Without loss of generality, we assume that*

1. *there are enough resources supplied to process all jobs:* $\sum_{\ell=1}^q \tilde{b}_\ell \geq \sum_{j \in \mathcal{J}} a_{ij}$;
2. *each job has at least one non-zero resource requirement:* $\forall j \in \mathcal{J} \sum_{i \in \mathcal{R}} a_{ij} > 0$; and
3. *at least one resource unit is supplied at time 0:* $\sum_{i \in \mathcal{R}} \tilde{b}_{i,0} > 0$.

The MATERIAL CONSUMPTION SCHEDULING PROBLEM is NP-hard even in case of one machine, only two supply dates ($q = 2$), and if the processing time of each job is

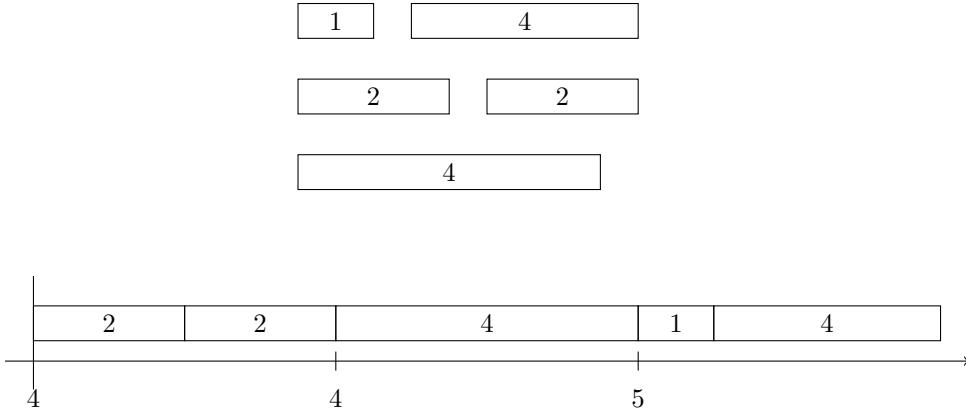


Figure 1: An example input (top) of the problem we consider with one resource type and a solution (bottom) with makespan 12. The processing time of each job is represented by its length and its resource requirement is represented by the number in its rectangle. There are three resource supplies at times 0, 4 and 8 which supply 4, 4, and 5 units of the resource, respectively; these are represented by the three lines with numbers on the time line. Note that the first two jobs have a combined resource requirement of 4 and hence can both be scheduled before the second supply. The third job “consumes” all resources from the second supply but the last two jobs start only after the third (and final) supply and hence their requirements can also be met.

the same as its resource requirement, i.e., $p_j = a_j, \forall j \in \mathcal{J}$ (Carlier 1984). While many variants of the MATERIAL CONSUMPTION SCHEDULING PROBLEM have been studied in the literature in terms of heuristics, polynomial-time approximation algorithms, or the detection of polynomial-time solvable special cases, we are not aware of any previous systematic studies concerning its multivariate complexity analysis. In other words, we seemingly for the first time, we study several natural problem-specific parameters and investigate how they influence the computational complexity of the problem. Doing so, we prove both parameterized hardness as well as encouraging fixed-parameter tractability results for this NP-hard problem.

Related work. Over the recent years, performing multivariate, parameterized complexity studies for fundamental scheduling problems became more and more popular (Bentert, van Bevern, and Niedermeier 2019; van Bevern et al. 2015; van Bevern, Niedermeier, and Suchý 2017; Bodlaender and van der Wegen 2020; Ganian, Hamm, and Mescoff 2020; Hermelin et al. 2017, 2019, 2020; Hermelin, Shabtay, and Talmon 2019; Knop and Koutecký 2018; Mnich and van Bevern 2018; Mnich and Wiese 2015). We contribute to this field by a seemingly first-time exploration of the MATERIAL CONSUMPTION SCHEDULING PROBLEM, focusing on one machine and the minimization of the makespan.

The problem was introduced in the 1980’s (Carlier 1984; Slowinski 1984). Indeed, even a bit earlier a problem where the jobs required non-renewable resources, but without any machine environment, was studied (Carlier and Rinnooy Kan 1982). The problem appears in several real-world applications, e.g., in the continuous casting stage of steel production (Herr and Goel 2016), in order picking in a platform with a distribution company (Belkaid et al. 2012), or in shoe production (Carrera, Ramdane-Cherif, and Portmann 2010).

Carlier (1984) proved several complexity results for dif-

ferent variants in the single-machine case, while Slowinski (1984) studied the parallel machine variant of the problem with preemptive jobs. Previous theoretical results mainly concentrate on the computational complexity and polynomial-time approximability of different variants; in this review we mainly focus on the most important results for the single-machine case and minimizing makespan as the objective. We remark that there are several recent results for variants with other objective functions (Bérczi, Király, and Omlor 2020; Györgyi and Kis 2019, 2020), with a more complex machine environment (Györgyi and Kis 2017), and with slightly different resource constraints (Davari et al. 2020).

Toker, Kondakci, and Erkip (1991) proved that the variant where the jobs require one non-renewable resource reduces to the 2-MACHINE FLOW SHOP PROBLEM provided that the single non-renewable resource has a unit supply in every time period. Later, Xie (1997) generalized this result for multiple resources. Grigoriev, Holthuijsen, and van de Klundert (2005) showed that the variant with unit processing times and two resources is NP-hard, and they also provided several polynomial-time 2-approximation algorithms for the general problem. There is also a PTAS for the variant with one resource and constant number of supply dates and an FPTAS for the case with $q = 2$ supply dates and one non-renewable resource only (Györgyi and Kis 2014). Györgyi and Kis (2015b) presented approximation-preserving reductions between problem variants in case of $q = 2$ and variants of the MULTIDIMENSIONAL KNAPSACK PROBLEM. These reductions have several consequences, e.g., it was shown that the problem is NP-hard if there are two resources, two supply dates and each job has a unit processing time, or that there is no FPTAS for the problem with two non-renewable resources and $q = 2$ supply dates, unless P = NP. Finally, there are three important extensions (Györgyi and Kis 2015a): (i) a PTAS for the variant where the number of re-

sources and the number of supplies dates are constants; (ii) a PTAS for the variant with only one resource and an arbitrary number of supply dates if the resource requirements are proportional to job processing times; and (iii) APX-hardness when the number of resources is part of the input.

Preliminaries and Notation. We use the standard three-fold $\alpha|\beta|\gamma$ notation (Graham et al. 1979), where α denotes the machine environment, β the further constraints like additional resources, and γ the objective function. We always consider a single machine, that is, 1 in the α field. The non-renewable resources are described by nr in the β field and $nr = r$ means that there are r different resource types. The only considered objective is the makespan, C_{\max} . For example, the MATERIAL CONSUMPTION SCHEDULING PROBLEM variant with a single machine, single resource type, and with the makespan as the objective is would be expressed as $1|nr = 1|C_{\max}$. We also sometimes consider the so-called *non-idling scheduling* (introduced by Chrétienne (2008)), in which a machine, indicated by NI in the α field, can only process all jobs continuously, without intermediate idling. As we make a simplifying assumption that the machine has to start processing jobs at time 0, we drop the optimization goal C_{\max} whenever considering non-idling scheduling. When there is just one resource ($nr = 1$), then we write a_j instead of $a_{1,j}$ and \tilde{b}_j instead of $\tilde{b}_{1,j}$, etc. to keep the notation simple. We also write $p_j = 1$ or $p_j = ca_j$ whenever, respectively, jobs have solely unit processing times or the resource requirements are proportional to the job processing times. Finally, we use “unary” to indicate that all numbers in an instance are encoded in unary. Thus, for example, $1, \text{NI}|p_j = 1, \text{unary}|$ —denotes a single, non-idling machine, unit-processing-time jobs and the unary encoding of all numbers.

The following lists the parameters employed in our complexity analysis.

n	number of jobs
q	number of supply periods
j	job index
ℓ	index of supply period
p_j	processing time of job j
$a_{i,j}$	resource requirement of job j from resource i
u_ℓ	the ℓ^{th} supply date
$\tilde{b}_{i,\ell}$	quantity supplied from resource i at u_ℓ
$b_{i,\ell}$	total resource supply from resource i over the first ℓ supplies, i.e., $\sum_{k=1}^{\ell} \tilde{b}_{ik}$

To simplify matters, we introduce the notations p_{\max} , a_{\max} , and b_{\max} for $\max_{j \in \mathcal{J}} p_j$, $\max_{j \in \mathcal{J}, i \in \mathcal{R}} a_{ij}$, and $\max_{\ell \in \{1, \dots, q\}, i \in \mathcal{R}} \tilde{b}_{il}$, respectively.

Primer on Multivariate Complexity. To analyze the parameterized complexity (Cygan et al. 2015; Downey and Fellows 2013; Flum and Grohe 2006; Niedermeier 2006) of MATERIAL CONSUMPTION SCHEDULING PROBLEM, we declare some part of the input the *parameter* (e.g., the number of supply periods). We call a parameterized problem *fixed-parameter tractable* if it is in the class FPT of prob-

lems solvable in time $f(\rho) \cdot |I|^{O(1)}$, where $|I|$ is the size of a given instance encoding, ρ is the value of the parameter, and f is an arbitrary computable (usually super-polynomial) function. Parameterized hardness (and completeness) is defined through parameterized reductions similar to classical polynomial-time many-one reductions. For this paper, it suffices to additionally ensure that the value of the parameter in the problem we reduce to depends only on the value of the parameter of the problem we reduce from. To obtain parameterized intractability, we use parameterized reduction from problems of the class W[1] which is widely believed to be a proper superclass of FPT.

The class XP contains all problems that can be solved in $\text{time } f(\rho)$ for a function f solely depending on the parameter ρ . While XP ensures polynomial-time solvability when ρ is a constant, FPT additionally ensures that the degree of the polynomial is independent of ρ . Unless $\text{P} = \text{NP}$, membership in XP can be excluded by showing that the problem is NP-hard for a constant parameter value (for short, we say the problem is para-NP-hard).

Our contributions. Most of our results are summarized in Table 1. We focus on the parameterized computational complexity of the MATERIAL CONSUMPTION SCHEDULING PROBLEM with respect to several parameters describing resource supplies. We show that the case of a single resource and jobs with unit processing time is polynomial-time solvable. However, if each job has a processing time proportional to the resource requirements, then the MATERIAL CONSUMPTION SCHEDULING PROBLEM becomes NP-hard even for a single resource and each supply providing one unit of the resource. Complementing an algorithm solving the MATERIAL CONSUMPTION SCHEDULING PROBLEM in polynomial time for a constant number q of supply dates, by proving W[1]-hardness, we show that the parameterization by the number of supply dates presumably does not yield fixed-parameter tractability. We circumvent the W[1]-hardness by combining the parameter number q of supply dates with the maximum requirement a_{\max} of a job, thereby obtaining fixed-parameter tractability for the combined parameter $q + a_{\max}$. Moreover, we show fixed-parameter tractability for the parameter u_{\max} that denotes the last resource supply time. Finally, we provide an outlook on cases with multiple resources and show that fixed-parameter tractability for $q + a_{\max}$ extends when we additionally take the number of resources into the combined parameter. For the MATERIAL CONSUMPTION SCHEDULING PROBLEM with an unbounded number of resources, we show intractability even for the case where all other previously discussed parameters are combined.

Due to the lack of space, missing proofs had to be deferred to an appendix.

Computational Complexity Limits

We start our investigation on the computational complexity of the MATERIAL CONSUMPTION SCHEDULING PROBLEM with outlining the limits of efficient computability. Setting up clear borders of tractability, we identify potential

q	b_{\max}	u_{\max}	a_{\max}	$a_{\max} + q$
$1 nr = 1, p_j = 1 C_{\max}$		P \ddagger		
$1 nr = 1, p_j = ca_j C_{\max}$	W[1]-h \diamond , XP \clubsuit	p-NP-h \blacksquare	FPT \diamond	XP \blacktriangle
$1 nr = 1, \text{unary} C_{\max}$	W[1]-h \diamond , XP \clubsuit	p-NP-h \blacksquare	FPT \diamond	XP \blacktriangle
$1 nr = 2, p_j = 1, \text{unary} C_{\max}$	W[1]-h \heartsuit , XP \clubsuit	p-NP-h \blacksquare	?	?
$1 nr = \text{const}, \text{unary} C_{\max}$	W[1]-h \heartsuit , XP \clubsuit	p-NP-h \blacksquare	?	?
$1 nr, p_j = 1, a_j = 1 C_{\max}$	p-NP-h \blacktriangledown	p-NP-h \blacksquare	W[1]-h \blacktriangledown	p-NP-h \blacktriangledown
				W[1]-h \blacktriangledown

Table 1: Results summary. The results are due to Theorem 3 (\ddagger), Theorem 4 (\diamond), Theorem 2 (\diamond), Theorem 1 (\blacksquare), (Györgyi and Kis 2014) (\clubsuit), Proposition 1 (\heartsuit), Proposition 2 (\spadesuit), Theorem 5 (\dagger), Theorem 6 (\blacktriangle), and Theorem 7 (\blacktriangledown). Note that W[1]-h stands for W[1]-hard, p-NP-h stands for para-NP-hard, and a question mark points to an open case.

scenarios suitable for seeking efficient solutions. This approach seems especially justified because the MATERIAL CONSUMPTION SCHEDULING PROBLEM is already NP-hard for a quite constrained scenario of unit processing times and two resources (Grigoriev, Holthuijsen, and van de Klundert 2005).

Both hardness results in this section use reductions from UNARY BIN PACKING. Given a number k of bins, bin size B , and n objects o_1, o_2, \dots, o_n of sizes s_1, s_2, \dots, s_n (encoded in unary) UNARY BIN PACKING asks to distribute the objects to the bin such that no bin exceeds its capacity. UNARY BIN PACKING is NP-hard and W[1]-hard parameterized by the number k of bins even if $\sum_{i=1}^n s_i = kB$ (Jansen et al. 2013).

We first focus on the case of a single resource, for which we find a strong intractability result. In the following theorem, we show that even if each supply comes with a single unit of a resource, then the problem is already NP-hard.

Theorem 1. $1|nr = 1, p_j = ca_j|C_{\max}$ is para-NP-hard with respect to the maximum number b_{\max} of resources supplied at once even if all numbers are encoded unary.

Proof. Given an instance I of UNARY BIN PACKING with $\sum_{i=1}^n s_i = kB$, we construct an instance I' of $1|nr = 1|C_{\max}$ with $b_{\max} = 1$ as described below.

We define n jobs $J_1 = (p_1, a_1), J_2 = (p_2, a_2), \dots, J_n = (p_n, a_n)$ such that $p_i = 2Bs_i$ and $a_i = 2s_i$. We also define a special job $J^* = \{p^*, a^*\}$, with $p^* = 2B$ and $a^* = 1$. Then, we set $2kB$ supply dates as follows. For each $i \in \{0, 1, \dots, k-1\}$ and $x \in \{0, \dots, 2B-1\}$, we create a supply date $q_i^x = (u_i^x, b_i^x) := ((2B + iB^2) - x, 1)$. We add a special supply date $q^* := (0, 1)$. Next, we show that I is a yes-instance if and only if there is a gapless schedule for I' , that is, $C_{\max} = 2(B^2 + B)$.

We first show that each solution to I can be efficiently transformed to a schedule with $C_{\max} = 2(B^2 + B)$. A yes-instance for I is a partition of the objects into k bins such that each bin is (exactly) full. Formally, there are k sets S_1, S_2, \dots, S_k such that $\bigcup_i S_i = O$, $S_i \cap S_j = \emptyset$ for all $i \neq j$, and $\sum_{o_g \in S_j} s_i = B$ for all j . We form a schedule for I' as follows. First, we schedule job j^* and then, continuously, all jobs corresponding to elements of set S_1, S_2 , and so on. The special supply q^* guarantees that the resource requirement of job j^* are met at time point 0. The remaining

jobs, corresponding to elements of the partitions, are scheduled at earliest at time point $2B$, when j^* is processed. The jobs representing each partition, by definition, require in total $2B$ resources and take, in total, $2B^2$ time. Thus, it is enough to ensure that in each point $2B + iB^2$, for $i \in \{0, 1, \dots, k-1\}$ there are at least $2B$ resources available. This is indeed true, because, for all $i \in \{0, 1, \dots, k-1\}$, each time point $2B + iB^2$, in which there is a supply of a single resource, is preceded with $2B - 1$ supplies of one resource. Furthermore, none of the preceding jobs can use the freshly supplied resources as the schedule is must be gapless and all processing times are multiplicities of $2B$. As a result, the schedule is feasible.

Now we show that a gapless schedule for I' implies that I is a yes-instance. Let S be a gapless schedule for I' . Observe that all processing times are multiplicities of $2B$ and therefore each job has to start at a time that is a multiplicity of $2B$. For each $i \in \{0, 1, \dots, k-1\}$, we show that there is no job that is scheduled to start before $2B + iB^2$ and ends after this time. We show this by induction over i . Since at time 0 there is only one resource available, job J^* , with processing time $2B$ must be scheduled first. Hence the statement holds for $i = 0$. Assuming that the statement holds for all $i < i'$ for some i' , we show that it also holds for i' . Assume towards a contradiction that there is a job J that starts before $t := 2B + i'B^2$ and ends after this time. Let S be the set of all jobs that were scheduled to start between $t_0 := 2B + (i' - 1)B^2$ and t . Recall that for each job $J_{i'} \in S$, we have that $p_{i'} = a_{i'}B$. Hence, since J ends after t , the number of resources used by S is larger than $\frac{t-t_0}{B} = B$. Since only $2B$ resources are available at time t , job J cannot be scheduled before time t or there is a gap in the schedule (such a gap would allow to use some of $2B$ supplied in $2B$ times unit just before time t), a contradiction. Hence, there is no job that starts before t and ends after it. Thus, the jobs can be partitioned into “phases,” that is, there are $k+1$ sets T_0, T_1, \dots, T_k such that $T_0 = \{J^*\}$, $\bigcup_{h>0} T_h = \mathcal{J} \setminus \{J^*\}$, $T_h \cap T_j = \emptyset$ for all $h \neq j$, and $\sum_{j_g \in T_g} p_j = 2B^2$ for all g . This corresponds to a bin packing where o_g belongs to bin $h > 0$ if and only if $J_g \in T_h$. \square

Note that Theorem 1 holds also when the processing times of the jobs to schedule are proportional to the resource consumption of the jobs. Additionally, Theorem 1 excludes

CONSUMPTION SCHEDULING PROBLEM, there is some optimal schedule where some job J starts being processed at some time t (in particular, the resource requirements of J are met at t). If, directly after the job the machine idles for some time, we can actually postpone processing J to the latest moment which still guarantees that J is ended before the next job is processed. Naturally, in any case, at the new starting time of J we can only have more resources than at the old starting time. Applying this observation exhaustively produces a solution that is clearly separated into idling time and busy time.

We will now further exploit the observation for the previous paragraph beyond only “moving” jobs without changing their mutual order. We first define a *domination relation* over jobs; in short, some job dominates another job if it is not shorter and at the same time it requires not more resources.

Definition 1. A job J_j *dominates* a job $J_{j'}$, if $p_j \geq p_{j'}$ and, for all $i \in \mathcal{R}$, $a_{i,j} \leq a_{i,j'}$.

Intuitively, when we deal with non-idling schedules, for a pair of jobs J_j and $J_{j'}$ where J_j dominates $J_{j'}$, it is better (or at least not worse) to schedule J_j before $J_{j'}$. Indeed, since among these two, J_j ’s requirements are not greater and its processing time is not smaller, surely after the machine stops processing J_j there will be at least as many resources available as if the machine had processed $J_{j'}$. We formalize this observation in the following lemma.

Lemma 2. For $1|nr|-$, there always exists a feasible schedule where for any pair J_j and $J_{j'}$ of jobs it holds that if J_j dominates $J_{j'}$, then J_j is processed before $J_{j'}$.

Applying Structured Solutions

We start with a polynomial-time algorithm that applies both Lemma 1 and Lemma 2 to solve a specific case of the MATERIAL CONSUMPTION SCHEDULING PROBLEM where each two jobs can be compared according to the domination relation (Definition 1). Indeed, if this is the case, then Lemma 2 exactly specifies the order in which the jobs should be scheduled.

Theorem 3. $1|nr|-$ and $1|nr|C_{\max}$ are solvable in, respectively, cubic and quadratic time if the domination relation is a weak order on a set of jobs. In particular, for the time u_{\max} of the last supply, $1|nr = 1, p_j = 1|C_{\max}$ and $1|nr = 1, a_j = 1|C_{\max}$ are solvable in $O(n \log n \log u_{\max})$ time and $1|nr = 1, p_j = 1|-$ and $1|nr = 1, a_j = 1|-$ are solvable in $O(n \log n)$ time.

Importantly, it is a simple task (requiring at most $O(n^2)$ comparisons) to identify the cases for which the polynomial algorithm above can be applied.

If the given jobs cannot be weakly ordered by domination, then the problem becomes NP-hard as shown in Theorem 1. This is to be expected: When jobs appear that are incomparable with respect to domination, one cannot efficiently decide which job, out of two, to schedule first: the one which requires fewer resource units but has a shorter processing time, or the one that requires more resource units but has a longer processing time. Indeed, it could be the case that sometimes one may want to schedule a shorter job with smaller resource consumption to save resources for later, or sometimes

pseudopolynomial algorithms for the case under consideration since the theorem statement is true also when all numbers are encoded in unary.

Theorem 1 motivates to study further problem-specific parameters. Observe that in the reduction presented in the proof of Theorem 1, we used an unbounded number of supply dates. Györgyi and Kis (2014) have shown a pseudopolynomial algorithm for $1|nr = 1|C_{\max}$ for the case that the number q of supplies is a constant. Thus, the question is whether we can even obtain fixed-parameter tractability for our problem by taking the number of supply dates as a parameter. Devising a reduction from UNARY BIN PACKING. We answer this question negatively in the following theorem.

Theorem 2. $1|nr = 1, p_j = a_j|C_{\max}$ parameterized by the number of supply dates is W[1]-hard even if all numbers are encoded unary.

So far, the theorems presented in this section show that our problem is (presumably) not fixed-parameter tractable neither with respect to the number of supply dates nor with respect to the maximum number of resources per supply. However, as we show in the following section, combining these two parameters allows for efficient solutions. Furthermore, we present other algorithms that, partially, allow us to successfully evade the hardness presented above.

(Parameterized) Tractability

Our analysis of efficient algorithms for our variant of the MATERIAL CONSUMPTION SCHEDULING PROBLEM starts with an introductory part presenting two lemmata exploiting structural properties of the problem’s solutions. Afterwards, we employ the lemmata and provide several tractability results, including polynomial-time solvability for one specific case.

Identifying Structured Solutions

A solution to the MATERIAL CONSUMPTION SCHEDULING PROBLEM is an ordered list of jobs to be executed on the machine(s). Additionally, the jobs need to be associated with their starting times. The starting times have to be chosen in such a way that no job starts when the machine is still processing another scheduled job and that each job requirement is met at the moment of starting the job. We show that, in fact, given an order of jobs, one can always compute the times of starting the jobs minimizing the makespan in polynomial time. We model this observation by presenting a polynomial-time Turing reduction between $1|nr = r|C_{\max}$ and $1|nr = r|-$ in Lemma 1. The crux of this lemma is to observe that there always exists an optimal solution to $1|nr = r|C_{\max}$ that is decomposable into two parts. First, when the machine is idling, and second, when the machine is continuously busy until all jobs are processed.

Lemma 1. There is a polynomial-time Turing reduction from $1|nr = r|C_{\max}$ to $1|nr = r|-$.

Let us intuitively stress the crucial observation backing Lemma 1 since we will extend it in the subsequent

Lemma 2. Assume that, for some instance of the MATERIAL

451 it is better to run a long job consuming, for example, all re-
 452 sources knowing that soon there will be another supply with
 453 sufficient resource units. Since NP-hardness (presumably)
 454 excludes polynomial-time solvability, we turn to parameterized
 455 complexity to get around the intractability.

456 The time u_{\max} of the last supply, seems promising as we
 457 could not show the MATERIAL CONSUMPTION SCHEDUL-
 458ING PROBLEM to be W[1]-hard for this parameterization.
 459 Indeed, we show that this parameterization yields fixed-
 460 parameter tractability. Intuitively, we demonstrate that our
 461 problem is tractable when the time until all resources are
 462 available is short.

463 **Theorem 4.** 1, NI|nr = 1|C_{max} parameterized by the
 464 time u_{\max} of the last supply is fixed-parameter tractable and
 465 can be solved in $O(2^{u_{\max}} \cdot n + n \log n)$ time.

466 *Proof.* We first sort all jobs by their processing time in $O(n)$
 467 time using bucket sort. We then sort all jobs with the
 468 same processing time by their resource requirement in over-
 469 all $O(n \log n)$ time. We then iterate over all subsets R
 470 of $[u_{\max}] = \{1, 2, \dots, u_{\max}\}$. We will refer to $R =$
 471 $\{r_1, r_2, \dots, r_k\}$, where $k = |R|$ and $r_i < r_j$ for $i < j$.
 472 For the sake of simplicity, we will use $r_0 = 0$. For each r_i
 473 in ascending order, we check whether there is a job with a
 474 processing time $r_i - r_{i-1}$ that was not scheduled before and
 475 if so, then we schedule the respective job that is first in each
 476 bucket (the job with the lowest resource requirement). After-
 477 wards we check, whether there is a job left that can be sched-
 478 uled at r_k and which has a processing time at least $u_{\max} - r_k$.
 479 Finally, we schedule all remaining jobs in an arbitrary order
 480 and check whether the total amount of resources suffices to
 481 schedule all jobs.

482 We will now prove that there is a valid gapless sched-
 483 ule if and only if all of these checks are met. Notice that
 484 if all checks are met, then our algorithm provides a valid
 485 gapless schedule. Now assume that there is a valid gapless
 486 schedule. We will show that our algorithm finds a (possi-
 487 bly different) valid gapless schedule. Let, without loss of
 488 generality, $J_{j_1}, J_{j_2}, \dots, J_{j_n}$ be a valid gapless schedule.
 489 Let further k be the index of the last job that is sched-
 490 uled latest at time u_{\max} . We now focus on the iteration
 491 where $R = \{0, p_{j_1}, p_{j_1} + p_{j_2}, \dots, \sum_{i=1}^k p_{j_i}\}$. If the algo-
 492 rithm schedules the jobs $J_{j_1}, J_{j_2}, \dots, J_{j_k}$, then it computes
 493 a valid gapless schedule and all checks are met. Otherwise it
 494 schedules some jobs different but by construction it always
 495 schedules a job with processing time p_{j_i} at position $i \leq k$.
 496 Due to Lemma 2 the schedule computed by the algorithm
 497 is also valid. Thus the algorithm computes a valid gapless
 498 schedule and all checks are met.

499 It remains to analyze the running time. The sorting steps
 500 in the beginning take $O(n \log n)$ time. There are $2^{u_{\max}}$ many
 501 iterations for R and each iteration takes $O(n)$ time as we can
 502 check in constant time for each r_i which job to schedule and
 503 this can be done at most n times (as afterwards there is no
 504 job left to schedule). Searching for the job that is scheduled
 505 at time r_k also takes $O(n)$ time as we can iterate over all
 506 remaining jobs and check in constant whether it fulfills both
 507 requirements. \square

508 Another possibility for fixed-parameter tractability via
 509 parameters measuring the resource supply structure comes
 510 from combining the parameters q and b_{\max} . Although both
 511 parameters alone yield intractability, combining them gives
 512 fixed-parameter tractability in an almost trivial way: By As-
 513 sumption 1, every job requires at least one resource so that
 514 $b_{\max} \cdot q$ is an upper bound for the number of job. Hence,
 515 with this parameter combination, we can try out all possi-
 516 ble schedules without idling (which by Lemma 1 extends to
 517 solving to 1, NI|nr = 1|C_{max}).

518 Motivated by this, we replace the parameter b_{\max} by the
 519 presumably much smaller (and hence practically much more
 520 useful) parameter a_{\max} . Indeed, we consider scenarios with
 521 only few resource supplies and jobs that require only small
 522 units of resources as practically relevant.

523 The following Theorem 5 employs the technique of
 524 Mixed Integer Linear Programming (MILP) (Bredereck
 525 et al. 2020) to positively answer our call for fixed-parameter
 526 tractability for the combined parameter (q, a_{\max}) .

527 **Theorem 5.** 1, NI|nr = 1|C_{max} is fixed-parameter
 528 tractable for the combined parameter $q + b_{\max}$, where q is
 529 the number of supplies and a_{\max} is the maximum resource
 530 requirement per job.

531 *Proof.* Applying the famous theorem of Lenstra (1983), we
 532 describe an integer linear program that uses only $f(q, a_{\max})$
 533 integer variables. Lenstra (1983) showed that an (mixed) in-
 534 teger linear programming is fixed-parameter tractable when
 535 parameterized by the number of integer variables (see also
 536 Frank and Tardos (1987) and Kannan (1987) for later im-
 537 provements). To make the description of the integer pro-
 538 gram significantly easier, we use an extension to integer lin-
 539 ear programs that allows concave transformations on vari-
 540 ables (Bredereck et al. 2020).

541 Our approach is based on two main observations. First,
 542 by Lemma 2 we can assume that there is always an optimal
 543 schedule that is consistent with the domination order. Sec-
 544 ond, within a phase (between two resource supplies), every
 545 job can be arbitrarily reordered. Roughly speaking, a solu-
 546 tion can be fully characterized by the number of jobs that
 547 have been started for each phase and each resource require-
 548 ment.

549 We use the following non-negative integer variables:

1. $x_{w,s}$ denoting the number of jobs requiring s resources
 550 and being started in phase w ,
 551
2. $x_{w,s}^{\Sigma}$ denoting the number of jobs requiring s resources
 552 and being started between phase 1 to w ,
 553
3. α_w denoting the number of resources available in the
 554 beginning of phase w ,
 555
4. d_w denoting the endpoint of phase w , that is, time point
 556 when the job started latest in phase w ends.
 557

558 Naturally, the objective is to minimize d_q .

First, ensure that $x_{w,s}^{\Sigma}$ are correctly computed from $x_{w,s}$
 559 by adding:

$$x_{w,s}^{\Sigma} = \sum_{w'=1}^w x_{w',s} \quad (1)$$

559 Second, we ensure that all jobs are scheduled somewhere. To
 560 this end, let $\#_s$ denote the number of jobs J_j with resource
 561 requirement $a_j = s$.

$$\forall s \in [a_{\max}] : \sum_{w \in [q]} x_{w,s} = \#_s \quad (2)$$

562 Third, we ensure that the α_w variables are set correctly,
 563 by setting $\alpha_1 = \tilde{b}_1$, and $\forall 2 \leq w \leq q$:

$$\alpha_w = \alpha_{w-1} + \tilde{b}_w - \sum_{s \in [a_{\max}]} x_{w-1,s} \cdot s. \quad (3)$$

564 Fourth, we ensure that we have always enough resources:

$$\forall 2 \leq w \leq q : \alpha_w \geq \tilde{b}_w \quad (4)$$

565 Next, we compute the endpoints d_w of each phase, assuming
 566 a schedule that respects the domination order. To this
 567 end, let $p_1^s, p_2^s, \dots, p_{\#_s}^s$ denote the processing times of jobs
 568 with resource requirement exactly s in decreasing order. Fur-
 569 ther, let $\tau_s(y)$ denote the processing time spend to schedule
 570 the y longest jobs with resource requirement exactly s , that
 571 is, we have $\tau_s(y) = \sum_{i=1}^y p_i^s$. Clearly, $\tau_s(x)$ is a concave
 572 function that can be precomputed for each $s \in [a_{\max}]$. To
 573 compute the endpoints, we add:

$$\forall w \in [q] : d_w = \sum_{s \in [a_{\max}]} \tau_s(x_{w,s}^{\Sigma}). \quad (5)$$

574 Since we assume gapless schedules, we ensure that there
 575 is no gap:

$$\forall 1 \leq w \leq q-1 : d_w \geq u_{w+1} - 1 \quad (6)$$

576 This completes the construction of the mixed ILP using
 577 concave transformations. The number of integer variables
 578 used in the ILP is $2q \cdot a_{\max}$ (for $x_{w,s}^{\Sigma}$ variables) plus $2q$ (for
 579 α_w and d_w variables). Moreover, the only concave trans-
 580 formations used in Constraint Set 5 are piecewise linear with
 581 only a polynomial number of pieces (in fact, the number
 582 of pieces is at most the number of jobs), as required to
 583 obtain fixed-parameter tractability of this extended class of
 584 ILPs (Bredereck et al. 2020). \square

585 Motivated by Theorem 5, we are interested in the com-
 586 putational complexity of the MATERIAL CONSUMPTION
 587 SCHEDULING PROBLEM for cases where only a_{\max} is
 588 small. When $a_{\max} = 1$, then we have polynomial-time solv-
 589 ability via Theorem 3. The next theorem shows that this
 590 extends to every constant value of a_{\max} . To obtain this re-
 591 sults, we develop a dynamic programming based algorithm
 592 for $1|NI|nr = r|$ – and apply Lemma 1.

593 **Theorem 6.** $1|nr = 1|C_{\max}$ can be solved in $O(q \cdot a_{\max} \cdot u_{\max} \cdot \log u_{\max} \cdot n^{2a_{\max}})$ time.

595 The question whether $1|nr = 1|C_{\max}$ is in FPT or W[1]-
 596 hard with respect to a_{\max} remains open.

A Glimpse on Multiple Resources

597 So far we focused on scenarios with only one non-renewable
 598 resource. In this section, we give a brief outlook on scenarios
 599 with multiple resources (still considering just one machine).
 600 Naturally, all hardness results transfer. For the tractability
 601 results, we identify several cases where tractability extends
 602 in some form, while other cases become significantly harder.
 603

604 We start our outlook with showing that already with two
 605 resources and unit processing times of the jobs, the MA-
 606 TERIAL CONSUMPTION SCHEDULING PROBLEM becomes
 607 computationally intractable, even when parameterized by
 608 the number of phases. Note that NP-hardness for $1|nr = 2, p_j = 1|C_{\max}$
 609 can also be transferred from Grigoriev, Holthuijsen, and van de Klundert (2005)[Theorem 4] (the
 610 statement is for a different optimization goal but the proof
 611 works).

613 **Proposition 1.** $1|nr = 2, p_j = 1|C_{\max}$ is W[1]-hard when
 614 parameterized by the number of supply dates even if all num-
 615 bers are encoded in unary.

616 Proposition 1 limits the hope for obtaining positive re-
 617 sults for the general case with multiple resources. Still, when
 618 adding the number of different resources to the combined
 619 parameter, we can extend our fixed-parameter tractability re-
 620 sult from Theorem 5. As we expect the number of different
 621 resources to be rather a small constant in many applications,
 622 we consider this result to be quite meaningful.

623 **Proposition 2.** $1|NI|nr = r|C_{\max}$ is fixed-parameter
 624 tractable for the combined parameter $q + a_{\max} + r$, where
 625 q is the number of supplies a_{\max} is the maximum resource
 626 requirement of a job.

627 Finally, by a reduction from INDEPENDENT SET we show
 628 that the MATERIAL CONSUMPTION SCHEDULING PRO-
 629BLEM is intractable even when combining all considered pa-
 630 rameters if the number of resources is unbounded.

631 **Theorem 7.** $1|nr, p_j = 1, a_j = 1|C_{\max}$ is NP-hard and
 632 W[1]-hard parameterized by u_{\max} even if $p_{\max} = a_{\max} =$
 633 $b_{\max} = 1$ and $q = 2$.

Conclusion

634 We provided a seemingly first thorough multivariate
 635 complexity analysis of the MATERIAL CONSUMPTION
 636 SCHEDULING PROBLEM on a single machine. Our main fo-
 637 cuses was on the case of one resource type ($nr = 1$).
 638

639 Particular open questions here refer to the parameter-
 640 ized complexity with respect to the single parameters a_{\max}
 641 and p_{\max} , their combination, and the closely related pa-
 642 rameter number of job types. Notably, this might be chal-
 643 lenging to answer because these questions are closely re-
 644 lated to long-standing open questions for BIN PACKING and
 645 $P||C_{\max}$ (Mnich and van Bevern 2018; Knop and Koutecký
 646 2018; Knop, Koutecký, and Mnich 2019). We have also seen
 647 that cases where the jobs can be ordered with respect to
 648 the domination ordering (Definition 1) are polynomial-time
 649 solvable. It seems promising to consider structural pa-
 650 rameters measuring the distance with this tractable case.

651 Our results for more than one resource type mean only
 652 first steps. They certainly invite to further investigations.

References

- References**

Belkaid, F.; Maliki, F.; Boudahri, F.; and Sari, Z. 2012. A Branch and Bound Algorithm to Minimize Makespan on Identical Parallel Machines with Consumable Resources. In *Advances in Mechanical and Electronic Engineering*, 217–221. Springer.

Bentert, M.; van Bevern, R.; and Niedermeier, R. 2019. Inductive k -independent graphs and c -colorable subgraphs in scheduling: a review. *Journal of Scheduling* 22(1): 3–20.

Bérczi, K.; Király, T.; and Omlor, S. 2020. Scheduling with Non-renewable Resources: Minimizing the Sum of Completion Times. In *Proceedings of the 6th International Symposium on Combinatorial Optimization*, 167–178. Springer.

van Bevern, R.; Mnich, M.; Niedermeier, R.; and Weller, M. 2015. Interval scheduling and colorful independent sets. *Journal of Scheduling* 18(5): 449–469.

van Bevern, R.; Niedermeier, R.; and Suchý, O. 2017. A parameterized complexity view on non-preemptively scheduling interval-constrained jobs: few machines, small looseness, and small slack. *Journal of Scheduling* 20(3): 255–265.

Bodlaender, H. L.; and van der Wegen, M. 2020. Parameterized Complexity of Scheduling Chains of Jobs with Delays. *CoRR* abs/2007.09023.

Bredereck, R.; Faliszewski, P.; Niedermeier, R.; Skowron, P.; and Talmon, N. 2020. Mixed integer programming with convex/concave constraints: Fixed-parameter tractability and applications to multicovering and voting. *Theoretical Computer Science* 814: 86–105.

Carlier, J. 1984. *Problèmes d'ordonnancements à contraintes de ressources: algorithmes et complexité*. Thèse d'état. Université Paris 6.

Carlier, J.; and Rinnooy Kan, A. H. G. 1982. Scheduling subject to nonrenewable resource constraints. *Operational Research Letters* 1: 52–55.

Carrera, S.; Ramdane-Cherif, W.; and Portmann, M.-C. 2010. Scheduling supply chain node with fixed component arrivals and two partially flexible deliveries. In *Proceedings of the 5th International Conference on Management and Control of Production and Logistics (MCPL '10)*, 6. IFAC Publisher.

Chrétienne, P. 2008. On single-machine scheduling without intermediate delays. *Discrete Applied Mathematics* 156(13): 2543–2550.

Cygan, M.; Fomin, F. V.; Kowalik, L.; Lokshtanov, D.; Marx, D.; Pilipczuk, M.; Pilipczuk, M.; and Saurabh, S. 2015. *Parameterized Algorithms*. Springer.

Davari, M.; Ranjbar, M.; De Causmaecker, P.; and Leus, R. 2020. Minimizing makespan on a single machine with release dates and inventory constraints. *European Journal of Operational Research* 286(1): 115–128.

Downey, R. G.; and Fellows, M. R. 2013. *Fundamentals of Parameterized Complexity*. Springer.

Flum, J.; and Grohe, M. 2006. *Parameterized Complexity Theory*. Springer.

Frank, A.; and Tardos, É. 1987. An application of simultaneous Diophantine approximation in combinatorial optimization. *Combinatorica* 7(1): 49–65.

Ganian, R.; Hamm, T.; and Mescoff, G. 2020. The Complexity Landscape of Resource-Constrained Scheduling. In *Proceedings of the 29th International Joint Conference on Artificial Intelligence (IJCAI '20)*, 1741–1747.

Graham, R. L.; Lawler, E. L.; Lenstra, J. K.; and Kan, A. R. 1979. Optimization and approximation in deterministic sequencing and scheduling: a survey. *Annals of Discrete Mathematics* 5: 287–326.

Grigoriev, A.; Holthuijsen, M.; and van de Klundert, J. 2005. Basic scheduling problems with raw material constraints. *Naval Research of Logistics* 52: 527–553.

Györgyi, P.; and Kis, T. 2014. Approximation schemes for single machine scheduling with non-renewable resource constraints. *Journal of Scheduling* 17: 135–144.

Györgyi, P.; and Kis, T. 2015a. Approximability of scheduling problems with resource consuming jobs. *Annals of Operations Research* 235(1): 319–336.

Györgyi, P.; and Kis, T. 2015b. Reductions between scheduling problems with non-renewable resources and knapsack problems. *Theoretical Computer Science* 565: 63–76.

Györgyi, P.; and Kis, T. 2017. Approximation schemes for parallel machine scheduling with non-renewable resources. *European Journal of Operational Research* 258(1): 113 – 123.

Györgyi, P.; and Kis, T. 2019. Minimizing total weighted completion time on a single machine subject to non-renewable resource constraints. *Journal of Scheduling* 22(6): 623–634.

Györgyi, P.; and Kis, T. 2020. New complexity and approximability results for minimizing the total weighted completion time on a single machine subject to non-renewable resource constraints. *arXiv preprint arXiv:2004.00972*.

Hermelin, D.; Kubitz, J.; Shabtay, D.; Talmon, N.; and Woeginger, G. J. 2017. Scheduling Two Agents on a Single Machine: A Parameterized Analysis of NP-hard Problems. *CoRR* abs/1709.04161.

Hermelin, D.; Manoussakis, G.; Pinedo, M.; Shabtay, D.; and Yedidsion, L. 2020. Parameterized Multi-Scenario Single-Machine Scheduling Problems. *Algorithmica* 82(9): 2644–2667.

Hermelin, D.; Pinedo, M.; Shabtay, D.; and Talmon, N. 2019. On the parameterized tractability of single machine scheduling with rejection. *European Journal of Operational Research* 273(1): 67–73.

Hermelin, D.; Shabtay, D.; and Talmon, N. 2019. On the parameterized tractability of the just-in-time flow-shop scheduling problem. *Journal of Scheduling* 22(6): 663–676.

- 758 Herr, O.; and Goel, A. 2016. Minimising total tardiness for
759 a single machine scheduling problem with family setups and
760 resource constraints. *European Journal of Operational Re-*
761 *search* 248(1): 123–135.
- 762 Jansen, K.; Kratsch, S.; Marx, D.; and Schlotter, I. 2013.
763 Bin packing with fixed number of bins revisited. *Journal of*
764 *Computer and System Sciences* 79(1): 39–49.
- 765 Kannan, R. 1987. Minkowski's convex body theorem and
766 integer programming. *Mathematics of Operations Research*
767 12(3): 415–440.
- 768 Knop, D.; and Koutecký, M. 2018. Scheduling meets n -fold
769 integer programming. *Journal of Scheduling* 21(5): 493–
770 503.
- 771 Knop, D.; Koutecký, M.; and Mnich, M. 2019. Combinatorial
772 n -fold integer programming and applications. *Mathe-*
773 *matical Programming*.
- 774 Lenstra, Jr, H. W. 1983. Integer Programming with a Fixed
775 Number of Variables. *Mathematics of Operations Research*
776 8(4): 538–548.
- 777 Mnich, M.; and van Bevern, R. 2018. Parameterized com-
778 plexity of machine scheduling: 15 open problems. *Comput-*
779 *ers & Operations Research* 100: 254–261.
- 780 Mnich, M.; and Wiese, A. 2015. Scheduling and fixed-
781 parameter tractability. *Mathematical Programming* 154(1–
782 2): 533–562.
- 783 Niedermeier, R. 2006. *Invitation to Fixed-Parameter Algo-*
784 *rithms*. Oxford University Press.
- 785 Slowinski, R. 1984. Preemptive scheduling of independent
786 jobs on parallel machines subject to financial constraints.
787 *European Journal of Operational Research* 15: 366–373.
- 788 Toker, A.; Kondakci, S.; and Erkip, N. 1991. Scheduling
789 under a non-renewable resource constraint. *Journal of the*
790 *Operational Research Society* 42: 811–814.
- 791 Xie, J. 1997. Polynomial algorithms for single machine
792 scheduling problems with financial constraints. *Operations*
793 *Research Letters* 21(1): 39–42.