

Dynasearch neighborhood for the earliness-tardiness scheduling problem with release dates and setup constraints

Francis Sourd *

Université Pierre et Marie Curie – CNRS –LIP6
4, place Jussieu — 75252 Paris Cedex 05, France

Abstract

The one-machine scheduling problem with sequence-dependent setup times and costs and earliness-tardiness penalties is addressed. This problem is NP-complete, so that local search approaches are very useful to efficiently find good feasible schedules. In this paper, we present an extension of the dynasearch neighborhood for this problem. Finding the best schedule in this neighborhood is shown to be NP-complete in the ordinary sense. However, the neighborhood can be efficiently explored in pseudo-polynomial time. Computational tests are finally presented.

1 Introduction

Local search algorithms are widely used as a practical approach for solving combinatorial optimization problems. Starting with a feasible solution, these algorithms iteratively try to improve the *current* solution by searching a better solution in the *neighborhood* of the current solution until a *local minimum* is found. The efficiency of these algorithms critically depends on the definition of the neighborhood: with a larger neighborhood, the quality of the local minimum is generally better but the computation time required to explore the neighborhood is longer. So, in practice, large neighborhood are not useful unless it can efficiently be explored. Such an approach is referred to as a *very large-scale neighborhood search technique*, this term is popularized in the survey by Ahuja et al. [1].

In the scheduling literature, Congram et al. [2] present a large neighborhood — called *dynasearch* — for the single-machine total weighted tardiness scheduling problem. This neighborhood can be seen as the composition of an arbitrary number of independent swap operations. Its size is exponential but a dynamic programming algorithm is proposed to compute the optimal schedule in the neighborhood in polynomial time. Based on this work, Grosso et al. [3] propose an enhanced dynasearch neighborhood that combines the swap operator used in [2] with a second “extract and reinsert” operator.

In this paper, we extend this approach to a more general and a more practical problem that appears in particular in manufacturing scheduling problems [4]. More precisely, our problem deals with setup times and setup costs, release dates, deadlines and a general

*Corresponding author. Fax: +1 (+33) 1 44 27 70 00 - Francis.Sourd@lip6.fr

end-time dependent costs that can for example model the usual earliness-tardiness penalties. The reader is referred to [5] for a presentation of the relevance of this problem and for related references in the just-in-time and setup scheduling literature. Mathematical programming formulations of the problem, a branch-and-bound algorithm and an efficient heuristic procedure are also presented in [5].

Section 2 introduces the scheduling problem. The associated dynasearch problem, which consists in finding the optimal solution in the so-called dynasearch neighborhood, is shown to be NP-complete in the ordinary sense. Section 3 is devoted to the presentation of a pseudo-polynomial algorithm to search the neighborhood and experimental results are eventually presented in Section 4.

2 Problem definition and complexity

A single machine has to process n tasks J_1, \dots, J_n such that at most one task is performed at any time. Preemption of tasks is not allowed but idle time can be inserted between two tasks. Each task J_i has a processing time p_i and belongs to a *group* (or *family*) $g_i \in \{1, \dots, q\}$ (with $q \leq n$). *Setup* or *changeover* times and costs, which are given as two $q \times q$ matrices, are associated to these groups. This means that in a schedule where J_j is processed immediately after J_i , there must be a setup time of at least $s(g_i, g_j)$ time units between the completion time of J_i , denoted by C_i , and the start time of J_j , which is $C_j - p_j$. During this setup period, no other task can be performed by the machine and we assume that the cost of the setup operation is $c(g_i, g_j) \geq 0$. We assume that there is no setup time and no setup cost between tasks belonging to the same group (that is $s(g, g) = 0$ and $c(g, g) = 0$) and the setup matrices satisfy the triangle inequalities (that is $s(g, g') \leq s(g, g'') + s(g'', g')$ and $c(g, g') \leq c(g, g'') + c(g'', g')$). A machine is said to be *in state* g as soon as the setup operation moving it into state g is finished and it remains in state g until another setup operation starts. Machine state is undefined during setup operations. Such setup times and costs are said to be *sequence-dependent* because they depend on both g and g' . When $s(g, g')$ and $c(g, g')$ ($g \neq g'$) can be expressed as a function in only the state g' , setups are said to be *sequence-independent*.

In addition to the setup constraints and costs, a cost $f_i(C_i)$ is due for each task. Such a cost depends on the completion time of J_i and is called *punctuality cost* in this paper. The idea is that the cost is low during the periods satisfying the manufacturer and the costumer and higher otherwise. More details about the possible use of such a cost function is given in [6]. Clearly, punctuality cost is a generalization of the *deviation* cost function $ET_i = \max(\alpha_i(d_i - C_i), \beta_i(C_i - d_i))$ used in earliness-tardiness scheduling. Note that release dates can easily be integrated in a function f_i .

In this paper, we are going to assume that the function f_i are piecewise linear, which is a very common and efficient way to represent and approximate functions in computer science. For simplicity, we assume that the cost functions are continuous. The number of segments of f_i will be denoted by $\|f_i\|$.

The problem is to minimize the sum of the punctuality and setup costs. In this paper, we assume that all the numerical time-related values (setup times, processing times, abscissas of piecewise linear functions) are integer. The problem is NP-complete in the

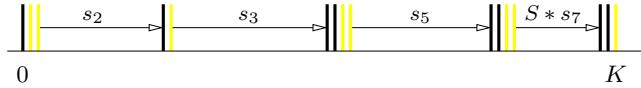


Figure 1: Complexity proof

strong sense — observe that it generalizes the one-machine earliness-tardiness scheduling problem. We now formally define the *dynasearch neighborhood*.

Let us consider an *initial sequence*. The tasks can be renumbered such that the sequence is (J_1, \dots, J_n) . Following [3], we define three operators which, given a pair of indices $i < j$, act on the sequence $\sigma = \alpha J_i \beta J_j \gamma$ as follows:

- SWAP _{i,j} ($\alpha J_i \beta J_j \gamma$) = $\alpha J_j \beta J_i \gamma$. When β is the empty sequence, the operator is also known as the *adjacent pairwise interchange*.
- EBSR _{i,j} ($\alpha J_i \beta J_j \gamma$) = $\alpha J_j J_i \beta \gamma$. EBSR stands for *extraction and backward-shifted reinsertion*.
- EFSR _{i,j} ($\alpha J_i \beta J_j \gamma$) = $\alpha \beta J_j J_i \gamma$. EFSR stands for *extraction and forward-shifted reinsertion*.

Two operators acting on position pairs $i < j$ and $k < l$ are called *independent* if $j < k$ or $l < i$. The *dynasearch sequence neighborhood* is the set of sequences that can be derived from the initial sequence through the application of a set of independent SWAP, EBSR and EFSR operators. Eventually, the *dynasearch neighborhood* is the set of schedules whose task sequences are in the sequence neighborhood. As the objective criterion of the problem is not a regular one, there is not an obvious one-to-one mapping between a sequence and a schedule — in other words, for a given sequence, the earliest schedule may not be the optimal one.

The *dynasearch problem* is, given a schedule or a sequence σ , to find a schedule in the dynasearch neighborhood of σ with minimal cost. The decision variant of this minimization problem is to find a schedule with a cost less than some given value K .

Theorem 1. *The dynasearch problem is NP-complete in the ordinary sense even if each punctuality cost function is the weighted tardiness indicator $w_i T_i = w_i \max(0, C_i - d)$ for some common due date d .*

Proof. The problem is clearly in NP. We transform PARTITION (which is NP-complete) to the dynasearch decision problem. Let an arbitrary instance of PARTITION given by the finite set $A = \{1, 2, \dots, k\}$ and the size $s_a \in \mathbb{Z}^+$ for each item $a \in A$. Let $K = \sum_{a \in A} s_a / 2$. We are going to build an instance of the dynasearch problem whose cost is less than or equal to K if and only if there exists a bipartition of A .

The scheduling problem has $n = 2k$ tasks J_1, J_2, \dots, J_n sequenced in this order. For any $a \in A$, the two jobs J_{2a-1} and J_{2a} are related to item a of PARTITION (J_i is related to $\lceil i/2 \rceil$). The main idea of the reduction is to partition A into B and $B - A$ with $B = \{a | C_{2a-1} < C_{2a}\}$ where C_{2a-1} and C_{2a} are respectively the completion time of J_{2a-1}

and J_{2a} in a schedule of the dynasearch neighborhood.

The processing time of any task J_i is 0. The punctuality cost function is $f_i(C_i) = (K + 1) \max(0, C_i - K)$ for each task. There are n setup groups, let $g_i = i$ be the group of J_i . Let $c(2a - 1, 2a) = s(2a, 2a - 1) = s_a$ and $s(2a - 1, 2a) = c(2a, 2a - 1) = 0$. Otherwise, that is if $\lceil i/2 \rceil \neq \lceil j/2 \rceil$,

$$c(i, j) = s(i, j) = \begin{cases} 0 & \text{if } i < j \\ K + 1 & \text{otherwise.} \end{cases}$$

Figure 1 illustrates the reduction for $k = 8$ items. The 16 jobs are sequenced in the order $J_1 J_2 J_4 J_3 J_6 J_5 J_7 J_8 J_10 J_9 J_{11} J_{12} J_{14} J_{13} J_{15} J_{16}$. We have $B = \{1, 4, 6, 8\}$. Setup times are non-null when $i > j$ and J_i precedes J_j .

We first prove that any schedule with cost less than or equal to K is in the dynasearch neighborhood. Since the cost functions are nondecreasing, we can assume that the schedule has no idle time so that its makespan is equal to the length of all the setup times, which is integer. As the cost is less than K , the makespan must be less than or equal to K . The setup time table shows that, if $a < a'$, any job related to a must be sequenced before any job related to a' . So, the sequence differs from the initial sequence by a number of swaps between adjacent jobs related to the same item. These swap operations are all independent so that this schedule is in the dynasearch neighborhood. The length of this schedule, which is less than K , is clearly $\sum_{a \notin B} s(2a, 2a - 1) = \sum_{a \notin B} s_a$ and the total cost is equal to the setup costs which are $\sum_{a \in B} s(2a, 2a - 1) = \sum_{a \in B} s_a$.

Therefore, there exists a schedule with cost at most K in the dynasearch neighborhood if and only if there is a bipartition of A . \square

In the next section, we show that the optimal schedule of the dynasearch problem can be computed in pseudo-polynomial time.

3 Dynasearch algorithm

Based on the approach introduced by [6], let us define $\Sigma_k^g(t)$ as the cost of optimally scheduling the tasks J_1, \dots, J_k subject to the constraints:

- the sequence is in the dynasearch neighborhood of (J_1, \dots, J_k) ,
- the group of the last task (*i.e.* the task in position k) is g ,
- the last task completes before time t .

Clearly, Σ_k^g is a real function. The value $\Sigma_k^g(t)$ is set to ∞ when there is no feasible schedule ending before t (for instance, if $t < \sum_{i=1}^k p_i$). We also define $\Sigma_0^g(t)$ as the cost of setting up the machine to state g before t without scheduling any task.

3.1 Recurrence equation

For $k \in \{0, 1, \dots, n\}$, let Σ_k be the vector of q functions $(\Sigma_k^1, \dots, \Sigma_k^q)$. Intuitively, the g^{th} coordinate of the vector represents the cost of scheduling the jobs J_1, \dots, J_k before t such

$$\Sigma_k = \min \left\{ \begin{array}{l} \Sigma_{k-1} \bullet J_k \\ \min_{1 \leq i \leq k} (\Sigma_{i-1} \bullet J_k \bullet J_{i+1} \bullet \dots \bullet J_{k-1} \bullet J_i) \\ \min_{1 \leq i \leq k} (\Sigma_{i-1} \bullet J_k \bullet J_i \bullet \dots \bullet J_{k-1}) \\ \min_{1 \leq i \leq k} (\Sigma_{i-1} \bullet J_{i+1} \bullet \dots \bullet J_k \bullet J_i) \end{array} \right\}$$

Figure 2: Computing Σ_k

that the machine is in state g at the end. If none of the jobs J_1, \dots, J_k is in group g , this cost is infinite. For any vector $\Sigma = (\Sigma^1, \dots, \Sigma^q)$ of q functions, we define the operator $\Sigma \bullet J_k$ that returns the vector of q functions whose g_k^{th} coordinate is

$$(\Sigma \bullet J_k)^{g_k}(t) = \min_{1 \leq g \leq q} \min_{t' < t} (\Sigma^g(t' - p_k - s(g, g_k)) + f_k(t')) + c(g, g_k)$$

and other function coordinates are infinite. It represents the cost of scheduling the tasks such that J_k is the last task. This expression is directly derived from the one of Sourd [6].

Figure 2 shows that Σ_k can be recursively derived from $\Sigma_0, \Sigma_1, \dots, \Sigma_{k-1}$. To understand the expression, we consider the sequence that correspond to the schedule yielding $\Sigma_k^g(t)$. This sequence has been obtained through a set of independent operators among SWAP, EBSR, EFSR. If J_k is the last operation of the sequence, then the cost of the schedule is clearly given by the first line in expression of Figure 2. Otherwise, that is the position of J_k has been modified by a SWAP, EBSR or EFSR operator respectively, then the cost of the schedule is given by the second, third or fourth line in Figure 2 respectively.

The cost of the best schedule in the dynasearch neighborhood is finally given by $\min_{1 \leq g \leq q} \min_{t > 0} \Sigma_n^g(t)$. Clearly, all the functions Σ_k^g are piecewise linear (see also [6]) and, since the processing times and the setup times are all integer, the abscissas of all the breakpoints of these functions are integer. This proves that the number of breakpoints is bounded by any constant T which is an upper bound on the makespan of the optimal schedule. For example, we can choose T as the sum of all the processing times and of $n \max_{gg'} s(g, g')$ and of the largest abscissa of all the cost functions f_i . This proves that the number of segments of the functions Σ_k^g is pseudo-polynomial in the input of the problem so that the dynasearch problem is also pseudo-polynomial. More precisely, computing $\Sigma \bullet J_k$ requires $O(qT)$ time [6] so that the total computation time is in $O(n^2qT)$.

We observe that the dynamic program can be solved either

- by storing all the values $\Sigma_k^g(t)$ for $k \in \{1, \dots, n\}$, $g \in \{1, \dots, q\}$, $t \in \{1, \dots, n\}$ in an array of size $O(nqT)$ or
- by storing the $O(nq)$ piecewise linear functions as the lists of their segments (see [6])

In practice, the second option is far more efficient and we adopt this approach in the rest of the paper.

3.2 Backtracking

As the dynasearch procedure is to be iterated, retrieving the sequence that yields the optimal schedule in the dynasearch neighborhood is a very necessary step of the algo-

rithm. Because the states of dynamic program are stored by the means of piecewise linear functions, the classical techniques used in dynamic programming must be adapted.

Clearly, the completion time of the last task in the optimal schedule can be set as the minimal t among all the pairs (g, t) minimizing $\Sigma_n^g(t)$. Let $(g(n), t(n))$ denote the corresponding pair, $g(n)$ is the group to which this last task belongs, but which is precisely the task? To answer this question, we must store in the segments of the functions Σ_k^g some information that records how the segment was created. The expression in Figure 2 shows that each function Σ_k^g is the minimum function of a finite set S of piecewise linear functions. So, each segment of Σ_k^g also belongs to at least one function of S .

So, let us specify how the segments are marked when computing Σ_k . Let us first consider the segments created by $\Sigma_{k-1} \bullet J_k$, which requires the computation of at most q piecewise linear functions. For each $g \in \{1, \dots, q\}$, all the segments of the function $t \mapsto \min_{t' < t} (\Sigma_{k-1}^g(t' - p_k - s(g, g_k)) + f_k(t')) + c(g, g_k)$ are marked with tag (NOP, g) . The first element NOP in the pair indicates that the last element (job J_k) is not concerned with any operator, so that the computation of Σ_k is derived from Σ_{k-1} , the second element g specify that this function is related to the g^{th} coordinate of Σ_{k-1} , that is the machine was previously in state g .

Similarly, the segments of the g_i^{th} coordinate of $\Sigma_{i-1} \bullet J_k \bullet J_{i+1} \bullet \dots \bullet J_{k-1} \bullet J_i$, are marked with tag (SWAP_{ik}, g) . SWAP _{i,k} indicates that the sequence between J_i and J_k is modified by the SWAP operator so that the function is derived from Σ_{i-1} and more precisely its g^{th} coordinate. Similar tags are used for the EBSR and EFSR operators.

When backtracking the dynamic program, we have to fix the completion time of all the jobs. We have already seen how to fix the completion time $t(n)$ of the last task. The name of the last task is revealed by the tag of the segment yielding $\Sigma_n^{g(n)}(t(n))$.

If the tag is (NOP, g) , then the last job is unchanged so it is J_n . Since J_n completes at $t(n)$, the previous job must complete before $\bar{t}(n-1) = t(n) - p_n - s(g, g(n))$ and belongs to the group $g(n-1) = g$. So we can go on the backtracking procedure by determining the completion time $t(n-1)$ of the task in position $n-1$, which is the earliest time $t \leq \bar{t}(n-1)$ that minimizes $\Sigma_{n-1}^{g(n-1)}$ on $[0, \bar{t}(n-1)]$.

If the tag is (SWAP_{in}, g) , we similarly know that the backtracking procedure must go on with the function $\Sigma_{i-1}^{g(i-1)}$ with $g(i-1) = g$ but we must also calculate the completion time $t(i-1)$ of the task in position $i-1$ in the new sequence. A *local* dynamic program is necessary for that. It first recomputes $\Sigma'_i = \Sigma_{i-1} \bullet J_n$, then $\Sigma'_{i+1} = \Sigma_{i-1} \bullet J_n \bullet J_{i+1}$ and so on until $\Sigma'_{n-1} = \Sigma_{i-1} \bullet J_n \bullet J_{i+1} \bullet \dots \bullet J_{n-1}$. As the last operation in the new optimal sequence is J_i , we have that the penultimate task — which is J_{n-1} if $i < n-1$ and J_n if $i = n-1$ — must end before $\bar{t}(n-1) = t(n) - p_i - s(g_{n-1}, g_i)$. So, the completion time $t(n-1)$ of the penultimate task is the earliest time before $\bar{t}(n-1)$ such that Σ'_{n-1} is minimum. Similarly, $t(n-2), \dots, t(i)$ are recursively computed. Then, we know that the job in position $i-1$ must end before $\bar{t}(i-1) = t(i) - p_n - c(g, g_n)$ and thus $t(i-1)$ is derived from the analysis of Σ_{i-1}^g on the interval $[0, \bar{t}(i-1)]$.

The cases where the tag is (EBSR_{in}, g) or (EFSR_{in}, g) are solved similarly. The backtracking procedure is iterated until all the completion times are computed.

3.3 Dominance rules and speedups

As for the implementations of the dynasearch procedures for the total weighted tardiness problem [2, 3], several dominance rules and speedups can be implemented for the proposed procedure. Unfortunately, dominance rules in presence of earliness costs and setup times and costs are significantly weaker than for the total weighted tardiness problem.

- If $p_i = p_j$, $g_i = g_j$ and $f'_i \geq f'_j$ (that is, in the earliness-tardiness case, $\alpha_i \leq \alpha_j$, $d_i \leq d_j$ and $\beta_i \geq \beta_j$), then J_i must precede J_j .
- If $g_i = g_{i+1}$ and $f_i(t + p_i) + f_{i+1}(t + p_i + p_{i+1}) \leq f_{i+1}(t + p_{i+1}) + f_i(t + p_i + p_{i+1})$ then J_i must precede J_{i+1} .
- Let \bar{C}_i be a lower bound of the completion time of the job in position i when jobs J_1, \dots, J_i are sequenced before all the other jobs. For example, we can take $\bar{C}_i = \sum_{i=1}^n p_i$. In the earliness-tardiness case, if $\max(d_i - p_i, d_{i+1} - p_{i+1}) \leq C_i$, then both J_i and J_{i+1} are late so that if $g_i = g_{i+1}$ and $p_i/\beta_i \leq p_{i+1}/\beta_{i+1}$ then J_i must precede J_{i+1} .

We also observe that all the functions $\Sigma_{i-1} \bullet J_{i+1} \bullet \dots \bullet J_k$ for $k > i$ can be recursively computed and stored once Σ_{i-1} has been computed. They are then reused (instead of recomputed from scratch) when computing the vectors Σ_k for $k > i$.

4 Experimental results

In order to evaluate the behavior of the dynasearch neighborhood, we compared it to a simple descent algorithm based on the *union* of the three neighborhoods SWAP, EBSR and EFSR — by contrast, the dynasearch neighborhood is seen as the *composition* of these neighborhoods. This algorithm is described in [5] and is used to find an initial feasible solution in the branch-and-bound procedure. Even if a single descent only finds a local optimum, the iteration of several descent procedure from randomly sequenced tasks is shown to be an efficient search algorithm, at least for small instances whose optimality can be computed. Indeed, when the algorithm is run n times the (global) optimum is found at least once for 95.2% of the 2160 instances with $n \in \{10, 12, 15\}$ tasks, the mean deviation is 0.08% and the maximal deviation is less than 10%.

With regard to the dynasearch algorithm, preliminary tests have shown that finding the best schedule in the dynasearch neighborhood is quite time-consuming. So, a descent algorithm that would run the dynasearch procedure at each step is not efficient at all. Therefore, in the tests presented below, the dynasearch descent procedure is started with the local optimum found by the simple descent procedure.

We used the test generator presented in [5] to create instances with $n \in \{30, 60, 100\}$ tasks and $q \in \{3, 10\}$ setup groups. In order to limit the complexity of the result tables, we fixed the other parameters which, according to preliminary tests, are not very significant for the computation times and solution quality :

- The *average tardiness* τ is equal to 0.8 and the *range factor* is set to 0.2.

- The *release date factor* (which is null in [5]) is set to 0.5. It means that the release dates are generated from the uniform distribution $[0, 0.5 \sum p_i]$.
- The setup times and costs are generated from the uniform distribution $[50, 100]$.

For each pair (n, q) , five instances were generated so that the benchmark set contains 30 different instances. Both algorithms are implemented in C++ and were run on a 1GHz Personal Computer under the Windows 2000 operating system. For each instance, the simple descent algorithm was run 20 times starting with a random sequence and the dynasearch descent algorithm was run starting with each one of the twenty local optima found by the simple descent.

Table 1 compares the simple descent algorithm and the dynasearch based algorithm. Each line of the table reports the behavior of the two algorithms for an instance, whose size (n, q) is indicated in the first two columns. The third column represents the best known solution for the instance, which is the solution found by the dynasearch procedure. The three columns related to the simple descent algorithm respectively reports the deviation between the best solution found by the algorithm (after 20 descents) and the best known solution, the mean deviation (for the twenty runs) and the mean running time. For the dynasearch algorithm, there is no “best” column because the best solution is by definition the solution of this algorithm.

The experimental tests show that the dynasearch procedure generally improves the solution found by the simple descent algorithm. However, it must be noted that computation times are longer. The improvement is significantly better when the number of jobs or the number of setup groups become larger.

For the weighted tardiness problems, Grosso et al. [3] show that the usage of EBSR and EFSR significantly improves the performance of the dynasearch procedure. We tested whether the same conclusion can be drawn for the earliness-tardiness problem with setups by running “limited” versions of our algorithms in which only the SWAP operator is active. Results for the five instances with $n = 60$ and $q = 10$ are presented in Table 2. Even if computation times are slightly decreased, performance is greatly unimproved: the optimum is never found and the deviation between the best solution found by either the simple descent procedure or the dynasearch algorithm is significantly larger. This result shows that combining SWAP neighborhoods with EBSR and EFSR neighborhoods is critically important to produce an efficient search algorithm.

5 Conclusion

In the paper, we have presented an extension of the dynasearch procedure for the one-machine problem with release dates, setup times and costs and weighted earliness-tardiness penalties. We have shown that the related dynasearch problem is NP-complete in the ordinary sense but it can be solved in pseudo-polynomial time. As the dynasearch problem for $1||\sum w_i T_i$ is known to be polynomial [2, 3], future research will be devoted to study the complexity classification of the dynasearch problems. In particular, the minimal open dynasearch problem is the problem associated to $1|r_i| \sum_i C_i$.

In our experiments, we have shown that the dynasearch procedure is able to improve the solution found by the simple descent algorithm so that it can be regarded in practice

n	q	Best solution	Local Search			Dynasearch	
			Best	Avg	time (s)	Avg	time (s)
30	3	33866	0.00	6.07	0.06	4.88	0.15
		27742	0.00	2.93	0.06	2.20	0.16
		33586	0.00	5.47	0.06	3.44	0.18
		39235	0.00	7.29	0.06	7.07	0.14
		40372	0.00	6.57	0.06	6.53	0.14
	10	46742	0.00	3.97	0.05	3.87	0.19
		48510	1.30	9.73	0.05	8.44	0.20
		58859	0.00	7.43	0.06	6.97	0.22
		53429	0.00	2.72	0.05	1.77	0.28
		65337	1.15	7.63	0.06	6.20	0.23
60	3	133702	0.00	5.85	0.80	5.81	1.50
		102482	0.00	4.26	0.71	4.23	1.44
		152685	0.00	2.98	0.70	2.50	1.35
		139446	0.00	3.56	0.61	3.38	1.62
		173338	0.39	5.60	0.60	5.47	1.24
	10	178464	0.00	5.94	0.58	4.99	1.88
		196661	0.00	7.05	0.57	5.71	2.18
		185427	1.56	6.62	0.52	5.27	2.92
		165989	0.00	6.98	0.56	6.46	1.80
		157513	0.57	8.20	0.57	5.70	2.93
100	3	356027	0.24	4.08	4.01	3.85	9.48
		340291	0.00	4.91	4.82	4.59	10.46
		316728	0.00	5.57	4.65	5.05	8.24
		322797	0.00	6.73	4.35	6.56	8.98
		330573	0.00	6.77	4.61	5.94	13.55
	10	504943	1.83	5.42	3.28	4.27	17.72
		436167	0.55	4.68	3.47	3.50	18.01
		509439	0.00	4.28	2.78	2.74	16.04
		468147	3.36	8.25	3.55	6.57	17.02
		415678	2.62	6.25	3.38	5.07	16.40

Table 1: Experimental results for SWAP, EBSR and EFSR based algorithms

n	q	Best	Simple descent			Dynasearch		
			Best	Avg	time (s)	Best	Avg	time (s)
60	10	178464	10.12	20.86	0.22	6.40	14.92	1.60
		196661	16.13	22.62	0.19	10.68	17.33	1.53
		185427	8.93	18.27	0.21	8.35	14.95	1.45
		165989	15.41	22.35	0.22	12.81	19.97	1.19
		157513	11.71	24.34	0.21	8.75	21.64	1.22

Table 2: Experimental results for SWAP-based algorithms

as a good way to escape from a local optimum. Finally, we have observed that the real key factor of the good performance of our algorithms is the combination — union or composition — of the three neighborhoods SWAP, EBSR and EFSR.

References

- [1] R.K. Ahuja, Ö. Ergun, J.B. Orlin, and A.P. Punnen, *A survey of very large-scale neighborhood search techniques*, Discrete Applied Mathematics **123** (2002), 75–103.
- [2] R.K. Congram, C.N. Potts, and S.L. van de Velde, *An iterated dynasearch algorithm for the single-machine total weighted tardiness scheduling problem*, INFORMS Journal on Computing **14** (2002), 52–67.
- [3] A. Grossso, F. Della Crocce, and R. Tadei, *An enhanced dynasearch neighborhood for the single-machine total weighted tardiness scheduling problem*, Operations Research Letters **32** (2004), 68–72.
- [4] W. Nuijten, T. Bousonville, F. Foccaci, D. Godard, and C. Le Pape, *Towards a real-life manufacturing scheduling problem and test bed*, Proceedings of PMS’04, 2004, submitted.
- [5] F. Sourd, *Earliness-tardiness scheduling with setup considerations*, Computers & Operations Research, to appear.
- [6] _____, *Scheduling a sequence of tasks with general completion costs*, Research report 2002/013, LIP6, 2002.